

# **Building End-to-End MLOps Pipelines with DVC**

Complete MLOps Implementation Guide

A Comprehensive Guide to Building Production-Ready  
ML Pipelines with DVC Automation, Experiment Tracking,  
and AWS Integration

**Sujil S**

[sujil19480@gmail.com](mailto:sujil19480@gmail.com)

December 25, 2025

## Contents

---

<b>1 Introduction to MLOps Pipelines</b>	<b>2</b>
1.1 What is MLOps? . . . . .	2
1.2 Course Agenda . . . . .	2
1.3 Basic ML Pipeline Components . . . . .	2
1.3.1 Pipeline Stages Explained . . . . .	2
1.4 Data Science vs MLOps Practice . . . . .	3
<b>2 Project Setup and Structure</b>	<b>4</b>
2.1 Initial Repository Setup . . . . .	4
2.1.1 Step 1: Create GitHub Repository . . . . .	4
2.1.2 Step 2: Project Directory Structure . . . . .	4
2.2 Initial Experiment Setup . . . . .	4
2.3 Creating the Source Directory . . . . .	5
2.4 Configuring .gitignore . . . . .	5
<b>3 Exploratory Data Analysis Phase</b>	<b>7</b>
3.1 Understanding the Dataset . . . . .	7
3.2 Notebook Workflow Overview . . . . .	7
3.3 Key Code Snippets from Notebook . . . . .	7
3.3.1 Basic Imports and Setup . . . . .	7
3.3.2 Data Loading and Cleaning . . . . .	7
3.3.3 Text Transformation Function . . . . .	8
3.3.4 Feature Engineering with TF-IDF . . . . .	8
3.3.5 Train-Test Split . . . . .	8
3.4 Model Training and Comparison . . . . .	9
3.4.1 Multiple Classifier Setup . . . . .	9
3.4.2 Model Evaluation Function . . . . .	9
<b>4 Data Ingestion Module</b>	<b>11</b>
4.1 Overview . . . . .	11
4.2 Complete Implementation . . . . .	11
4.3 Key Components Explained . . . . .	13
4.3.1 Logging Setup . . . . .	13
4.3.2 Function Breakdown . . . . .	13
4.4 Running the Module . . . . .	13
<b>5 Data Pre-Processing Module</b>	<b>15</b>
5.1 Overview . . . . .	15
5.2 Complete Implementation . . . . .	15
5.3 Text Transformation Pipeline . . . . .	17
5.4 Running the Module . . . . .	18
<b>6 Feature Engineering Module</b>	<b>19</b>
6.1 Overview . . . . .	19
6.2 Complete Implementation . . . . .	19
6.3 Understanding TF-IDF . . . . .	21
6.3.1 Why TF-IDF? . . . . .	21
6.4 max_features Parameter . . . . .	21
6.5 Running the Module . . . . .	22

<b>7 Model Building Module</b>	<b>23</b>
7.1 Overview . . . . .	23
7.2 Complete Implementation . . . . .	23
7.3 Model Selection: RandomForest . . . . .	25
7.4 Hyperparameters Explained . . . . .	25
7.5 Model Persistence with Pickle . . . . .	26
7.6 Running the Module . . . . .	26
<b>8 Model Evaluation Module</b>	<b>27</b>
8.1 Overview . . . . .	27
8.2 Complete Implementation . . . . .	27
8.3 Evaluation Metrics Explained . . . . .	29
8.3.1 Accuracy . . . . .	29
8.3.2 Precision . . . . .	29
8.3.3 Recall (Sensitivity) . . . . .	29
8.3.4 AUC (Area Under ROC Curve) . . . . .	29
8.4 Why These Metrics? . . . . .	30
8.5 Metrics JSON Format . . . . .	30
8.6 Running the Module . . . . .	30
<b>9 Testing the Complete Pipeline Manually</b>	<b>31</b>
9.1 Running All Components Sequentially . . . . .	31
9.2 Expected Directory Structure After Execution . . . . .	31
9.3 Verification Checklist . . . . .	31
9.4 Initial Git Commit . . . . .	32
<b>10 Setting Up DVC Pipeline (Without Parameters)</b>	<b>33</b>
10.1 Understanding DVC Pipelines . . . . .	33
10.2 Initialize DVC . . . . .	33
10.3 Creating dvc.yaml File . . . . .	33
10.4 Understanding dvc.yaml Structure . . . . .	34
10.4.1 Stage Components . . . . .	34
10.5 Stage Dependency Chain . . . . .	35
10.6 Running the DVC Pipeline . . . . .	35
10.7 Understanding dvc.lock File . . . . .	35
10.8 Visualizing the Pipeline . . . . .	36
10.9 Key DVC Concepts . . . . .	37
10.9.1 The .dvc Directory . . . . .	37
10.9.2 dvc.yaml vs dvc.lock vs .dvc files . . . . .	37
10.10 Intelligent Re-Running . . . . .	37
10.11 Committing Pipeline to Git . . . . .	38
<b>11 Adding Configurable Parameters</b>	<b>39</b>
11.1 Why Parameterize? . . . . .	39
11.2 Creating params.yaml . . . . .	39
11.3 Updating dvc.yaml with Parameters . . . . .	39
11.4 Loading Parameters in Python Code . . . . .	40
11.5 Updated Module Implementations . . . . .	40
11.5.1 Data_Ingestion.py Changes . . . . .	41
11.5.2 Feature_Engineering.py Changes . . . . .	41
11.5.3 Model_Building.py Changes . . . . .	42
11.6 Testing Parameterized Pipeline . . . . .	42

11.7 Benefits of Parameterization . . . . .	42
11.8 Commit Parameterized Pipeline . . . . .	43
<b>12 Experiment Tracking with DVCLive</b>	<b>44</b>
12.1 What is DVCLive? . . . . .	44
12.2 Installation . . . . .	44
12.3 Updating Model_Evaluation.py . . . . .	44
12.4 Understanding DVCLive Code . . . . .	45
12.4.1 The Live Context Manager . . . . .	45
12.4.2 Logging Metrics . . . . .	45
12.4.3 Logging Parameters . . . . .	46
12.5 Running Experiments . . . . .	46
12.6 Viewing Experiments . . . . .	46
12.6.1 Command Line . . . . .	46
12.6.2 VS Code Extension . . . . .	47
12.7 The dvclive/ Directory . . . . .	47
12.8 Running Multiple Experiments . . . . .	47
12.8.1 Experiment 1: Baseline . . . . .	47
12.8.2 Experiment 2: More Trees . . . . .	48
12.8.3 Experiment 3: Different Features . . . . .	48
12.9 Comparing Experiments . . . . .	48
12.10 Applying an Experiment . . . . .	48
12.11 Removing Experiments . . . . .	49
12.12 Best Practices for Experimentation . . . . .	49
<b>13 AWS S3 Integration for Remote Storage</b>	<b>50</b>
13.1 Why Remote Storage? . . . . .	50
13.2 AWS Setup Prerequisites . . . . .	50
13.3 Step-by-Step AWS Configuration . . . . .	50
13.3.1 Step 1: Create IAM User . . . . .	50
13.3.2 Step 2: Create S3 Bucket . . . . .	51
13.3.3 Step 3: Install AWS CLI and DVC S3 Support . . . . .	51
13.3.4 Step 4: Configure AWS Credentials . . . . .	51
13.3.5 Step 5: Add S3 Remote to DVC . . . . .	52
13.4 Pushing Data to S3 . . . . .	52
13.4.1 Initial Push . . . . .	52
13.5 Understanding S3 Storage Structure . . . . .	52
13.6 Pulling Data from S3 . . . . .	53
13.6.1 Cloning Project on New Machine . . . . .	53
13.7 Complete Experiment Workflow with S3 . . . . .	53
13.8 Key Concepts: Git vs DVC vs S3 . . . . .	53
13.9 Configuring Remote in .dvc/config . . . . .	54
13.10 Verifying S3 Upload . . . . .	54
13.11 Alternative: Using DVC with Other Storage . . . . .	54
<b>14 What does dvc push actually push?</b>	<b>55</b>
14.1 What Exactly Does dvc push Upload? . . . . .	55
14.1.1 How DVC Decides What to Push . . . . .	55
14.1.2 What Gets Pushed . . . . .	55
14.1.3 What Does NOT Get Pushed . . . . .	55
14.1.4 Example: Multiple Experiments . . . . .	55
14.1.5 Pushing Data for All Commits . . . . .	56

14.2 Summary . . . . .	56
<b>15 Complete MLOps Workflow Summary</b>	<b>57</b>
15.1 Running Multiple Experiments with DVCLive . . . . .	57
15.2 Correct Workflow: Compare and Apply Best Experiment . . . . .	57
15.3 Why This Matters for Rollback . . . . .	58
15.4 Rollback Procedure . . . . .	59
15.5 Complete Command Reference . . . . .	60
15.6 Understanding DVCLive and Local Experiments . . . . .	60
15.7 Workflow Visualization . . . . .	61
15.8 When to Use MLflow . . . . .	62
<b>16 Common Issues and Troubleshooting</b>	<b>63</b>
16.1 DVC Issues . . . . .	63
16.1.1 Issue: "Stage didn't change, skipping" . . . . .	63
16.1.2 Issue: "Failed to push data" . . . . .	63
16.1.3 Issue: "File not found" during dvc pull . . . . .	63
16.2 Git Issues . . . . .	64
16.2.1 Issue: Large files in Git . . . . .	64
16.2.2 Issue: Merge conflicts in dvc.lock . . . . .	64
16.3 Python Issues . . . . .	64
16.3.1 Issue: Import errors . . . . .	64
16.3.2 Issue: NLTK data not found . . . . .	65
16.4 AWS Issues . . . . .	65
16.4.1 Issue: "Access Denied" on S3 . . . . .	65
16.4.2 Issue: Wrong region . . . . .	65
16.5 Experiment Tracking Issues . . . . .	65
16.5.1 Issue: Experiments not showing . . . . .	65
16.5.2 Issue: DVCLive not logging . . . . .	65
<b>17 MLOps Best Practices</b>	<b>66</b>
17.1 Code Organization Best Practices . . . . .	66
17.1.1 Modular Design Principles . . . . .	66
17.1.2 Logging Best Practices . . . . .	67
17.1.3 Error Handling Patterns . . . . .	67
17.2 Data Management Best Practices . . . . .	68
17.2.1 Directory Structure Philosophy . . . . .	68
17.2.2 Data Versioning Strategy . . . . .	68
17.2.3 Data Validation . . . . .	69
17.3 Experiment Management Best Practices . . . . .	70
17.3.1 Naming Conventions . . . . .	70
17.3.2 Parameter Management . . . . .	71
17.3.3 Metric Tracking . . . . .	71
17.4 Collaboration Best Practices . . . . .	72
17.4.1 Documentation Standards . . . . .	72
17.4.2 Code Review Process . . . . .	73
17.4.3 Team Communication . . . . .	73
17.5 Security and Compliance . . . . .	74
17.5.1 Credential Management . . . . .	74
17.5.2 Data Privacy . . . . .	75
17.6 Performance Optimization . . . . .	75
17.6.1 Caching Strategies . . . . .	75

17.6.2	Memory Management . . . . .	76
17.7	Common Pitfalls to Avoid . . . . .	76
<b>18</b>	<b>Quick Reference Guide</b>	<b>78</b>
18.1	Essential DVC Commands . . . . .	78
18.2	Complete Workflow Cheat Sheet . . . . .	79
18.2.1	Initial Project Setup . . . . .	79
18.2.2	Full Experiment Workflow . . . . .	80
18.2.3	Rollback to a Promoted (Committed) Experiment . . . . .	80
18.2.4	AWS S3 Setup Workflow . . . . .	81
18.3	Project Structure Template . . . . .	82
18.4	Key File Purposes . . . . .	83
18.5	Git vs DVC Tracking . . . . .	83
18.6	params.yaml Template . . . . .	83
18.7	dvc.yaml Template . . . . .	84
18.8	Common Command Combinations . . . . .	85
<b>19</b>	<b>Conclusion and Key Takeaways</b>	<b>86</b>
19.1	What You've Accomplished . . . . .	86
19.2	The MLOps Journey: From Notebook to Production . . . . .	86
19.3	Core Principles You've Learned . . . . .	87
19.3.1	1. Version Everything . . . . .	87
19.3.2	2. Automate Repetitive Tasks . . . . .	87
19.3.3	3. Make Everything Reproducible . . . . .	87
19.3.4	4. Track Experiments Systematically . . . . .	87
19.4	The Complete Workflow Visualization . . . . .	88
19.5	Key Lessons Learned . . . . .	88

## 1 Introduction to MLOps Pipelines

### 1.1 What is MLOps?

**MLOps (Machine Learning Operations)** is a set of practices that combines Machine Learning, DevOps, and Data Engineering to deploy and maintain ML systems in production reliably and efficiently.

#### Core MLOps Principles

- **Automation:** Automate ML workflows and deployments
- **Reproducibility:** Ensure experiments can be recreated
- **Versioning:** Track code, data, and models
- **Monitoring:** Track model performance in production
- **Collaboration:** Enable team collaboration on ML projects

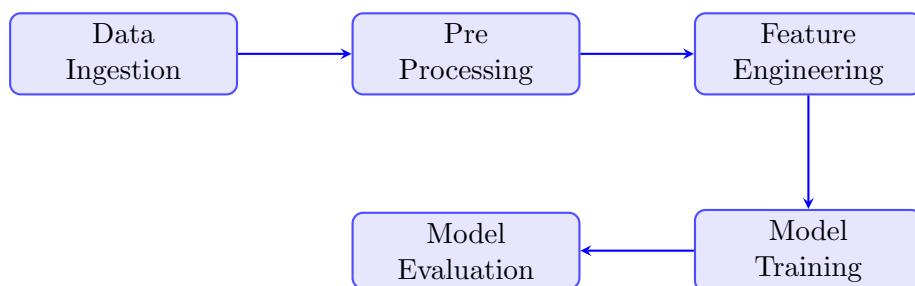
### 1.2 Course Agenda

This comprehensive guide covers the following key topics:

1. **End-to-End ML Pipeline:** Building complete ML workflows with logging and exception handling
2. **DVC Automation:** Automating pipelines using DVC YAML configuration
3. **Parameterization:** Adding configurable parameters for experimentation
4. **Experiment Tracking:** Using DVCLive for tracking experiments
5. **AWS Integration:** Setting up AWS S3 for data versioning

### 1.3 Basic ML Pipeline Components

A typical machine learning pipeline consists of the following stages:



#### 1.3.1 Pipeline Stages Explained

1. **Data Ingestion:** Collecting and loading raw data from various sources
2. **Pre-Processing:** Cleaning, transforming, and preparing data
3. **Feature Engineering:** Creating meaningful features from raw data
4. **Model Training:** Training machine learning models on prepared data
5. **Model Evaluation:** Assessing model performance with metrics

## 1.4 Data Science vs MLOps Practice

Data Science Focus	MLOps Focus
Pre-processing techniques	Coding best practices
Feature engineering methods	Robust pipeline architecture
Model hyperparameter tuning	Experiment tracking
Grid search and optimization	Automation with YAML
Algorithm selection	AWS/Cloud integration
Model performance	Reproducibility and versioning

### Important Note

MLOps extends data science by adding engineering rigor: version control, automation, testing, monitoring, and deployment practices that make ML systems production-ready.

## 2 Project Setup and Structure

### 2.1 Initial Repository Setup

#### 2.1.1 Step 1: Create GitHub Repository

```
# On GitHub: Create new repository named "MLOPS-DVC-Project"
# Clone to local machine
git clone https://github.com/username/MLOPS-DVC-Project.git
cd MLOPS-DVC-Project
```

#### 2.1.2 Step 2: Project Directory Structure

Create the following directory structure:

```
MLOPS-DVC-Project/
  Experiments/
    spam.csv
    mynotebook.ipynb
  src/
    Data_Ingestion.py
    Data_Pre_Processing.py
    Feature_Engineering.py
    Model_Building.py
    Model_Evaluation.py
  data/
    raw/
      interim/
        processed/
  models/
  reports/
  logs/
  .gitignore
  dvc.yaml
  params.yaml
  README.md
```

### 2.2 Initial Experiment Setup

The `Experiments/` folder contains the initial exploration work:

- **spam.csv**: Raw dataset for spam classification
- **mynotebook.ipynb**: Jupyter notebook with exploratory data analysis

#### Purpose of Experiments Folder

The experiments folder serves as a sandbox for:

- Exploratory data analysis (EDA)
- Trying different preprocessing techniques
- Testing various model architectures

- Validating assumptions before production code

## 2.3 Creating the Source Directory

```
# Create src directory
mkdir src
cd src
```

The `src/` directory will contain all production-ready Python scripts:

- Each component as a separate module
- Proper logging configuration
- Exception handling
- Modular and reusable code

## 2.4 Configuring .gitignore

Create a `.gitignore` file to exclude unnecessary files:

### Initial .gitignore Configuration

```
# Data directories
data/
models/
reports/

# Python
__pycache__/
*.pyc
*.pyo
*.pyd
.Python
venv/
env/

# Logs
logs/
*.log

# Jupyter
.ipynb_checkpoints/

# IDE
.vscode/
.idea/

# OS
.DS_Store
Thumbs.db
```

```
# DVC  
/dvc.lock
```

### Warning

**Important:** Always add `data/`, `models/`, and `reports/` to `.gitignore`. DVC will handle versioning these large files, not Git.

## 3 Exploratory Data Analysis Phase

### 3.1 Understanding the Dataset

The spam classification dataset contains SMS messages labeled as spam or ham (not spam).

#### Dataset Structure

##### Columns:

- v1: Label (spam/ham)
- v2: SMS text message
- Unnamed: 2, 3, 4: Empty columns to be dropped

### 3.2 Notebook Workflow Overview

The `mynotebook.ipynb` follows this workflow:

1. **Import Libraries:** NumPy, Pandas, Matplotlib, NLTK
2. **Load Data:** Read CSV file
3. **Data Cleaning:** Drop unnecessary columns, rename columns
4. **Preprocessing:** Encode labels, remove duplicates
5. **Feature Engineering:** Text transformation, TF-IDF
6. **Model Training:** Train multiple classifiers
7. **Model Evaluation:** Compare accuracy and precision

### 3.3 Key Code Snippets from Notebook

#### 3.3.1 Basic Imports and Setup

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5
6 from wordcloud import WordCloud
7 import nltk
8 from nltk.corpus import stopwords
9
10 # Download NLTK data
11 nltk.download('stopwords')
12 nltk.download('punkt')
```

#### 3.3.2 Data Loading and Cleaning

```

1 # Load data
2 df = pd.read_csv('spam.csv')
3
4 # Drop unnecessary columns
```

```

5 df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],
6         inplace=True)
7
8 # Rename columns
9 df.rename(columns={'v1': 'target', 'v2': 'text'}, inplace=True)
10
11 # Check duplicates
12 print(f"Duplicates: {df.duplicated().sum()}")
13 df = df.drop_duplicates(keep='first')

```

### 3.3.3 Text Transformation Function

```

1 from nltk.stem.porter import PorterStemmer
2 import string
3
4 ps = PorterStemmer()
5
6 def transform_text(text):
7     # Lowercase transformation
8     text = text.lower()
9
10    # Tokenization
11    text = nltk.word_tokenize(text)
12
13    # Remove special characters
14    text = [word for word in text if word.isalnum()]
15
16    # Remove stopwords and punctuation
17    text = [word for word in text
18            if word not in stopwords.words('english')
19            and word not in string.punctuation]
20
21    # Stemming
22    text = [ps.stem(word) for word in text]
23
24    return " ".join(text)
25
26 # Apply transformation
27 df['transformed_text'] = df['text'].apply(transform_text)

```

### 3.3.4 Feature Engineering with TF-IDF

```

1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3 # Create TF-IDF features
4 tfidf = TfidfVectorizer(max_features=500)
5 X = tfidf.fit_transform(df['transformed_text']).toarray()
6 y = df['target'].values

```

### 3.3.5 Train-Test Split

```

1 from sklearn.model_selection import train_test_split
2
3 X_train, X_test, y_train, y_test = train_test_split(

```

```

4     X, y, test_size=0.20, random_state=2
5 )

```

## 3.4 Model Training and Comparison

### 3.4.1 Multiple Classifier Setup

```

1 from sklearn.linear_model import LogisticRegression
2 from sklearn.svm import SVC
3 from sklearn.naive_bayes import MultinomialNB
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.neighbors import KNeighborsClassifier
6 from sklearn.ensemble import (RandomForestClassifier,
7                               AdaBoostClassifier,
8                               BaggingClassifier,
9                               ExtraTreesClassifier,
10                              GradientBoostingClassifier)
11 from xgboost import XGBClassifier
12
13 # Initialize classifiers
14 clfs = {
15     'SVC': SVC(kernel="sigmoid", gamma=1.0),
16     'KNN': KNeighborsClassifier(),
17     'NB': MultinomialNB(),
18     'DT': DecisionTreeClassifier(max_depth=5),
19     'LR': LogisticRegression(solver='liblinear', penalty='l1'),
20     'RF': RandomForestClassifier(n_estimators=50, random_state=2),
21     'Adaboost': AdaBoostClassifier(n_estimators=50, random_state=2),
22     'Bgc': BaggingClassifier(n_estimators=50, random_state=2),
23     'ETC': ExtraTreesClassifier(n_estimators=50, random_state=2),
24     'GBDT': GradientBoostingClassifier(n_estimators=50,
25                                         random_state=2),
26     'xgb': XGBClassifier(n_estimators=50, random_state=2)
27 }

```

### 3.4.2 Model Evaluation Function

```

1 from sklearn.metrics import accuracy_score, precision_score
2
3 def train_classifier(clf, X_train, y_train, X_test, y_test):
4     clf.fit(X_train, y_train)
5     y_pred = clf.predict(X_test)
6     accuracy = accuracy_score(y_test, y_pred)
7     precision = precision_score(y_test, y_pred)
8     return accuracy, precision
9
10 # Train and evaluate all classifiers
11 for name, clf in clfs.items():
12     accuracy, precision = train_classifier(
13         clf, X_train, y_train, X_test, y_test
14     )
15     print(f"\nFor: {name}")
16     print(f"Accuracy: {accuracy:.4f}")
17     print(f"Precision: {precision:.4f}")

```

**Important Note**

The notebook phase is exploratory. Once you identify the best approaches, refactor the code into production-ready modules in the `src/` directory with proper logging, error handling, and modularity.

## 4 Data Ingestion Module

### 4.1 Overview

The Data Ingestion module is responsible for:

- Loading raw data from source (URL, file, database)
- Basic preprocessing (dropping columns, renaming)
- Splitting data into train and test sets
- Saving processed data to designated directories

### 4.2 Complete Implementation

src/Data\_Ingestion.py

```
1 import pandas as pd
2 import os
3 from sklearn.model_selection import train_test_split
4 import logging
5
6 # Ensure the "logs" directory exists
7 log_dir = 'logs'
8 os.makedirs(log_dir, exist_ok=True)
9
10 # Logging configuration
11 logger = logging.getLogger('Data_Ingestion')
12 logger.setLevel('DEBUG')
13
14 console_handler = logging.StreamHandler()
15 console_handler.setLevel('DEBUG')
16
17 log_file_path = os.path.join(log_dir, 'Data_Ingestion.log')
18 file_handler = logging.FileHandler(log_file_path)
19 file_handler.setLevel('DEBUG')
20
21 formatter = logging.Formatter(
22     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
23 )
24 console_handler.setFormatter(formatter)
25 file_handler.setFormatter(formatter)
26
27 logger.addHandler(console_handler)
28 logger.addHandler(file_handler)
29
30
31 def load_data(data_url: str) -> pd.DataFrame:
32     """Load data from a CSV file."""
33     try:
34         df = pd.read_csv(data_url)
35         logger.debug('Data loaded from %s', data_url)
36         return df
37     except pd.errors.ParserError as e:
38         logger.error('Failed to parse the CSV file: %s', e)
39         raise
40     except Exception as e:
```

```
41     logger.error('Unexpected error while loading data: %s',
42     e)
43     raise
44
45 def preprocess_data(df: pd.DataFrame) -> pd.DataFrame:
46     """Preprocess the data."""
47     try:
48         df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed:
49             4'],
50                 inplace=True)
51         df.rename(columns={'v1': 'target', 'v2': 'text'},
52                 inplace=True)
53         logger.debug('Data preprocessing completed')
54         return df
55     except KeyError as e:
56         logger.error('Missing column in dataframe: %s', e)
57         raise
58     except Exception as e:
59         logger.error('Unexpected error during preprocessing: %s
60             , e')
61         raise
62
63 def save_data(train_data: pd.DataFrame,
64               test_data: pd.DataFrame,
65               data_path: str) -> None:
66     """Save train and test datasets."""
67     try:
68         raw_data_path = os.path.join(data_path, 'raw')
69         os.makedirs(raw_data_path, exist_ok=True)
70
71         train_data.to_csv(
72             os.path.join(raw_data_path, "train.csv"),
73             index=False
74         )
75         test_data.to_csv(
76             os.path.join(raw_data_path, "test.csv"),
77             index=False
78         )
79         logger.debug('Train and test data saved to %s',
80                     raw_data_path)
81     except Exception as e:
82         logger.error('Error while saving data: %s', e)
83         raise
84
85 def main():
86     try:
87         test_size = 0.2
88         data_url = 'https://raw.githubusercontent.com/...'
89
90         df = load_data(data_url=data_url)
91         final_df = preprocess_data(df)
92
93         train_data, test_data = train_test_split(
```

```
94         final_df, test_size=test_size, random_state=2
95     )
96
97     save_data(train_data, test_data, data_path='./data')
98 except Exception as e:
99     logger.error('Failed to complete data ingestion: %s', e)
100    print(f"Error: {e}")
101
102 if __name__ == '__main__':
103     main()
```

## 4.3 Key Components Explained

### 4.3.1 Logging Setup

- **Logger Name:** Data\_Ingestion
- **Level:** DEBUG (captures all messages)
- **Handlers:**
  - Console handler: Outputs to terminal
  - File handler: Writes to logs/Data\_Ingestion.log
- **Format:** Timestamp, logger name, level, message

### 4.3.2 Function Breakdown

#### 1. load\_data():

- Loads CSV from URL or file path
- Handles parsing errors
- Logs successful load

#### 2. preprocess\_data():

- Drops unnecessary columns
- Renames columns to meaningful names
- Handles missing column errors

#### 3. save\_data():

- Creates directory structure
- Saves train and test CSVs
- Logs save location

#### 4. main():

- Orchestrates the entire process
- Handles top-level exceptions
- Entry point for script execution

## 4.4 Running the Module

```
# Navigate to project root
cd MLOPS-DVC-Project

# Run data ingestion
python src/Data_Ingestion.py
```

**Expected Output:**

- logs/Data\_Ingestion.log created
- data/raw/train.csv created
- data/raw/test.csv created
- Console displays debug messages

**Important Note**

Always test each module individually before integrating into the DVC pipeline. This ensures each component works correctly in isolation.

## 5 Data Pre-Processing Module

### 5.1 Overview

The Data Pre-Processing module handles:

- Label encoding (converting text labels to numeric)
- Removing duplicate entries
- Text transformation (lowercasing, tokenization, stemming)
- Removing stopwords and punctuation

### 5.2 Complete Implementation

#### src/Data\_Pre\_Processing.py - Part 1

```
1 import os
2 import logging
3 import pandas as pd
4 from sklearn.preprocessing import LabelEncoder
5 from nltk.stem.porter import PorterStemmer
6 from nltk.corpus import stopwords
7 import string
8 import nltk
9
10 nltk.download('stopwords')
11 nltk.download('punkt')
12
13 # Ensure the "logs" directory exists
14 log_dir = 'logs'
15 os.makedirs(log_dir, exist_ok=True)
16
17 # Setting up logger
18 logger = logging.getLogger('Data_Pre_Processing')
19 logger.setLevel('DEBUG')
20
21 console_handler = logging.StreamHandler()
22 console_handler.setLevel('DEBUG')
23
24 log_file_path = os.path.join(log_dir, 'Data_Pre_Processing.log')
25 file_handler = logging.FileHandler(log_file_path)
26 file_handler.setLevel('DEBUG')
27
28 formatter = logging.Formatter(
29     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
30 )
31 console_handler.setFormatter(formatter)
32 file_handler.setFormatter(formatter)
33
34 logger.addHandler(console_handler)
35 logger.addHandler(file_handler)
36
37
38 def transform_text(text):
39     """
```

```

40     Transform text: lowercase, tokenize, remove stopwords,
41     punctuation, and stem.
42     """
43
44     ps = PorterStemmer()
45
46     # Convert to lowercase
47     text = text.lower()
48
49     # Tokenize
50     text = nltk.word_tokenize(text)
51
52     # Remove non-alphanumeric tokens
53     text = [word for word in text if word.isalnum()]
54
55     # Remove stopwords and punctuation
56     text = [word for word in text
57             if word not in stopwords.words('english')
58             and word not in string.punctuation]
59
60     # Stem the words
61     text = [ps.stem(word) for word in text]
62
63     # Join back into string
64     return " ".join(text)

```

## src/Data\_Pre\_Processing.py - Part 2

```

1 def preprocess_df(df, text_column='text',
2                   target_column='target'):
3     """
4     Preprocess DataFrame: encode target, remove duplicates,
5     transform text.
6     """
7     try:
8         logger.debug('Starting preprocessing for DataFrame')
9
10        # Encode the target column
11        encoder = LabelEncoder()
12        df[target_column] = encoder.fit_transform(df[
13            target_column])
14        logger.debug('Target column encoded')
15
16        # Remove duplicate rows
17        df = df.drop_duplicates(keep='first')
18        logger.debug('Duplicates removed')
19
20        # Apply text transformation
21        df.loc[:, text_column] = df[text_column].apply(
22            transform_text
23        )
24        logger.debug('Text column transformed')
25
26    return df
27 except KeyError as e:
28     logger.error('Column not found: %s', e)

```

```

28         raise
29     except Exception as e:
30         logger.error('Error during text normalization: %s', e)
31         raise
32
33
34 def main(text_column='text', target_column='target'):
35     """
36     Main function: load raw data, preprocess, save processed.
37     """
38     try:
39         # Load data from data/raw
40         train_data = pd.read_csv('./data/raw/train.csv')
41         test_data = pd.read_csv('./data/raw/test.csv')
42         logger.debug('Data loaded properly')
43
44         # Transform the data
45         train_processed = preprocess_df(train_data,
46                                         text_column,
47                                         target_column)
48         test_processed = preprocess_df(test_data,
49                                         text_column,
50                                         target_column)
51
52         # Store in data/interim
53         data_path = os.path.join("./data", "interim")
54         os.makedirs(data_path, exist_ok=True)
55
56         train_processed.to_csv(
57             os.path.join(data_path, "train_processed.csv"),
58             index=False
59         )
60         test_processed.to_csv(
61             os.path.join(data_path, "test_processed.csv"),
62             index=False
63         )
64
65         logger.debug('Processed data saved to %s', data_path)
66     except FileNotFoundError as e:
67         logger.error('File not found: %s', e)
68     except pd.errors.EmptyDataError as e:
69         logger.error('No data: %s', e)
70     except Exception as e:
71         logger.error('Failed to complete data transformation: %s', e)
72         print(f"Error: {e}")
73
74
75 if __name__ == '__main__':
76     main()

```

### 5.3 Text Transformation Pipeline

The `transform_text()` function implements a comprehensive NLP preprocessing pipeline:

1. **Lowercasing:** Converts all text to lowercase

- "Hello World" → "hello world"
2. **Tokenization:** Splits text into individual words
- "hello world" → ["hello", "world"]
3. **Alphanumeric Filtering:** Removes special characters
- ["hello", "world", "!"] → ["hello", "world"]
4. **Stopword Removal:** Removes common words
- ["the", "quick", "fox"] → ["quick", "fox"]
5. **Stemming:** Reduces words to root form
- ["running", "runs", "ran"] → ["run", "run", "run"]

## 5.4 Running the Module

```
# Run data preprocessing
python src/Data_Pre_Processing.py
```

### Expected Output:

- logs/Data.Pre.Processing.log created
- data/interim/train\_processed.csv created
- data/interim/test\_processed.csv created

### Why Interim Directory?

The data flow follows this structure:

- data/raw/: Original, untouched data
- data/interim/: Partially processed data
- data/processed/: Final features ready for modeling

This separation allows tracking transformations at each stage.

## 6 Feature Engineering Module

### 6.1 Overview

The Feature Engineering module:

- Applies TF-IDF (Term Frequency-Inverse Document Frequency)
- Converts text data into numerical feature vectors
- Creates fixed-size feature matrices for modeling
- Saves processed features with labels

### 6.2 Complete Implementation

src/Feature\_Engineering.py

```

1 import pandas as pd
2 import os
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 import logging
5
6 # Ensure the "logs" directory exists
7 log_dir = 'logs'
8 os.makedirs(log_dir, exist_ok=True)
9
10 # Logging configuration
11 logger = logging.getLogger('Feature_Engineering')
12 logger.setLevel('DEBUG')
13
14 console_handler = logging.StreamHandler()
15 console_handler.setLevel('DEBUG')
16
17 log_file_path = os.path.join(log_dir, 'Feature_Engineering.log')
18 file_handler = logging.FileHandler(log_file_path)
19 file_handler.setLevel('DEBUG')
20
21 formatter = logging.Formatter(
22     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
23 )
24 console_handler.setFormatter(formatter)
25 file_handler.setFormatter(formatter)
26
27 logger.addHandler(console_handler)
28 logger.addHandler(file_handler)
29
30
31 def load_data(file_path: str) -> pd.DataFrame:
32     """Load data from a CSV file."""
33     try:
34         df = pd.read_csv(file_path)
35         df.fillna(' ', inplace=True)
36         logger.debug('Data loaded and NaNs filled from %s',
37                      file_path)
38         return df
39     except pd.errors.ParserError as e:

```

```
40         logger.error('Failed to parse the CSV file: %s', e)
41         raise
42     except Exception as e:
43         logger.error('Unexpected error while loading data: %s',
44                     e)
45         raise
46
47     def apply_tfidf(train_data: pd.DataFrame,
48                     test_data: pd.DataFrame,
49                     max_features: int) -> tuple:
50         """Apply TF-IDF to the data."""
51         try:
52             vectorizer = TfidfVectorizer(max_features=max_features)
53
54             X_train = train_data['text'].values
55             y_train = train_data['target'].values
56             X_test = test_data['text'].values
57             y_test = test_data['target'].values
58
59             X_train_tfidf = vectorizer.fit_transform(X_train)
60             X_test_tfidf = vectorizer.transform(X_test)
61
62             train_df = pd.DataFrame(X_train_tfidf.toarray())
63             train_df['label'] = y_train
64
65             test_df = pd.DataFrame(X_test_tfidf.toarray())
66             test_df['label'] = y_test
67
68             logger.debug('TF-IDF applied and data transformed')
69             return train_df, test_df
70         except Exception as e:
71             logger.error('Error during TF-IDF transformation: %s',
72                         e)
73             raise
74
75     def save_data(df: pd.DataFrame, file_path: str) -> None:
76         """Save the dataframe to a CSV file."""
77         try:
78             os.makedirs(os.path.dirname(file_path), exist_ok=True)
79             df.to_csv(file_path, index=False)
80             logger.debug('Data saved to %s', file_path)
81         except Exception as e:
82             logger.error('Unexpected error while saving data: %s',
83                         e)
84             raise
85
86     def main():
87         try:
88             max_features = 50
89
90             train_data = load_data('./data/interim/train_processed.
91             csv')
92             test_data = load_data('./data/interim/test_processed.
```

```

    csv')

92     train_df, test_df = apply_tfidf(train_data, test_data,
93                                     max_features)

94
95     save_data(train_df,
96                 os.path.join("./data", "processed",
97                             "train_tfidf.csv"))
98     save_data(test_df,
99                 os.path.join("./data", "processed",
100                            "test_tfidf.csv"))
101
102 except Exception as e:
103     logger.error('Failed to complete feature engineering: %s', e)
104     print(f"Error: {e}")
105
106
107 if __name__ == '__main__':
108     main()

```

### 6.3 Understanding TF-IDF

**TF-IDF (Term Frequency-Inverse Document Frequency)** measures the importance of a word in a document relative to a collection of documents.

#### TF-IDF Formula

**TF (Term Frequency):** How often a word appears in a document

$$TF(t, d) = \frac{\text{Number of times term } t \text{ appears in document } d}{\text{Total number of terms in document } d}$$

**IDF (Inverse Document Frequency):** How rare/common a word is across all documents

$$IDF(t) = \log \left( \frac{\text{Total number of documents}}{\text{Number of documents containing term } t} \right)$$

**TF-IDF Score:**

$$\text{TF-IDF}(t, d) = TF(t, d) \times IDF(t)$$

#### 6.3.1 Why TF-IDF?

- **Reduces weight of common words:** Words like "the", "is" get lower scores
- **Increases weight of distinctive words:** Rare words get higher scores
- **Creates fixed-size vectors:** Essential for machine learning models
- **Captures semantic importance:** Better than simple word counts

### 6.4 max\_features Parameter

The `max_features` parameter controls dimensionality:

- **max\_features=50:** Creates 50 features (top 50 important words)

- **Higher values:** More features, more information, higher complexity
- **Lower values:** Fewer features, less information, faster training

#### Important Note

The `max_features` value is a hyperparameter that can be tuned. We'll later make this configurable via `params.yaml` for easy experimentation.

## 6.5 Running the Module

```
# Run feature engineering
python src/Feature_Engineering.py
```

#### Expected Output:

- `logs/Feature_Engineering.log` created
- `data/processed/train_tfidf.csv` created (50 features + 1 label column)
- `data/processed/test_tfidf.csv` created (50 features + 1 label column)

## 7 Model Building Module

### 7.1 Overview

The Model Building module:

- Loads processed feature data
- Trains a RandomForest classifier
- Saves the trained model as a pickle file
- Logs the training process

### 7.2 Complete Implementation

src/Model\_Building.py

```
1 import os
2 import numpy as np
3 import pandas as pd
4 import pickle
5 import logging
6 from sklearn.ensemble import RandomForestClassifier
7
8 # Ensure the "logs" directory exists
9 log_dir = 'logs'
10 os.makedirs(log_dir, exist_ok=True)
11
12 # Logging configuration
13 logger = logging.getLogger('Model_Building')
14 logger.setLevel('DEBUG')
15
16 console_handler = logging.StreamHandler()
17 console_handler.setLevel('DEBUG')
18
19 log_file_path = os.path.join(log_dir, 'Model_Building.log')
20 file_handler = logging.FileHandler(log_file_path)
21 file_handler.setLevel('DEBUG')
22
23 formatter = logging.Formatter(
24     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
25 )
26 console_handler.setFormatter(formatter)
27 file_handler.setFormatter(formatter)
28
29 logger.addHandler(console_handler)
30 logger.addHandler(file_handler)
31
32
33 def load_data(file_path: str) -> pd.DataFrame:
34     """Load data from a CSV file."""
35     try:
36         df = pd.read_csv(file_path)
37         logger.debug('Data loaded from %s with shape %s',
38                      file_path, df.shape)
39         return df
40     except pd.errors.ParserError as e:
```

```

41         logger.error('Failed to parse the CSV file: %s', e)
42         raise
43     except FileNotFoundError as e:
44         logger.error('File not found: %s', e)
45         raise
46     except Exception as e:
47         logger.error('Unexpected error while loading data: %s',
48                     e)
49         raise
50
51 def train_model(X_train: np.ndarray,
52                 y_train: np.ndarray,
53                 params: dict) -> RandomForestClassifier:
54     """Train the RandomForest model."""
55     try:
56         if X_train.shape[0] != y_train.shape[0]:
57             raise ValueError(
58                 "Number of samples in X_train and y_train must
59                 match")
60
61         logger.debug('Initializing RandomForest with params: %s
62                 ,',
63                 params)
64         clf = RandomForestClassifier(
65             n_estimators=params['n_estimators'],
66             random_state=params['random_state']
67         )
68
69         logger.debug('Model training started with %d samples',
70                     X_train.shape[0])
71         clf.fit(X_train, y_train)
72         logger.debug('Model training completed')
73
74         return clf
75     except ValueError as e:
76         logger.error('ValueError during model training: %s', e)
77         raise
78     except Exception as e:
79         logger.error('Error during model training: %s', e)
80         raise
81
82 def save_model(model, file_path: str) -> None:
83     """Save the trained model to a file."""
84     try:
85         os.makedirs(os.path.dirname(file_path), exist_ok=True)
86
87         with open(file_path, 'wb') as file:
88             pickle.dump(model, file)
89         logger.debug('Model saved to %s', file_path)
90     except FileNotFoundError as e:
91         logger.error('File path not found: %s', e)
92         raise
93     except Exception as e:

```

```

94         logger.error('Error while saving model: %s', e)
95         raise
96
97
98     def main():
99         try:
100             params = {'n_estimators': 25, 'random_state': 2}
101
102             train_data = load_data('./data/processed/train_tfidf.
103             csv')
103             X_train = train_data.iloc[:, :-1].values
104             y_train = train_data.iloc[:, -1].values
105
106             clf = train_model(X_train, y_train, params)
107
108             model_save_path = 'models/model.pkl'
109             save_model(clf, model_save_path)
110         except Exception as e:
111             logger.error('Failed to complete model building: %s', e
112         )
112         print(f"Error: {e}")
113
114
115 if __name__ == '__main__':
116     main()

```

### 7.3 Model Selection: RandomForest

#### Why RandomForest?

- **Robust:** Handles overfitting well through ensemble learning
- **Feature Importance:** Provides insights into which features matter
- **No Feature Scaling:** Works well without normalization
- **Handles Imbalanced Data:** Good for spam/ham classification
- **Parallel Training:** Can utilize multiple CPU cores

### 7.4 Hyperparameters Explained

- **n\_estimators:** Number of decision trees in the forest

- More trees = Better performance but slower training
- Typical range: 50-200

- **random\_state:** Ensures reproducibility

- Same value = Same results across runs
- Essential for experiment tracking

## 7.5 Model Persistence with Pickle

**Pickle** serializes Python objects to disk:

- **Saves entire model:** Including learned parameters
- **Quick loading:** No need to retrain
- **Version control friendly:** When combined with DVC

### Warning

#### Important .gitignore Rule:

Add `models/` to `.gitignore!` Git should NOT track model files (they're binary and large). DVC will handle model versioning.

```
# In .gitignore
models/
```

## 7.6 Running the Module

```
# Run model building
python src/Model_Building.py
```

#### Expected Output:

- `logs/Model_Building.log` created
- `models/model.pkl` created
- Console displays training progress

## 8 Model Evaluation Module

### 8.1 Overview

The Model Evaluation module:

- Loads the trained model
- Loads test data
- Makes predictions
- Calculates evaluation metrics
- Saves metrics to JSON file

### 8.2 Complete Implementation

src/Model\_Evaluation.py

```
1 import os
2 import numpy as np
3 import pandas as pd
4 import pickle
5 import json
6 from sklearn.metrics import (accuracy_score, precision_score,
7                               recall_score, roc_auc_score)
8 import logging
9
10 # Ensure the "logs" directory exists
11 log_dir = 'logs'
12 os.makedirs(log_dir, exist_ok=True)
13
14 # Logging configuration
15 logger = logging.getLogger('Model_Evaluation')
16 logger.setLevel('DEBUG')
17
18 console_handler = logging.StreamHandler()
19 console_handler.setLevel('DEBUG')
20
21 log_file_path = os.path.join(log_dir, 'Model_Evaluation.log')
22 file_handler = logging.FileHandler(log_file_path)
23 file_handler.setLevel('DEBUG')
24
25 formatter = logging.Formatter(
26     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
27 )
28 console_handler.setFormatter(formatter)
29 file_handler.setFormatter(formatter)
30
31 logger.addHandler(console_handler)
32 logger.addHandler(file_handler)
33
34
35 def load_model(file_path: str):
36     """Load the trained model from a file."""
37     try:
38         with open(file_path, 'rb') as file:
```

```
39         model = pickle.load(file)
40     logger.debug('Model loaded from %s', file_path)
41     return model
42 except FileNotFoundError:
43     logger.error('File not found: %s', file_path)
44     raise
45 except Exception as e:
46     logger.error('Unexpected error while loading model: %s',
47 , e)
48     raise
49
50 def load_data(file_path: str) -> pd.DataFrame:
51     """Load data from a CSV file."""
52     try:
53         df = pd.read_csv(file_path)
54         logger.debug('Data loaded from %s', file_path)
55         return df
56     except pd.errors.ParserError as e:
57         logger.error('Failed to parse the CSV file: %s', e)
58         raise
59     except Exception as e:
60         logger.error('Unexpected error while loading data: %s',
61 , e)
62         raise
63
64 def evaluate_model(clf, X_test: np.ndarray,
65                     y_test: np.ndarray) -> dict:
66     """Evaluate the model and return metrics."""
67     try:
68         y_pred = clf.predict(X_test)
69         y_pred_proba = clf.predict_proba(X_test)[:, 1]
70
71         accuracy = accuracy_score(y_test, y_pred)
72         precision = precision_score(y_test, y_pred)
73         recall = recall_score(y_test, y_pred)
74         auc = roc_auc_score(y_test, y_pred_proba)
75
76         metrics_dict = {
77             'accuracy': accuracy,
78             'precision': precision,
79             'recall': recall,
80             'auc': auc
81         }
82         logger.debug('Model evaluation metrics calculated')
83         return metrics_dict
84     except Exception as e:
85         logger.error('Error during model evaluation: %s', e)
86         raise
87
88
89 def save_metrics(metrics: dict, file_path: str) -> None:
90     """Save the evaluation metrics to a JSON file."""
91     try:
92         os.makedirs(os.path.dirname(file_path), exist_ok=True)
```

```

93     with open(file_path, 'w') as file:
94         json.dump(metrics, file, indent=4)
95     logger.debug('Metrics saved to %s', file_path)
96 except Exception as e:
97     logger.error('Error while saving metrics: %s', e)
98     raise
99
100
101
102 def main():
103     try:
104         clf = load_model('./models/model.pkl')
105         test_data = load_data('./data/processed/test_tfidf.csv',
106 )
107         X_test = test_data.iloc[:, :-1].values
108         y_test = test_data.iloc[:, -1].values
109
110         metrics = evaluate_model(clf, X_test, y_test)
111
112         save_metrics(metrics, 'reports/metrics.json')
113     except Exception as e:
114         logger.error('Failed to complete model evaluation: %s',
115 e)
116         print(f"Error: {e}")
117
118 if __name__ == '__main__':
119     main()

```

## 8.3 Evaluation Metrics Explained

### 8.3.1 Accuracy

Percentage of correct predictions:

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

### 8.3.2 Precision

Of all predicted spam, how many are actually spam:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

### 8.3.3 Recall (Sensitivity)

Of all actual spam, how many did we catch:

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

### 8.3.4 AUC (Area Under ROC Curve)

Measures model's ability to distinguish between classes:

- AUC = 1.0: Perfect classifier
- AUC = 0.5: Random guessing
- Higher is better

#### 8.4 Why These Metrics?

Metric	When to Prioritize
Accuracy	Balanced datasets, overall performance
Precision	When false positives are costly (e.g., marking legitimate emails as spam)
Recall	When false negatives are costly (e.g., missing actual spam)
AUC	Overall model discrimination ability, imbalanced datasets

#### 8.5 Metrics JSON Format

##### reports/metrics.json Example

```
{
    "accuracy": 0.9732,
    "precision": 0.9821,
    "recall": 0.9456,
    "auc": 0.9887
}
```

##### Warning

##### Add reports/ to .gitignore:

```
# In .gitignore
reports/
```

DVC will track metrics files, not Git.

#### 8.6 Running the Module

```
# Run model evaluation
python src/Model_Evaluation.py
```

##### Expected Output:

- logs/Model\_Evaluation.log created
- reports/metrics.json created with all metrics
- Console displays evaluation progress

## 9 Testing the Complete Pipeline Manually

### 9.1 Running All Components Sequentially

Before automating with DVC, verify each component works correctly:

```
# Step 1: Data Ingestion
python src/Data_Ingestion.py

# Step 2: Data Pre-Processing
python src/Data_Pre_Processing.py

# Step 3: Feature Engineering
python src/Feature_Engineering.py

# Step 4: Model Building
python src/Model_Building.py

# Step 5: Model Evaluation
python src/Model_Evaluation.py
```

### 9.2 Expected Directory Structure After Execution

```
MLOPS-DVC-Project/
  data/
    raw/
      train.csv
      test.csv
    interim/
      train_processed.csv
      test_processed.csv
    processed/
      train_tfidf.csv
      test_tfidf.csv
  models/
    model.pkl
  reports/
    metrics.json
  logs/
    Data_Ingestion.log
    Data_Pre_Processing.log
    Feature_Engineering.log
    Model_Building.log
    Model_Evaluation.log
```

### 9.3 Verification Checklist

1. **Data Files:** All CSV files created in correct directories
2. **Model File:** model.pkl exists in models/
3. **Metrics File:** metrics.json contains all 4 metrics

4. **Log Files:** Each component has its log file
5. **No Errors:** Check log files for any ERROR messages

## 9.4 Initial Git Commit

```
# Check .gitignore is properly configured
cat .gitignore

# Stage all source code
git add .gitignore
git add src/
git add Experiments/
git add README.md

# Commit
git commit -m "Initial commit: Add all pipeline components"

# Push to remote
git push origin main
```

### Important Note

At this point, Git tracks:

- Source code (`src/`)
- Experiments (`Experiments/`)
- Configuration files (`.gitignore`)

Git does NOT track:

- Data files (`data/`)
- Models (`models/`)
- Reports (`reports/`)
- Logs (`logs/`)

DVC will handle these large files in the next section.

## 10 Setting Up DVC Pipeline (Without Parameters)

### 10.1 Understanding DVC Pipelines

DVC Pipelines automate ML workflows by:

- Defining stages and dependencies
- Tracking inputs and outputs
- Automatically detecting changes
- Running only what's necessary
- Creating reproducible workflows

### 10.2 Initialize DVC

```
# Initialize DVC in the project
dvc init
```

This creates:

- .dvc/ directory: DVC configuration and cache
- .dvcignore: Similar to .gitignore for DVC
- Updates .gitignore to exclude DVC cache

### 10.3 Creating dvc.yaml File

The `dvc.yaml` file defines the entire pipeline:

#### dvc.yaml (Basic Version)

```
1 stages:
2   data_ingestion:
3     cmd: python src/Data_Ingestion.py
4     deps:
5       - src/Data_Ingestion.py
6     outs:
7       - data/raw
8
9   data_preprocessing:
10    cmd: python src/Data_Pre_Processing.py
11    deps:
12      - data/raw
13      - src/Data_Pre_Processing.py
14    outs:
15      - data/interim
16
17  feature_engineering:
18    cmd: python src/Feature_Engineering.py
19    deps:
20      - data/interim
21      - src/Feature_Engineering.py
22    outs:
23      - data/processed
```

```
24
25     model_building:
26         cmd: python src/Model_Building.py
27         deps:
28             - data/processed
29             - src/Model_Building.py
30         outs:
31             - models/model.pkl
32
33     model_evaluation:
34         cmd: python src/Model_Evaluation.py
35         deps:
36             - models/model.pkl
37             - src/Model_Evaluation.py
38         metrics:
39             - reports/metrics.json
```

## 10.4 Understanding dvc.yaml Structure

### 10.4.1 Stage Components

Each stage has:

1. **cmd**: Command to execute

```
1 cmd: python src/Data_Ingestion.py
2
```

2. **deps**: Dependencies (files that, if changed, trigger re-run)

```
1 deps:
2   - src/Data_Ingestion.py
3   - data/raw
4
```

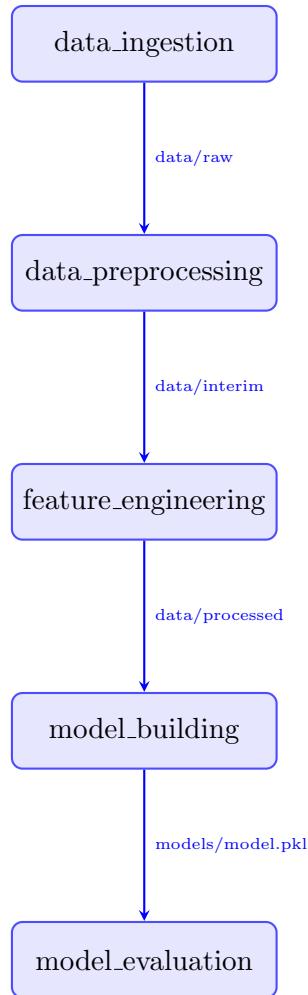
3. **outs**: Outputs (files/directories created by stage)

```
1 outs:
2   - data/interim
3
```

4. **metrics**: Metric files (special output for tracking)

```
1 metrics:
2   - reports/metrics.json
3
```

## 10.5 Stage Dependency Chain



## 10.6 Running the DVC Pipeline

```
# Run the entire pipeline
dvc repro
```

**What happens during dvc repro:**

1. DVC reads `dvc.yaml`
2. Checks if dependencies changed
3. Runs necessary stages in order
4. Tracks all outputs
5. Creates `dvc.lock` file

## 10.7 Understanding `dvc.lock` File

After running `dvc repro`, DVC creates `dvc.lock`:

### What is dvc.lock?

`dvc.lock` is analogous to `package-lock.json` or `requirements.txt.lock`:

- **Locks exact versions:** Stores MD5 hashes of all files
- **Ensures reproducibility:** Same lock = same results
- **Tracked by Git:** Commit this file!
- **Auto-generated:** Never edit manually

### Sample `dvc.lock` Entry

```

1 data_ingestion:
2   cmd: python src/Data_Ingestion.py
3   deps:
4     - path: src/Data_Ingestion.py
5       md5: a1b2c3d4e5f6g7h8i9j0
6   outs:
7     - path: data/raw
8       md5: x1y2z3a4b5c6d7e8f9g0

```

## 10.8 Visualizing the Pipeline

```
# Visualize pipeline as a DAG (Directed Acyclic Graph)
dvc dag
```

### Example Output:

```

+-----+
| data_ingestion      |
+-----+
*
*
*
+-----+
| data_preprocessing  |
+-----+
*
*
*
+-----+
| feature_engineering |
+-----+
*
*
*
+-----+
| model_building      |
+-----+
*
*
```

```

*  

+-----+  

| model_evaluation |  

+-----+

```

## 10.9 Key DVC Concepts

### 10.9.1 The .dvc Directory

- **.dvc/cache/**: Stores all versions of tracked data
- **.dvc/config**: DVC configuration (remotes, etc.)
- **.dvc/.gitignore**: Prevents Git from tracking cache

#### Cache Structure

The `.dvc/cache/` contains:

- **Both tokens AND data**: Unlike remote storage
- **All versions**: Every experiment's data
- **Content-addressable**: Files named by MD5 hash
- **Deduplicated**: Identical files stored once

### 10.9.2 dvc.yaml vs dvc.lock vs .dvc files

File	Purpose
dvc.yaml	Defines pipeline stages, dependencies, outputs
dvc.lock	Locks exact file versions (MD5 hashes) for reproducibility
<name>.dvc	Created by <code>dvc add</code> , tracks individual files/folders

#### Warning

##### Key Difference:

- **dvc add**: Manually track files, creates `<name>.dvc`
- **dvc.yaml outputs**: Automatically tracked, uses `dvc.lock`

With pipelines, you DON'T need `dvc add` for outputs!

## 10.10 Intelligent Re-Running

DVC only re-runs changed stages:

```
# If you run dvc repro again without changes
dvc repro

# Output:
Stage 'data_ingestion' didn't change, skipping
Stage 'data_preprocessing' didn't change, skipping
Stage 'feature_engineering' didn't change, skipping
```

```
Stage 'model_building' didn't change, skipping
Stage 'model_evaluation' didn't change, skipping
Data and pipelines are up to date.
```

### When does DVC re-run a stage?

- **Code changes:** Modify `src/` files
- **Dependency changes:** Input data changes
- **Parameter changes:** (covered in next section)
- **Manual force:** `dvc repro --force`

## 10.11 Committing Pipeline to Git

```
# Stage DVC files
git add dvc.yaml dvc.lock .dvc/.gitignore .dvc/config

# Commit
git commit -m "Add DVC pipeline automation"

# Push
git push origin main
```

### Important Note

#### What's in Git vs DVC now?

##### Git tracks:

- Source code (`src/`)
- Pipeline definition (`dvc.yaml`)
- File version locks (`dvc.lock`)
- DVC configuration (`.dvc/config`)

##### DVC tracks (in `.dvc/cache/`):

- Data (`data/`)
- Models (`models/`)
- Reports (`reports/`)

## 11 Adding Configurable Parameters

### 11.1 Why Parameterize?

Hardcoded values in code are problematic:

- **Not experiment-friendly:** Changing values requires code edits
- **Not version-controlled:** Hard to track what values were used
- **Not DVC-aware:** DVC can't detect parameter changes
- **Not reproducible:** Unclear what parameters produced results

**Solution:** Centralize parameters in `params.yaml`

### 11.2 Creating params.yaml

#### params.yaml

```
1 data_ingestion:
2   test_size: 0.15
3
4 feature_engineering:
5   max_features: 45
6
7 model_building:
8   n_estimators: 20
9   random_state: 2
```

### 11.3 Updating dvc.yaml with Parameters

#### dvc.yaml (With Parameters)

```
1 stages:
2   data_ingestion:
3     cmd: python src/Data_Ingestion.py
4     deps:
5       - src/Data_Ingestion.py
6     params:
7       - data_ingestion.test_size
8     outs:
9       - data/raw
10
11   data_preprocessing:
12     cmd: python src/Data_Pre_Processing.py
13     deps:
14       - data/raw
15       - src/Data_Pre_Processing.py
16     outs:
17       - data/interim
18
19   feature_engineering:
20     cmd: python src/Feature_Engineering.py
21     deps:
```

```

22     - data/interim
23     - src/Feature_Engineering.py
24   params:
25     - feature_engineering.max_features
26   outs:
27     - data/processed
28
29   model_building:
30     cmd: python src/Model_Building.py
31   deps:
32     - data/processed
33     - src/Model_Building.py
34   params:
35     - model_building.n_estimators
36     - model_building.random_state
37   outs:
38     - models/model.pkl
39
40   model_evaluation:
41     cmd: python src/Model_Evaluation.py
42   deps:
43     - models/model.pkl
44     - src/Model_Evaluation.py
45   metrics:
46     - reports/metrics.json

```

## 11.4 Loading Parameters in Python Code

Add a `load_params()` function to each module:

### Parameter Loading Function

```

1 import yaml
2
3 def load_params(params_path: str) -> dict:
4     """Load parameters from a YAML file."""
5     try:
6         with open(params_path, 'r') as file:
7             params = yaml.safe_load(file)
8             logger.debug('Parameters retrieved from %s',
params_path)
9         return params
10    except FileNotFoundError:
11        logger.error('File not found: %s', params_path)
12        raise
13    except yaml.YAMLError as e:
14        logger.error('YAML error: %s', e)
15        raise
16    except Exception as e:
17        logger.error('Unexpected error: %s', e)
18        raise

```

## 11.5 Updated Module Implementations

### 11.5.1 Data\_Ingestion.py Changes

```

1 # Add at top
2 import yaml
3
4 # Add load_params function (shown above)
5
6 # Update main()
7 def main():
8     try:
9         # Load parameters
10        params = load_params(params_path='params.yaml')
11        test_size = params['data_ingestion']['test_size']
12
13        data_url = 'https://raw.githubusercontent.com/...'
14        df = load_data(data_url=data_url)
15        final_df = preprocess_data(df)
16
17        # Use parameterized test_size
18        train_data, test_data = train_test_split(
19            final_df, test_size=test_size, random_state=2
20        )
21
22        save_data(train_data, test_data, data_path='./data')
23    except Exception as e:
24        logger.error('Failed to complete data ingestion: %s', e)
25        print(f"Error: {e}")

```

### 11.5.2 Feature\_Engineering.py Changes

```

1 # Add at top
2 import yaml
3
4 # Add load_params function
5
6 # Update main()
7 def main():
8     try:
9         # Load parameters
10        params = load_params(params_path='params.yaml')
11        max_features = params['feature_engineering']['max_features']
12
13        train_data = load_data('./data/interim/train_processed.csv')
14        test_data = load_data('./data/interim/test_processed.csv')
15
16        # Use parameterized max_features
17        train_df, test_df = apply_tfidf(train_data, test_data,
18                                        max_features)
19
20        save_data(train_df, './data/processed/train_tfidf.csv')
21        save_data(test_df, './data/processed/test_tfidf.csv')
22    except Exception as e:
23        logger.error('Failed to complete feature engineering: %s', e)
24        print(f"Error: {e}")

```

### 11.5.3 Model\_Building.py Changes

```

1 # Add at top
2 import yaml
3
4 # Add load_params function
5
6 # Update main()
7 def main():
8     try:
9         # Load parameters
10        params = load_params('params.yaml')['model_building']
11
12        train_data = load_data('./data/processed/train_tfidf.csv')
13        X_train = train_data.iloc[:, :-1].values
14        y_train = train_data.iloc[:, -1].values
15
16        # Use parameterized model parameters
17        clf = train_model(X_train, y_train, params)
18
19        model_save_path = 'models/model.pkl'
20        save_model(clf, model_save_path)
21    except Exception as e:
22        logger.error('Failed to complete model building: %s', e)
23        print(f"Error: {e}")

```

## 11.6 Testing Parameterized Pipeline

```
# Run pipeline with new parameters
dvc repro
```

DVC now detects parameter changes:

- If you change params.yaml, DVC knows
- Only affected stages re-run
- Parameters tracked in dvc.lock

## 11.7 Benefits of Parameterization

### 1. Easy Experimentation:

```

1 # Change this:
2 model_building:
3     n_estimators: 20
4
5 # To this:
6 model_building:
7     n_estimators: 50
8
9 # Then run: dvc repro
10
```

2. **DVC Awareness:** DVC detects changes and re-runs appropriate stages
3. **Version Control:** Parameters tracked with code in Git

4. **Reproducibility:** params.yaml + dvc.lock = exact reproduction

5. **Documentation:** Clear record of hyperparameters used

## 11.8 Commit Parameterized Pipeline

```
# Stage changes
git add params.yaml dvc.yaml dvc.lock src/

# Commit
git commit -m "Add parameter configuration with params.yaml"

# Push
git push origin main
```

### Important Note

#### Important Behavior:

If you run dvc repro without any changes:

```
Stage 'data_ingestion' didn't change, skipping
Stage 'data_preprocessing' didn't change, skipping
Stage 'feature_engineering' didn't change, skipping
Stage 'model_building' didn't change, skipping
Stage 'model_evaluation' didn't change, skipping
Data and pipelines are up to date.
```

DVC is smart enough to skip unnecessary work!

## 12 Experiment Tracking with DVCLive

### 12.1 What is DVCLive?

DVCLive is DVC's experiment tracking library:

- **Logs metrics:** Accuracy, loss, custom metrics
- **Logs parameters:** Hyperparameters used
- **Creates plots:** Training curves, confusion matrices
- **Integrates with DVC:** Automatic experiment versioning
- **Lightweight:** No external services needed

### 12.2 Installation

```
pip install dvclive
```

### 12.3 Updating Model\_Evaluation.py

Add DVCLive integration to track experiments:

#### Model\_Evaluation.py with DVCLive

```

1 import os
2 import numpy as np
3 import pandas as pd
4 import pickle
5 import json
6 from sklearn.metrics import (accuracy_score, precision_score,
7                               recall_score, roc_auc_score)
8 import logging
9 from dvclive import Live # Add DVCLive
10 import yaml # Add YAML
11
12 # ... (previous logging setup remains same) ...
13
14 # Add load_params function
15 def load_params(params_path: str) -> dict:
16     """Load parameters from a YAML file."""
17     try:
18         with open(params_path, 'r') as file:
19             params = yaml.safe_load(file)
20             logger.debug('Parameters retrieved from %s',
21                         params_path)
22             return params
23     except FileNotFoundError:
24         logger.error('File not found: %s', params_path)
25         raise
26     except yaml.YAMLError as e:
27         logger.error('YAML error: %s', e)
28         raise
29     except Exception as e:
30         logger.error('Unexpected error: %s', e)
31         raise

```

```

31
32 # ... (other functions remain same) ...
33
34 def main():
35     try:
36         # Load parameters
37         params = load_params('params.yaml')
38
39         clf = load_model('./models/model.pkl')
40         test_data = load_data('./data/processed/test_tfidf.csv',
41 )
42
43         X_test = test_data.iloc[:, :-1].values
44         y_test = test_data.iloc[:, -1].values
45
46         metrics = evaluate_model(clf, X_test, y_test)
47         save_metrics(metrics, 'reports/metrics.json')
48
49         # DVCLive logging
50         with Live(save_dvc_exp=True) as live:
51             # Log metrics
52             live.log_metric('accuracy', metrics['accuracy'])
53             live.log_metric('precision', metrics['precision'])
54             live.log_metric('recall', metrics['recall'])
55             live.log_metric('auc', metrics['auc'])
56
57             # Log parameters
58             live.log_params(params)
59
60         except Exception as e:
61             logger.error('Failed to complete model evaluation: %s',
62             e)
63             print(f"Error: {e}")
64
65 if __name__ == '__main__':
66     main()

```

## 12.4 Understanding DVCLive Code

### 12.4.1 The Live Context Manager

```

1 with Live(save_dvc_exp=True) as live:
2     # Logging code here

```

- **save\_dvc\_exp=True:** Saves experiment in DVC
- **Context manager:** Automatically handles setup/cleanup
- **Creates dvclive/ directory:** Stores experiment data

### 12.4.2 Logging Metrics

```

1 live.log_metric('accuracy', metrics['accuracy'])
2 live.log_metric('precision', metrics['precision'])

```

Each metric is:

- Recorded in `dvclive/` folder
- Tracked by DVC for comparison
- Viewable in experiment tables

### 12.4.3 Logging Parameters

```
1 live.log_params(params)
```

Logs all parameters from `params.yaml`:

- Test size
- Max features
- Model hyperparameters

## 12.5 Running Experiments

```
# Run an experiment (instead of dvc repro)
dvc exp run
```

**What happens:**

1. Pipeline executes
2. DVCLive logs metrics and parameters
3. `dvclive/` directory created
4. Experiment saved in DVC's experiment database
5. Unique experiment ID generated

## 12.6 Viewing Experiments

### 12.6.1 Command Line

```
# Show all experiments
dvc exp show
```

**Example Output:**

```
=====
Experiment      accuracy  precision  recall   auc
=====
workspace      0.9732    0.9821    0.9456   0.9887
exp-a1b2c      0.9651    0.9745    0.9382   0.9801
exp-c3d4e      0.9689    0.9778    0.9421   0.9845
exp-e5f6g      0.9712    0.9802    0.9444   0.9869
=====
```

### 12.6.2 VS Code Extension

Install DVC Extension in VS Code:

1. Open VS Code
2. Go to Extensions (Ctrl+Shift+X)
3. Search "DVC"
4. Install "DVC" by Iterative

**Features:**

- Visual experiment comparison table
- Click to view experiment details
- Sort by metrics
- Apply experiments to workspace

## 12.7 The dvclive/ Directory

After running `dvc exp run`, a `dvclive/` directory is created:

```
dvclive/
  metrics.json
  params.yaml
  plots/
```

**Important: dvclive/ is Temporary!**

**Key Concept:** The `dvclive/` folder is like a whiteboard:

- **Current experiment only:** Contains latest run's data
- **Gets overwritten:** Next run replaces contents
- **DVC saves copies:** Each experiment snapshot saved internally
- **Don't add to Git:** Let DVC handle it

Think of it as:

- `dvclive/` = temporary workspace
- DVC's internal storage = permanent photo album
- `dvc exp show` = views the photo album, not the workspace

## 12.8 Running Multiple Experiments

### 12.8.1 Experiment 1: Baseline

```
1 # params.yaml
2 model_building:
3   n_estimators: 20
4   random_state: 2
```

```
dvc exp run
```

### 12.8.2 Experiment 2: More Trees

```
1 # params.yaml
2 model_building:
3   n_estimators: 50
4   random_state: 2
```

```
dvc exp run
```

### 12.8.3 Experiment 3: Different Features

```
1 # params.yaml
2 feature_engineering:
3   max_features: 100
4
5 model_building:
6   n_estimators: 50
7   random_state: 2
```

```
dvc exp run
```

## 12.9 Comparing Experiments

```
# Show all experiments with metrics
dvc exp show

# Show differences between experiments
dvc exp diff exp-a1b2c exp-c3d4e
```

## 12.10 Applying an Experiment

If Experiment 2 has the best metrics:

```
# Apply experiment to workspace
dvc exp apply exp-c3d4e
```

**What happens:**

- Workspace updated with that experiment's code and parameters
- Files restored to that experiment's state
- **Note:** Experiment ID no longer exists after applying

### Warning

#### Critical Understanding:

`dvc exp apply` does NOT create a Git commit!  
 It's temporary—like copying experiment → workspace.  
 To make it permanent:

1. Apply experiment: `dvc exp apply <exp-id>`
2. Commit to Git: `git add . && git commit`

Experiments are NOT Git commits—they’re separate DVC entities.

## 12.11 Removing Experiments

```
# Remove specific experiment
dvc exp remove exp-a1b2c

# Remove all experiments except workspace
dvc exp gc
```

## 12.12 Best Practices for Experimentation

1. **Run experiments frequently:** After each significant change
2. **Use descriptive parameters:** Clear naming in `params.yaml`
3. **Compare before committing:** Use `dvc exp show` to choose best
4. **Apply + Commit best experiment:** Make it permanent in Git
5. **Clean up:** Remove failed/uninteresting experiments

## 13 AWS S3 Integration for Remote Storage

---

### 13.1 Why Remote Storage?

Local DVC cache (`.dvc/cache/`) limitations:

- **Machine-specific:** Lost if machine crashes
- **Not shareable:** Team members can't access
- **Storage constraints:** Limited by local disk space
- **No backup:** Single point of failure

Remote storage (S3) benefits:

- **Cloud backup:** Data safe in AWS
- **Team collaboration:** Everyone accesses same data
- **Scalable:** Unlimited storage
- **Reproducible:** Pull exact data versions anywhere

### 13.2 AWS Setup Prerequisites

1. AWS Account
2. IAM User with S3 permissions
3. Access Key ID
4. Secret Access Key

### 13.3 Step-by-Step AWS Configuration

#### 13.3.1 Step 1: Create IAM User

1. Log into AWS Console
2. Navigate to IAM → Users
3. Click "Create User"
4. Set username (e.g., `dvc-user`)
5. Enable "Programmatic access"
6. Attach policy: `AmazonS3FullAccess`
7. Save Access Key ID and Secret Access Key

#### Warning

##### Security Warning:

- NEVER commit AWS credentials to Git
- Store securely in environment variables

- Use IAM roles in production
- Consider using AWS CLI profiles

### 13.3.2 Step 2: Create S3 Bucket

1. Navigate to S3 in AWS Console
2. Click "Create bucket"
3. Enter unique bucket name (e.g., `my-dvc-storage-bucket`)
4. Select region (e.g., `us-east-1`)
5. Keep default settings (or configure as needed)
6. Click "Create bucket"

#### Important Note

##### Bucket Naming Rules:

- Must be globally unique across all AWS
- 3-63 characters long
- Lowercase letters, numbers, hyphens only
- Cannot start/end with hyphen

### 13.3.3 Step 3: Install AWS CLI and DVC S3 Support

```
# Install AWS CLI tools
pip install awscli

# Install DVC with S3 support
pip install dvc[s3]
```

### 13.3.4 Step 4: Configure AWS Credentials

```
# Configure AWS CLI
aws configure
```

You'll be prompted for:

```
AWS Access Key ID [None]: <your-access-key-id>
AWS Secret Access Key [None]: <your-secret-access-key>
Default region name [None]: us-east-1
Default output format [None]: json
```

This creates:

- `~/.aws/credentials`: Stores access keys
- `~/.aws/config`: Stores configuration

### 13.3.5 Step 5: Add S3 Remote to DVC

```
# Add S3 bucket as DVC remote
dvc remote add -d dvcstore s3://my-dvc-storage-bucket

# Verify remote configuration
dvc remote list
```

**Expected Output:**

```
dvcstore    s3://my-dvc-storage-bucket
```

**Command Breakdown:**

- **dvc remote add:** Adds a remote storage location
- **-d:** Sets as default remote
- **dvcstore:** Name for this remote (can be anything)
- **s3://...:** S3 bucket URL

## 13.4 Pushing Data to S3

### 13.4.1 Initial Push

```
# Push all tracked data to S3
dvc push
```

**What happens:**

1. DVC reads `dvc.lock` to identify tracked files
2. Compares local cache with S3
3. Uploads missing/changed files
4. Stores files using content-addressable hashes

**Expected Output:**

Collecting	8.00 [00:00, 2.50entry/s]
Pushing	8.00 [00:05, 1.60file/s]
8 files pushed	

## 13.5 Understanding S3 Storage Structure

In your S3 bucket, files are stored like:

```
my-dvc-storage-bucket/
  files/
    md5/
      a1/
        b2c3d4e5f6g7h8i9j0k1l2m3n4o5p6
      x1/
        y2z3a4b5c6d7e8f9g0h1i2j3k4l5m6
      ...
      ...
```

- **Content-addressable:** Files named by MD5 hash
- **Deduplicated:** Same content = same hash = stored once
- **Efficient:** Only changed files uploaded

## 13.6 Pulling Data from S3

### 13.6.1 Cloning Project on New Machine

```
# Clone Git repository
git clone https://github.com/username/MLOPS-DVC-Project.git
cd MLOPS-DVC-Project

# Pull data from S3
dvc pull
```

**What happens:**

1. DVC reads `dvc.lock`
2. Downloads required files from S3
3. Populates local cache
4. Restores files to workspace

## 13.7 Complete Experiment Workflow with S3

### Full Workflow: Local → Git + DVC → S3

```
# 1. Run experiment
dvc exp run

# 2. If satisfied, push data to S3
dvc push

# 3. Commit code to Git
git add .
git commit -m "Experiment: increased n_estimators to 50"

# 4. Push code to GitHub
git push origin main
```

## 13.8 Key Concepts: Git vs DVC vs S3

System	Stores	Purpose
<b>Git</b>	Code, <code>dvc.yaml</code> , <code>dvc.lock</code> , <code>params.yaml</code>	Version control for code and metadata
<b>Local DVC Cache</b>	All data versions locally	Fast access, temporary
<b>S3 Remote</b>	All data versions in cloud	Backup, sharing, collaboration

### Data Flow Visualization

**Workspace → DVC Cache → S3 Remote**

- `dvc add / dvc exp run`: Workspace → DVC Cache
- `dvc commit`: Finalizes tracked outputs (no new cache unless manually modified)
- `dvc push`: DVC Cache → Remote Storage (S3/GCS/etc.)
- `dvc pull`: Remote Storage → DVC Cache
- `dvc checkout`: DVC Cache → Workspace

### 13.9 Configuring Remote in `.dvc/config`

After running `dvc remote add`, check `.dvc/config`:

#### `.dvc/config`

```
[core]
remote = dvcstore

['remote "dvcstore"]'
url = s3://my-dvc-storage-bucket
```

This file is tracked by Git, so team members automatically use the same remote!

### 13.10 Verifying S3 Upload

1. Go to AWS S3 Console
2. Navigate to your bucket
3. Check `files/md5/` directory
4. Verify files exist (they'll be named by hash)

### 13.11 Alternative: Using DVC with Other Storage

DVC supports multiple storage backends:

- **Google Cloud Storage**: `dvc remote add -d myremote gs://bucket`
- **Azure Blob**: `dvc remote add -d myremote azure://container`
- **Google Drive**: `dvc remote add -d myremote gdrive://folder-id`
- **SSH/SFTP**: `dvc remote add -d myremote ssh://server/path`
- **Local/Network**: `dvc remote add -d myremote /mnt/storage`

## 14 What does dvc push actually push?

### 14.1 What Exactly Does dvc push Upload?

#### Key Rule

dvc push uploads **only the data objects referenced by the current workspace state**. It does **not** automatically upload all past experiments or all cached data.

#### 14.1.1 How DVC Decides What to Push

When you run dvc push, DVC answers one question:

*Which data objects are required to reproduce the current workspace?*

To determine this, DVC inspects:

- `dvc.lock` (if using `dvc.yaml` pipelines)
- `.dvc` files (if using `dvc add`)

Each of these files contains **hashes** that point to objects in `.dvc/cache/`. Only those cache objects are considered “needed”.

#### 14.1.2 What Gets Pushed

- Data files, models, or artifacts stored in `.dvc/cache/`
- Only the cache objects whose hashes appear in:
  - the current `dvc.lock`, and/or
  - the current `.dvc` files

**Important:** DVC does *not* push hash files or tokens themselves— it pushes the actual data objects referenced by those hashes.

#### 14.1.3 What Does NOT Get Pushed

- Data from previous experiments that are not applied
- Data referenced only by discarded or local experiments
- Experiment history stored in `.git/refs/exp/`
- All cache contents by default

#### 14.1.4 Example: Multiple Experiments

Assume you ran three experiments:

- Experiment 1 produced data version A
- Experiment 2 produced data version B
- Experiment 3 produced data version C

After applying Experiment 3, the current `dvc.lock` contains hashes pointing only to version C.

```
dvc push
```

**Result:**

- Only data version C is uploaded to remote storage
- Versions A and B remain local

#### 14.1.5 Pushing Data for All Commits

If you want to upload data referenced by **all Git commits** (not just the current one), use:

```
dvc push --all-commits
```

This scans Git history and uploads cache objects referenced by every committed state.

**Note:** This still does not push local experiment history—only committed states.

## 14.2 Summary

**Warning****Key Takeaways:**

- `dvc push` uploads data referenced by the current workspace only
- It relies on `dvc.lock` and/or `.dvc` files
- Unapplied experiments are not pushed
- To push all committed data, use `dvc push --all-commits`

## 15 Complete MLOps Workflow Summary

### 15.1 Running Multiple Experiments with DVCLive

#### How DVC Experiments Work

When you run multiple experiments with `dvc exp run`:

##### What Happens:

1. Run Experiment 1 → DVC stores it locally in `.git/refs/exps/`
2. Run Experiment 2 → DVC stores it separately, Exp 1 remains
3. Run Experiment 3 → DVC stores it separately, Exp 1 & 2 remain
4. Use `dvc exp show` → See all experiments with metrics
5. Use `dvc exp diff` → Compare experiments

**Result:** All experiments are tracked locally and can be compared!

##### DVC + DVCLive Benefits:

- DVCLive automatically logs metrics, parameters, and plots
- DVC assigns each experiment a unique identifier
- Compare multiple experiments side-by-side
- Local experiment history preserved

#### Warning

**Important Limitation:** Experiment History is Local Only

##### The Constraint:

- DVC experiments live in `.git/refs/exps/` (local Git refs)
- `dvc push` does NOT upload experiment history to remote
- Only the **applied/promoted** experiment is committed to Git
- Team members cannot see your experiment history

**For centralized experiment tracking across teams,** use MLflow or similar tools.  
(We'll cover this in advanced sections.)

### 15.2 Correct Workflow: Compare and Apply Best Experiment

#### Proper Workflow for Multiple Experiments

```
# ===== RUN MULTIPLE EXPERIMENTS =====
# Experiment 1
vim params.yaml # e.g., n_estimators: 20
dvc exp run -n "baseline-20"
```

```

# Experiment 2
vim params.yaml # e.g., n_estimators: 50
dvc exp run -n "test-50"

# Experiment 3
vim params.yaml # e.g., n_estimators: 100
dvc exp run -n "test-100"

# ===== COMPARE EXPERIMENTS =====
dvc exp show
# See all experiments with metrics side-by-side

dvc exp diff baseline-20 test-50
# Compare specific experiments

# ===== APPLY BEST EXPERIMENT =====
# Choose best performing experiment (e.g., test-50)
dvc exp apply test-50

# ===== COMMIT THE BEST VERSION =====
# 1. Push data to S3 (uploads cached outputs)
dvc push

# 2. Commit to Git (saves code + metadata)
git add dvc.lock params.yaml
git commit -m "Apply best experiment: n_estimators=50, accuracy=0.92"
git push origin main

```

### What Gets Committed

#### One Git Commit = One Promoted Experiment:

- Only the **applied** experiment becomes permanent
- Other experiments remain local (not pushed to remote)
- Git history shows only promoted experiments
- Each team member's local experiments stay private

### 15.3 Why This Matters for Rollback

What You CAN Rollback	What You CANNOT Rollback
Promoted experiments (via Git commits)	Unpromoted local experiments
Code, data, models of committed state	Alternative experiments (not preserved)
git checkout + dvc pull works	Unpromoted experiments (may be discarded)
Parameters and metrics in Git history	Local experiment refs (not preserved, may be discarded)

### Warning

#### Key Understanding: Git Commits Are Singular

When you `dvc exp apply` and commit:

- One Git commit = one experiment outcome
- Unpromoted experiments are not preserved and may be discarded
- Rolling back restores only committed experiments
- To preserve all experiments, you need external tracking (e.g., MLflow)

## 15.4 Rollback Procedure

To reproduce a previously committed experiment:

```
# 1. View commit history
git log --oneline

# Example output:
# a1b2c3d Apply best: n_estimators=100, accuracy=0.94
# d4e5f6g Apply best: n_estimators=50, accuracy=0.92
# h7i8j9k Baseline: n_estimators=20, accuracy=0.88

# 2. Checkout desired experiment
git checkout d4e5f6g

# 3. Pull corresponding data from S3
dvc pull

# 4. Verify
python src/Model_Evaluation.py
```

**What happens:**

1. `git checkout`: Restores code + `dvc.lock` for that commit
2. `dvc pull`: Uses `dvc.lock` to fetch exact data from S3
3. Everything restored: code, data, model, parameters

### Rollback is Only Possible For...

- Experiments that were **applied** with `dvc exp apply`
- Experiments that were committed to Git
- Experiments whose data was pushed with `dvc push`

Unpromoted local experiments are not preserved and may be discarded!

## 15.5 Complete Command Reference

### Experiment Workflow

```
# 1. Run multiple experiments
dvc exp run -n "exp-name" # Repeat with different params

# 2. Compare experiments
dvc exp show             # View all experiments
dvc exp diff exp1 exp2   # Compare two experiments

# 3. Apply best experiment
dvc exp apply exp-name    # Promote to workspace

# 4. Commit the promoted experiment
dvc push                  # Upload cached data to S3
git add dvc.lock params.yaml
git commit -m "message"
git push                   # Upload to GitHub
```

### To Rollback

1. git checkout <hash> # Restore code
2. dvc pull # Restore data

### Advanced: Push All Commits' Data

```
# Upload data for all Git commits (not just current)
dvc push --all-commits
```

## 15.6 Understanding DVCLive and Local Experiments

### The Experiment Lab Analogy

Think of DVC experiments like a lab notebook:

#### Current Workspace (dvclive/):

- Shows the active experiment's metrics and plots
- Gets overwritten when you run a new experiment
- Temporary workspace (don't add to Git)

#### Experiment History (Local Git Refs):

- DVC stores all experiments in .git/refs/exps/
- Each experiment keeps its own metrics and parameters
- `dvc exp show` displays all experiments
- **Local only** - not pushed to GitHub/S3

#### Committed Experiments (Git History):

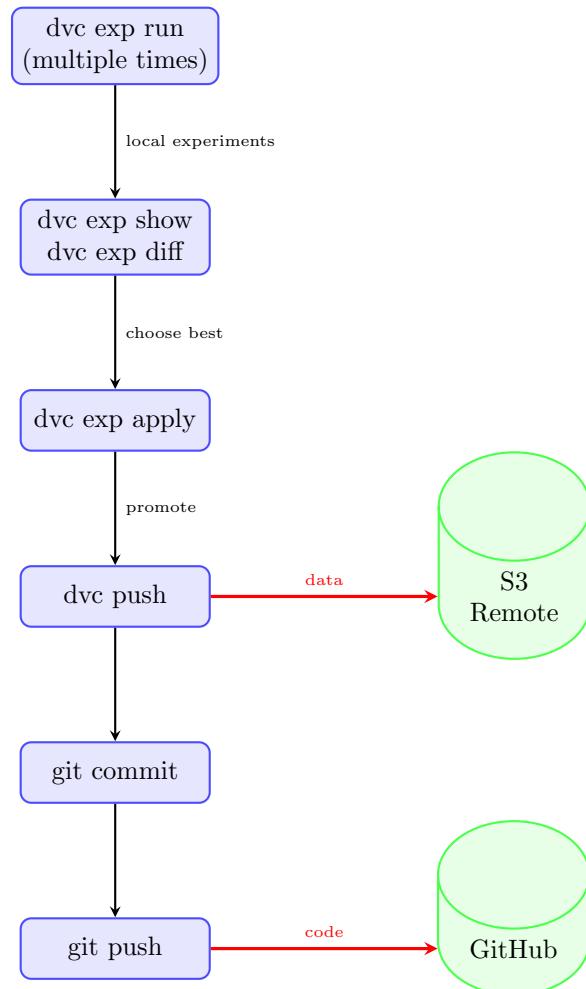
- Only applied experiments enter Git history
- Permanent and shareable across team
- Can be rolled back anytime

### Warning

#### Key Points:

- `dvc live/` folder shows current state only
- DVC remembers all experiments locally
- Only promoted experiments survive in remote/Git
- For centralized experiment tracking, use MLflow

## 15.7 Workflow Visualization



## 15.8 When to Use MLflow

### Limitations of DVC Experiments

DVC experiments are excellent for local iteration, but have limitations:

#### What DVC Provides:

- Local experiment tracking and comparison
- Data and model versioning
- Reproducibility of promoted experiments

#### What DVC Does NOT Provide:

- Centralized experiment history across team
- Remote storage of all experiment runs
- Web UI for comparing experiments
- Experiment history persistence beyond local refs

**Solution:** For centralized experiment tracking, use MLflow alongside DVC. (This will be covered in advanced sections.)

## 16 Common Issues and Troubleshooting

---

### 16.1 DVC Issues

#### 16.1.1 Issue: "Stage didn't change, skipping"

**Problem:** DVC won't re-run stages even after changes.

**Solutions:**

```
# Force re-run entire pipeline  
dvc repro --force  
  
# Force re-run specific stage  
dvc repro --force model_building  
  
# Clear cache and re-run  
dvc remove <stage>  
dvc repro
```

#### 16.1.2 Issue: "Failed to push data"

**Problem:** dvc push fails with S3 errors.

**Solutions:**

1. Check AWS credentials:

```
aws s3 ls s3://your-bucket-name
```

2. Verify remote configuration:

```
dvc remote list  
cat .dvc/config
```

3. Check bucket permissions in AWS Console

4. Reconfigure AWS:

```
aws configure
```

#### 16.1.3 Issue: "File not found" during dvc pull

**Problem:** dvc pull can't find files in S3.

**Cause:** You never ran dvc push for that version.

**Solution:**

- Cannot recover if data never pushed
- Re-run the experiment
- Follow complete workflow (commit → push)

## 16.2 Git Issues

### 16.2.1 Issue: Large files in Git

**Problem:** Accidentally committed data/models to Git.

**Solution:**

```
# Remove from Git tracking
git rm -r --cached data/
git rm -r --cached models/

# Update .gitignore
echo "data/" >> .gitignore
echo "models/" >> .gitignore

# Commit removal
git commit -m "Remove large files from Git tracking"
git push origin main
```

### 16.2.2 Issue: Merge conflicts in dvc.lock

**Problem:** Multiple team members modified pipeline.

**Solution:**

```
# Accept their version
git checkout --theirs dvc.lock

# Or accept your version
git checkout --ours dvc.lock

# Then re-run pipeline
dvc repro

# Commit resolved lock
git add dvc.lock
git commit -m "Resolve dvc.lock conflict"
```

## 16.3 Python Issues

### 16.3.1 Issue: Import errors

**Problem:** Cannot import modules.

**Solution:**

```
# Install requirements
pip install -r requirements.txt

# Or install individually
pip install pandas scikit-learn nltk pyyaml dvclive
```

### 16.3.2 Issue: NLTK data not found

**Problem:** stopwords or punkt not found.

**Solution:**

```
1 import nltk
2 nltk.download('stopwords')
3 nltk.download('punkt')
```

## 16.4 AWS Issues

### 16.4.1 Issue: "Access Denied" on S3

**Problem:** IAM user lacks permissions.

**Solution:**

1. Go to IAM in AWS Console
2. Find your user
3. Attach policy: `AmazonS3FullAccess`
4. Or create custom policy with `s3:GetObject`, `s3:PutObject`, `s3>ListBucket`

### 16.4.2 Issue: Wrong region

**Problem:** Bucket in different region than configured.

**Solution:**

```
# Check bucket region in S3 console
# Update AWS config
aws configure set region <correct-region>
```

## 16.5 Experiment Tracking Issues

### 16.5.1 Issue: Experiments not showing

**Problem:** `dvc exp show` displays no experiments.

**Cause:** Used `dvc repro` instead of `dvc exp run`.

**Solution:**

- Use `dvc exp run` for experiments
- `dvc repro` is for production pipeline execution

### 16.5.2 Issue: DVCLive not logging

**Problem:** No metrics in `dvclive/`.

**Solution:**

1. Verify DVCLive installed: `pip install dvclive`
2. Check code has `with Live(save_dvc_exp=True):`
3. Ensure metrics logged: `live.log_metric(...)`
4. Check for exceptions in logs

## 17 MLOps Best Practices

### 17.1 Code Organization Best Practices

#### 17.1.1 Modular Design Principles

##### 1. Single Responsibility Principle:

- One component per file
- Each module does one thing well
- Clear separation of concerns

##### Good vs Bad Structure

###### No Bad - Everything in one file:

```
pipeline.py (2000 lines)
    - Data loading
    - Preprocessing
    - Feature engineering
    - Model training
    - Evaluation
```

###### Yes Good - Modular structure:

```
src/
--- Data_Ingestion.py      (200 lines)
--- Data_Pre_Processing.py (250 lines)
--- Feature_Engineering.py (180 lines)
--- Model_Building.py     (150 lines)
--- Model_Evaluation.py    (200 lines)
```

##### 2. Reusable Functions:

```
1 # Good: Reusable utility functions
2 def load_params(params_path: str) -> dict:
3     """Load parameters from YAML file."""
4     with open(params_path, 'r') as f:
5         return yaml.safe_load(f)
6
7 def setup_logger(name: str, log_file: str) -> logging.Logger:
8     """Configure and return a logger."""
9     logger = logging.getLogger(name)
10    # ... setup code
11    return logger
12
```

##### 3. Type Hints and Docstrings:

```
1 def train_model(
2     X_train: np.ndarray,
3     y_train: np.ndarray,
4     params: dict
5 ) -> RandomForestClassifier:
6     """
```

```

7     Train a RandomForest classifier.
8
9     Args:
10    X_train: Training features of shape (n_samples, n_features)
11    y_train: Training labels of shape (n_samples,)
12    params: Dictionary containing model hyperparameters
13
14    Returns:
15    Trained RandomForestClassifier instance
16
17    Raises:
18    ValueError: If X_train and y_train shapes don't match
19    """
20    # Implementation
21    pass
22

```

### 17.1.2 Logging Best Practices

#### 1. Appropriate Log Levels:

```

1 logger.debug('Detailed information for debugging')
2 logger.info('General informational messages')
3 logger.warning('Warning messages for potential issues')
4 logger.error('Error messages for failures')
5 logger.critical('Critical errors requiring immediate attention')
6

```

#### 2. Structured Logging:

```

1 # Good: Structured, informative logging
2 logger.info(f'Training started with {X_train.shape[0]} samples')
3 logger.info(f'Parameters: {params}')
4 logger.debug(f'Feature shape: {X_train.shape}')
5 logger.info(f'Training completed in {elapsed_time:.2f}s')
6
7 # Bad: Vague logging
8 logger.info('Training done')
9

```

#### 3. Log File Organization:

```

logs/
+-- Data_Ingestion_2024-12-20.log
+-- Data_Pre_Processing_2024-12-20.log
+-- Feature_Engineering_2024-12-20.log
+-- Model_Building_2024-12-20.log
+-- Model_Evaluation_2024-12-20.log

```

### 17.1.3 Error Handling Patterns

```

1 def load_data(file_path: str) -> pd.DataFrame:
2     """Load data with comprehensive error handling."""

```

```

3     try:
4         # Attempt operation
5         df = pd.read_csv(file_path)
6         logger.info(f'Successfully loaded data from {file_path}')
7         return df
8
9     except FileNotFoundError:
10        # Specific exception handling
11        logger.error(f'File not found: {file_path}')
12        raise
13
14    except pd.errors.ParserError as e:
15        # Another specific exception
16        logger.error(f'Failed to parse CSV: {e}')
17        raise
18
19    except Exception as e:
20        # Catch-all for unexpected errors
21        logger.error(f'Unexpected error loading data: {e}')
22        raise
23
24    finally:
25        # Cleanup code (if needed)
26        logger.debug('Load data operation completed')

```

## 17.2 Data Management Best Practices

### 17.2.1 Directory Structure Philosophy

#### Three-Tier Data Structure

**data/raw/**: Original, immutable data

- Never modify these files
- Treat as read-only
- Can always reproduce from source

**data/interim/**: Intermediate processing

- Partially processed data
- Reusable across experiments
- Checkpoint for long pipelines

**data/processed/**: Final features

- Ready for model training
- Feature matrices
- Final transformations applied

### 17.2.2 Data Versioning Strategy

#### 1. Always Use DVC for Data:

```
# Either track data directly:  
dvc add data/processed  
  
# OR track data via pipelines:  
dvc.yaml:  
  outs:  
    - data/processed
```

## 2. Never Commit Large Files to Git:

### Warning

Files larger than 100MB should NEVER go into Git:

- Git repositories become slow
- Clone times increase dramatically
- GitHub blocks files >100MB
- Use DVC instead!

## 3. Avoid Manual Data Version Directories:

- Prefer DVC hashes over folder-based versioning
- Use Git commits and DVC cache for true version history
- Manual version folders are optional and discouraged at scale

### 17.2.3 Data Validation

```
1 def validate_data(  
2     df: pd.DataFrame,  
3     expected_columns: list,  
4     min_rows: int = 100  
5 ) -> bool:  
6     """  
7         Validate data meets expectations.  
8  
9     Checks:  
10        - Required columns exist  
11        - Minimum row count met  
12        - No completely empty columns  
13        - Correct data types  
14    """  
15    logger.info('Starting data validation')  
16  
17    # Check columns  
18    missing_cols = set(expected_columns) - set(df.columns)  
19    if missing_cols:  
20        logger.error(f'Missing columns: {missing_cols}')  
21        raise ValueError(f"Missing required columns: {missing_cols}")  
22  
23    # Check row count  
24    if len(df) < min_rows:
```

```

25         logger.error(f'Only {len(df)} rows, expected {min_rows}')
26         raise ValueError(f"Insufficient rows: {len(df)} < {min_rows}
27     ")
28
29     # Check for empty columns
30     empty_cols = df.columns[df.isnull().all()].tolist()
31     if empty_cols:
32         logger.warning(f'Empty columns found: {empty_cols}')
33
34     # Check data types
35     type_issues = []
36     expected_types = {'text': 'object', 'target': ('int64', 'float64'
37     )}
38     for col, expected_dtype in expected_types.items():
39         if df[col].dtype not in expected_dtype:
40             type_issues.append(f'{col}: {df[col].dtype}')
41
42     if type_issues:
43         logger.error(f'Type mismatches: {type_issues}')
44         raise TypeError(f"Data type issues: {type_issues}")
45
46     logger.info('Data validation passed')
47     return True

```

## 17.3 Experiment Management Best Practices

### 17.3.1 Naming Conventions

#### 1. Descriptive Commit Messages:

##### Good Commit Messages

Yes "Promote: n\_estimators=100, accuracy=0.97"  
 Yes "Feature: Add TF-IDF with bigrams, precision +0.03"  
 Yes "Fix: Correct data leakage in preprocessing pipeline"  
 Yes "Baseline: RandomForest with default params, acc=0.95"

##### Warning

##### Bad Commit Messages:

No "Update model"  
 No "Fix bug"  
 No "Changes"  
 No "WIP"  
 No "asdfasdfs"

#### 2. Experiment Naming Pattern:

`exp-<date>-<feature>-<value>`

Examples:

`exp-2024-12-20-n_estimators-100`

```
exp-2024-12-20-max_features-200
exp-2024-12-20-baseline
```

### 17.3.2 Parameter Management

#### 1. All Hyperparameters in params.yaml:

```
1 # No Bad: Hardcoded in code
2 clf = RandomForestClassifier(n_estimators=50, max_depth=10)
3
4 # Yes Good: From params.yaml
5 params = load_params('params.yaml')['model_building']
6 clf = RandomForestClassifier(**params)
7
```

#### 2. Document Parameter Choices:

```
1 # params.yaml with comments
2 model_building:
3   n_estimators: 50          # Increased from 20, improved accuracy
4   max_depth: 10            # Prevent overfitting on small dataset
5   min_samples_split: 5     # Baseline value
6   random_state: 2          # For reproducibility
7
```

#### 3. Track Parameter Changes:

```
# In commit messages
git commit -m "Exp: n_estimators 20→50, max_depth 5→10"
```

### 17.3.3 Metric Tracking

#### 1. Track Multiple Metrics:

```
1 with Live(save_dvc_exp=True) as live:
2   # Primary metrics
3   live.log_metric('accuracy', accuracy)
4   live.log_metric('precision', precision)
5   live.log_metric('recall', recall)
6   live.log_metric('f1_score', f1)
7   live.log_metric('auc', auc)
8
9   # Secondary metrics
10  live.log_metric('training_time', elapsed_time)
11  live.log_metric('num_features', X_train.shape[1])
12  live.log_metric('num_samples', X_train.shape[0])
13
14  # Log all parameters
15  live.log_params(params)
16
```

#### 2. Business-Relevant Metrics:

```
1 # Don't just track ML metrics
2 # Track business impact
3 business_metrics = {
4     'false_positive_rate': fp_rate,
5     'false_negative_rate': fn_rate,
6     'cost_per_error': calculate_cost(fp, fn),
7     'expected_savings': calculate_savings(tp, tn)
8 }
9
```

## 17.4 Collaboration Best Practices

### 17.4.1 Documentation Standards

#### Essential README.md Sections

```
# Project Name

## Overview
Brief description of the ML problem and solution.

## Setup Instructions
1. Clone repository
2. Create virtual environment
3. Install dependencies
4. Configure AWS credentials
5. Run pipeline

## Project Structure
Explain directory organization.

## Data Description
- Source of data
- Features and target
- Data collection process

## Pipeline Stages
Describe each stage of the pipeline.

## Experiments
Table of experiments with results.

## How to Run
- Training: 'dvc repro'
- Experiments: 'dvc exp run'
- Rollback: Instructions

## Troubleshooting
Common issues and solutions.

## Team Members
```

Contact information.

```
## License
Project license.
```

#### 17.4.2 Code Review Process

##### 1. Pull Request Template:

```
## Changes
- What was changed
- Why it was changed

## Experiment Results
- Metric improvements
- Parameter changes

## Testing
- What was tested
- Test results

## Checklist
- [ ] Code follows style guide
- [ ] Tests pass
- [ ] Documentation updated
- [ ] dvc repro runs successfully (only for projects with dvc.yaml)
```

##### 2. Review Checklist:

- Code is modular and readable
- Proper error handling
- Logging is comprehensive
- Parameters in params.yaml
- No hardcoded values
- Data not committed to Git
- Tests exist and pass
- Documentation updated

#### 17.4.3 Team Communication

##### 1. Experiment Log:

Date	Experimenter	Changes	Result
2024-12-20	Alice	n_estimators=100	acc=0.97
2024-12-19	Bob	Add bigrams	acc=0.96
2024-12-18	Alice	Baseline	acc=0.95

##### 2. Weekly Sync Meetings:

- Review experiment results
- Discuss blockers
- Plan next experiments
- Share learnings

### 3. Slack/Communication Channels:

```
#ml-experiments: Share experiment results
#ml-questions: Ask technical questions
#ml-alerts: Pipeline failures, important updates
```

## 17.5 Security and Compliance

### 17.5.1 Credential Management

#### Warning

##### NEVER Commit These to Git:

- AWS Access Keys / Secret Keys
- API tokens
- Database passwords
- Private keys
- OAuth tokens
- Any secrets or credentials

#### Secure Credential Management

```
1 # Yes Good: Use environment variables
2 import os
3
4 AWS_ACCESS_KEY = os.getenv('AWS_ACCESS_KEY_ID')
5 AWS_SECRET_KEY = os.getenv('AWS_SECRET_ACCESS_KEY')
6 API_KEY = os.getenv('API_KEY')
7
8 # Load from .env file (never commit .env!)
9 from dotenv import load_dotenv
10 load_dotenv()
11
12 # No Bad: Hardcoded credentials
13 AWS_ACCESS_KEY = "AKIAIOSFODNN7EXAMPLE" # DON'T DO THIS!
```

##### .env file (add to .gitignore):

```
AWS_ACCESS_KEY_ID=your_access_key
AWS_SECRET_ACCESS_KEY=your_secret_key
API_KEY=your_api_key
```

### 17.5.2 Data Privacy

#### 1. Sensitive Data Handling:

```

1 def anonymize_data(df: pd.DataFrame) -> pd.DataFrame:
2     """Remove PII (Personally Identifiable Information)."""
3     # Remove or hash sensitive columns
4     df = df.drop(columns=['email', 'phone', 'ssn'])
5
6     # Hash user IDs
7     df['user_id'] = df['user_id'].apply(
8         lambda x: hashlib.sha256(str(x).encode()).hexdigest()
9     )
10
11     return df
12

```

#### 2. Data Access Controls:

- Use IAM roles with minimal permissions
- Separate dev/staging/prod environments
- Audit data access logs
- Implement data retention policies

## 17.6 Performance Optimization

### 17.6.1 Caching Strategies

#### 1. DVC Cache Optimization:

```

1 # Use hardlinks for faster cache
2 dvc config cache.type hardlink,symlink
3
4 # Set cache directory on faster disk
5 dvc cache dir /path/to/ssd/cache
6
7 # Shared cache for team
8 dvc config cache.dir /shared/dvc/cache
9 dvc config cache.shared group
10

```

#### 2. Partial Pipeline Execution:

```

1 # Run only specific stage
2 dvc repro model_building
3
4 # Run from specific stage onwards
5 dvc repro --downstream model_building
6

```

#### 3. Parallel Processing:

```

1 # Use all CPU cores
2 from sklearn.ensemble import RandomForestClassifier
3
4 clf = RandomForestClassifier(
5     n_estimators=100,

```

```

6      n_jobs=-1 # Use all cores
7 )
8
9 # Parallel data loading
10 import multiprocessing as mp
11
12 with mp.Pool(processes=4) as pool:
13     results = pool.map(process_file, file_list)
14

```

## 17.6.2 Memory Management

```

1 def load_data_chunked(file_path: str, chunksize: int = 10000):
2     """Load large CSV in chunks to manage memory."""
3     chunks = []
4     for chunk in pd.read_csv(file_path, chunksize=chunksize):
5         # Process chunk
6         processed = preprocess_chunk(chunk)
7         chunks.append(processed)
8
9     return pd.concat(chunks, ignore_index=True)
10
11 # Clear memory after use
12 import gc
13 del large_dataframe
14 gc.collect()

```

## 17.7 Common Pitfalls to Avoid

### Warning

#### Top 10 Common Mistakes in MLOps:

1. Committing large data or models to Git instead of using DVC
2. Promoting or committing experiments prematurely instead of comparing them first
3. Hardcoding parameters instead of using `params.yaml`
4. Not pushing data to remote storage, making commits unreproducible
5. Editing `dvc.lock` manually
6. Committing AWS credentials, API keys, or secrets to version control
7. Using `dvc repro` when experiment tracking with `dvc exp run` is intended
8. Ignoring logs and metrics, which are critical for debugging and evaluation
9. Skipping data validation at pipeline stages
10. Failing to document experiment assumptions and results

### Best Practice Checklist

Before committing code, verify:

- All parameters are defined in `params.yaml`
- No hardcoded configuration values in code
- Logging is configured in all critical modules
- Error handling is implemented
- Type hints and docstrings are added where appropriate
- Data files are excluded via `.gitignore`
- No credentials or secrets are present in the codebase
- Pipelines run successfully (`dvc repro`, if applicable)
- Best experiment is applied and committed (not all experiments)
- Documentation and experiment notes are updated

## 18 Quick Reference Guide

### 18.1 Essential DVC Commands

#### Pipeline Management

```
dvc init          # Initialize DVC in repository
dvc repro         # Run entire pipeline
dvc repro --force # Force re-run all stages
dvc repro <stage> # Run specific stage
dvc dag           # Visualize pipeline as DAG
dvc status        # Check pipeline status
dvc status -c     # Check status with cloud comparison
```

#### Data Management

```
dvc add <file>      # Track file/directory with DVC
dvc commit          # Save workspace changes to cache
dvc push            # Upload data to remote storage
dvc pull            # Download data from remote storage
dvc checkout        # Restore files from cache
dvc fetch           # Download without checkout
dvc gc              # Garbage collect unused cache
```

#### Remote Storage

```
dvc remote add -d <name> <url>      # Add default remote
dvc remote list          # List all remotes
dvc remote modify <name> <key> <v>    # Modify remote config
dvc remote remove <name>          # Remove remote
dvc remote rename <old> <new>       # Rename remote

# Examples:
dvc remote add -d s3store s3://bucket
dvc remote add -d gdrive gdrive://folder-id
dvc remote add -d local /mnt/storage
```

#### Experiment Tracking

```
dvc exp run          # Run experiment (saves automatically)
dvc exp show         # Show all experiments in table
dvc exp show --no-pager # Show without pagination
dvc exp diff         # Compare experiments
dvc exp diff <exp1> <exp2> # Compare specific experiments
dvc exp apply <id>      # Apply experiment to workspace
dvc exp remove <id>      # Remove specific experiment
dvc exp gc           # Clean up all experiments
dvc exp list         # List all experiments
```

### Git Integration

```
git add dvc.yaml dvc.lock params.yaml # Stage DVC files
git commit -m "message"           # Commit to Git
git push origin main             # Push to GitHub

# Typical workflow
git add .
git commit -m "Experiment: description"
git push
```

## 18.2 Complete Workflow Cheat Sheet

### 18.2.1 Initial Project Setup

#### Step 1: Project Initialization

```
# 1. Create and clone Git repository
git clone https://github.com/username/project.git
cd project

# 2. Create virtual environment
python -m venv venv
source venv/bin/activate # Linux/Mac
venv\Scripts\activate # Windows

# 3. Install dependencies
pip install pandas scikit-learn nltk pyyaml
pip install dvc dvclive
pip install dvc[s3]      # For AWS S3
pip install awscli       # AWS CLI

# 4. Initialize DVC
dvc init

# 5. Create directory structure
mkdir -p src data/raw data/interim data/processed models reports logs

# 6. Configure .gitignore
cat >> .gitignore << EOF
data/
models/
reports/
logs/
venv/
__pycache__/
*.pyc
dvclive/
EOF

# 7. Initial commit
```

```
git add .
git commit -m "Initial project setup"
git push origin main
```

### 18.2.2 Full Experiment Workflow

#### Step 2: Running Experiments

```
# === RUN MULTIPLE EXPERIMENTS ===

# 1. Modify parameters
vim params.yaml

# 2. Run experiment
dvc exp run

# 3. View and compare results
dvc exp show

# === REPEAT STEPS 1{3 FOR MULTIPLE EXPERIMENTS ===

# 4. Choose the best experiment

# 5. Apply (promote) the best experiment
dvc exp apply <exp-id>

# 6. Commit promoted state to Git
git add dvc.lock params.yaml
git commit -m "Promote best experiment: n_estimators=50, acc=0.97"

# 7. Push data to remote storage
dvc push

# 8. Push code to GitHub
git push origin main
```

### 18.2.3 Rollback to a Promoted (Committed) Experiment

#### Step 3: Rollback Procedure

```
# 1. View commit history
git log --oneline

# Example output:
# a1b2c3d Promote best: n_estimators=100, acc=0.98
# d4e5f6g Promote best: n_estimators=50, acc=0.97
# h7i8j9k Promote best: n_estimators=20, acc=0.95

# 2. Checkout desired promoted experiment (Git commit)
git checkout d4e5f6g
```

```
# 3. Pull corresponding data from remote storage
dvc pull

# 4. Verify the restored experiment
python src/Model_Evaluation.py
cat reports/metrics.json

# 5. Return to the latest version (optional)
git checkout main
dvc pull
```

#### 18.2.4 AWS S3 Setup Workflow

##### Step 4: AWS Integration

```
# 1. Install AWS tools
pip install dvc[s3]
pip install awscli

# 2. Configure AWS credentials
aws configure
# Enter: Access Key ID
# Enter: Secret Access Key
# Enter: Region (e.g., us-east-1)
# Enter: Output format (json)

# 3. Add S3 remote to DVC
dvc remote add -d dvcstore s3://your-bucket-name

# 4. Verify configuration
dvc remote list
# Output: dvcstore    s3://your-bucket-name

# 5. Test connection
dvc push

# 6. Commit DVC config to Git
git add .dvc/config
git commit -m "Add S3 remote storage"
git push origin main
```

### 18.3 Project Structure Template

```
MLOPS-DVC-Project/
|
+-- Experiments/          # Exploration phase
|   +- spam.csv
|   +- mynotebook.ipynb
|
+-- src/                  # Production code
|   +- Data_Ingestion.py
|   +- Data_Pre_Processing.py
|   +- Feature_Engineering.py
|   +- Model_Building.py
|   +- Model_Evaluation.py
|
+-- data/                 # Data files (DVC tracked)
|   +- raw/                # Original data
|   +- interim/             # Intermediate processing
|   +- processed/           # Final features
|
+-- models/                # Trained models (DVC tracked)
|   +- model.pkl
|
+-- reports/               # Metrics and reports (DVC tracked)
|   +- metrics.json
|
+-- logs/                  # Application logs
|   +- Data_Ingestion.log
|   +- Data_Pre_Processing.log
|   +- Feature_Engineering.log
|   +- Model_Building.log
|   +- Model_Evaluation.log
|
+-- dvclive/                # DVCLive temporary files
|   +- metrics.json
|   +- params.yaml
|   +- plots/
|
+-- .dvc/                  # DVC configuration
|   +- cache/               # Local data cache
|   +- config/               # Remote storage config
|   +- .gitignore
|
+-- dvc.yaml                # Pipeline definition
+-- dvc.lock                 # Pipeline lock file
+-- params.yaml              # Hyperparameters
+-- .gitignore                # Git ignore rules
+-- .dvcignore                # DVC ignore rules
+-- requirements.txt          # Python dependencies
+-- README.md                 # Project documentation
```

## 18.4 Key File Purposes

File	Purpose
<b>dvc.yaml</b>	Defines pipeline stages, dependencies, outputs, parameters
<b>dvc.lock</b>	Locks exact file versions using MD5 hashes
<b>params.yaml</b>	Centralized hyperparameter configuration
<b>.gitignore</b>	Tells Git what NOT to track (data, models, logs)
<b>.dvcignore</b>	Tells DVC what NOT to track
<b>.dvc/config</b>	DVC remote storage configuration
<b>.dvc/cache/</b>	Local cache storing all data versions
<b>dvclive/</b>	Temporary experiment metrics (current run only)
<b>requirements.txt</b>	Python package dependencies

## 18.5 Git vs DVC Tracking

Git Tracks	DVC Tracks
Source code ( <code>src/*.py</code> )	Data files ( <code>data/**</code> )
<code>dvc.yaml</code> (pipeline definition)	Model files ( <code>models/*.pkl</code> )
<code>dvc.lock</code> (version locks)	Reports ( <code>reports/*.json</code> )
<code>params.yaml</code> (hyperparameters)	Large binary files
<code>.gitignore</code> , <code>.dvcignore</code>	Outputs defined in <code>dvc.yaml</code>
<code>.dvc/config</code> (DVC settings)	Intermediate data artifacts
Documentation ( <code>README.md</code> )	Feature matrices
<code>requirements.txt</code>	Preprocessed datasets

## 18.6 params.yaml Template

### Complete params.yaml Example

```

1 # Data Ingestion Parameters
2 data_ingestion:
3     test_size: 0.15
4     random_state: 2
5     data_url: "https://raw.githubusercontent.com/.../spam.csv"
6
7 # Data Preprocessing Parameters
8 data_preprocessing:
9     remove_stopwords: true
10    apply_stemming: true
11    lowercase: true
12
13 # Feature Engineering Parameters
14 feature_engineering:
15    max_features: 45
16    ngram_range: [1, 2]
17    min_df: 2
18    max_df: 0.95
19
20 # Model Building Parameters
21 model_building:
22    model_type: "RandomForest"
23    n_estimators: 20
24    max_depth: null

```

```
25     min_samples_split: 2
26     min_samples_leaf: 1
27     random_state: 2
28     n_jobs: -1
29
30 # Model Evaluation Parameters
31 model_evaluation:
32     metrics:
33         - accuracy
34         - precision
35         - recall
36         - f1_score
37         - auc
```

## 18.7 dvc.yaml Template

### Complete dvc.yaml with All Features

```
1 stages:
2     data_ingestion:
3         cmd: python src/Data_Ingestion.py
4         deps:
5             - src/Data_Ingestion.py
6         params:
7             - data_ingestion.test_size
8             - data_ingestion.random_state
9         outs:
10            - data/raw
11
12     data_preprocessing:
13         cmd: python src/Data_Pre_Processing.py
14         deps:
15             - data/raw
16             - src/Data_Pre_Processing.py
17         params:
18             - data_preprocessing
19         outs:
20            - data/interim
21
22     feature_engineering:
23         cmd: python src/Feature_Engineering.py
24         deps:
25             - data/interim
26             - src/Feature_Engineering.py
27         params:
28             - feature_engineering.max_features
29             - feature_engineering.ngram_range
30         outs:
31            - data/processed
32
33     model_building:
34         cmd: python src/Model_Building.py
35         deps:
36             - data/processed
```

```
37      - src/Model_Building.py
38  params:
39    - model_building
40  outs:
41    - models/model.pkl
42
43  model_evaluation:
44    cmd: python src/Model_Evaluation.py
45    deps:
46      - models/model.pkl
47      - data/processed
48      - src/Model_Evaluation.py
49    params:
50      - model_evaluation.metrics
51    metrics:
52      - reports/metrics.json:
53        cache: false
54    plots:
55      - reports/plots/confusion_matrix.png
```

## 18.8 Common Command Combinations

### Frequently Used Command Sequences

```
# Complete experiment cycle
dvc exp run && dvc commit && dvc push && git add . && \
git commit -m "Experiment" && git push

# Quick status check
dvc status && git status

# Full pipeline re-run
dvc repro --force && dvc push && git add dvc.lock && \
git commit -m "Pipeline re-run" && git push

# Sync with team
git pull && dvc pull

# Clean up experiments
dvc exp gc && dvc gc

# View pipeline and experiments
dvc dag && dvc exp show
```

## 19 Conclusion and Key Takeaways

### 19.1 What You've Accomplished

Throughout this comprehensive guide, you've built a complete, production-ready MLOps pipeline. Let's review the major achievements:

#### 1. End-to-End ML Pipeline:

- Data Ingestion with error handling
- Data Pre-Processing with text transformation
- Feature Engineering using TF-IDF
- Model Building with RandomForest
- Model Evaluation with comprehensive metrics

#### 2. Pipeline Automation:

- Automated workflow using DVC
- Intelligent dependency tracking
- Reproducible experiments
- Parameter-driven configuration

#### 3. Version Control Mastery:

- Git for code versioning
- DVC for data versioning
- AWS S3 for cloud backup
- Complete rollback capability

#### 4. Experiment Tracking:

- DVCLive integration
- Metric comparison across experiments
- Parameter tracking
- Visual experiment comparison

#### 5. Production Readiness:

- Comprehensive logging
- Error handling
- Modular architecture
- Team collaboration setup

### 19.2 The MLOps Journey: From Notebook to Production

#### The Transformation

##### Where You Started:

- Jupyter notebooks
- Local experiments

- Manual tracking
- Irreproducible results
- Solo development

**Where You Are Now:**

- Production-ready code
- Automated pipelines
- Systematic tracking
- Fully reproducible
- Collaboration-ready

### 19.3 Core Principles You've Learned

#### 19.3.1 1. Version Everything

Component	Tool	What's Tracked
Code	Git	Python scripts, configs
Data	DVC	Raw, interim, processed
Models	DVC	Trained model files
Parameters	Git	params.yaml
Metrics	DVC	Evaluation results
Dependencies	Git	requirements.txt

#### 19.3.2 2. Automate Repetitive Tasks

- **Before:** Manually run 5 Python scripts in sequence
- **After:** `dvc repro` runs entire pipeline

#### 19.3.3 3. Make Everything Reproducible

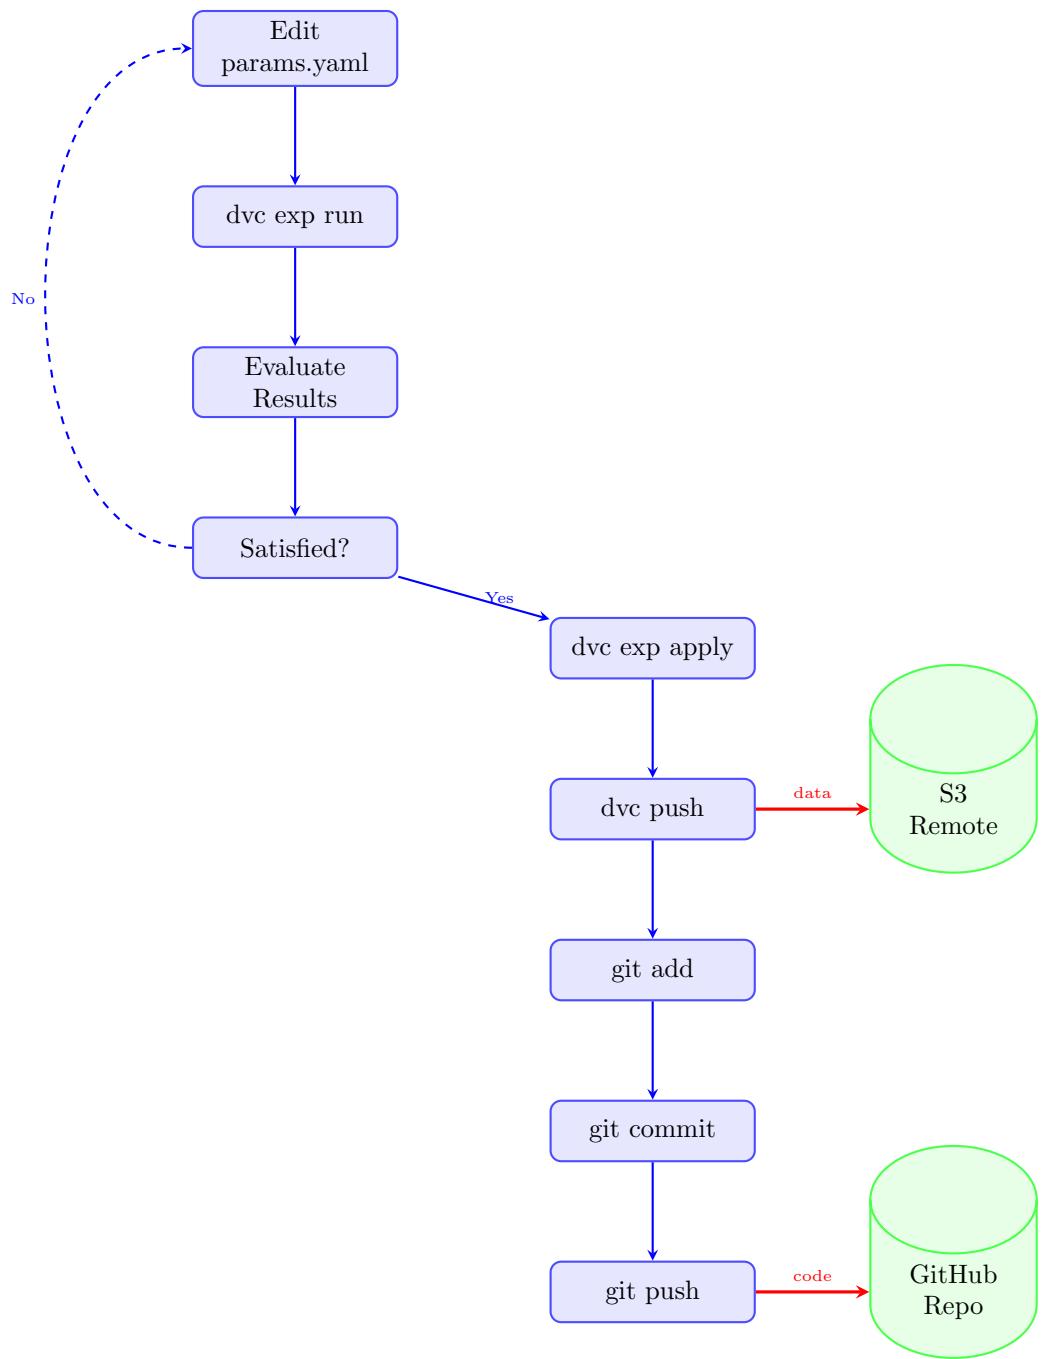
The Golden Rule:

```
git checkout <commit-hash>
dvc pull
# => Exact same code, data, and results
```

#### 19.3.4 4. Track Experiments Systematically

Experiment	Parameters	Accuracy	Decision
Exp 1	n_estimators=20	0.950	Baseline
Exp 2	n_estimators=50	0.970	Improved
Exp 3	n_estimators=100	0.972	Selected
Exp 4	n_estimators=200	0.971	Overfitting

## 19.4 The Complete Workflow Visualization



## 19.5 Key Lessons Learned

### 1. Data is as Important as Code:

- Models depend on data
- Data changes over time
- Version control applies to both

### 2. Automation Saves Time:

- Initial setup takes time
- Long-term benefits are huge

- Consistency across team

### 3. Logging is Essential:

- Debug issues faster
- Understand what happened
- Track pipeline progress

### 4. Parameters in Config Files:

- Easy experimentation
- Clear documentation
- Version controlled

### 5. Rollback Capability is Powerful:

- Revert to any **promoted and committed** experiment
- Compare past and current results reliably
- Enables safe and reproducible experimentation