

# Python Logging Module

## Complete Reference Guide

For MLOps and Production Systems

A Comprehensive Guide to Logging,  
Best Practices, and Real-World Implementation

Sujil S

[sujil19480@gmail.com](mailto:sujil19480@gmail.com)

December 25, 2025

## Contents

---

<b>1</b>	<b>Introduction to Python Logging</b>	<b>2</b>
1.1	What is Logging? . . . . .	2
1.1.1	Why Logging Matters . . . . .	2
1.2	Why Use Logging in MLOps? . . . . .	2
1.3	Logging vs Print Statements . . . . .	3
<b>2</b>	<b>Core Logging Components</b>	<b>4</b>
2.1	The Logging Architecture . . . . .	4
2.2	Logger Objects . . . . .	4
2.2.1	Creating a Logger . . . . .	4
2.2.2	Logger Methods . . . . .	4
2.3	Handler Objects . . . . .	4
2.3.1	Common Handler Types . . . . .	5
2.4	Formatter Objects . . . . .	5
2.4.1	Common Format Attributes . . . . .	5
2.5	Filter Objects . . . . .	6
<b>3</b>	<b>Logging Levels</b>	<b>7</b>
3.1	The Five Standard Levels . . . . .	7
3.2	DEBUG Level (10) . . . . .	7
3.3	INFO Level (20) . . . . .	7
3.4	WARNING Level (30) . . . . .	8
3.5	ERROR Level (40) . . . . .	8
3.6	CRITICAL Level (50) . . . . .	9
3.7	Setting Logging Levels . . . . .	9
<b>4</b>	<b>Handlers in Detail</b>	<b>11</b>
4.1	StreamHandler (Console Handler) . . . . .	11
4.1.1	Basic Usage . . . . .	11
4.1.2	Directing to stderr . . . . .	11
4.2	FileHandler . . . . .	11
4.2.1	Basic Usage . . . . .	11
4.2.2	File Modes . . . . .	12
4.3	RotatingFileHandler . . . . .	12
4.3.1	Usage Example . . . . .	12
4.4	TimedRotatingFileHandler . . . . .	13
4.4.1	Usage Example . . . . .	13
4.4.2	Common 'when' Values . . . . .	13
4.5	Multiple Handlers . . . . .	14
<b>5</b>	<b>Formatters and Message Formatting</b>	<b>16</b>
5.1	Creating Formatters . . . . .	16
5.1.1	Basic Formatter . . . . .	16
5.1.2	Custom Date Format . . . . .	16
5.2	Common Format Patterns . . . . .	16
5.2.1	Simple Format (Development) . . . . .	16
5.2.2	Standard Format (General Purpose) . . . . .	16
5.2.3	Detailed Format (Production) . . . . .	17
5.2.4	MLOps Format (Comprehensive) . . . . .	17

5.3	Complete Format Attributes Reference . . . . .	17
<b>6</b>	<b>Logging Configuration</b>	<b>18</b>
6.1	Basic Configuration with basicConfig() . . . . .	18
6.1.1	Simple Setup . . . . .	18
6.1.2	Detailed Configuration . . . . .	18
6.1.3	Multiple Handlers with basicConfig() . . . . .	18
6.2	Manual Configuration (Recommended) . . . . .	19
6.2.1	Complete Setup Example . . . . .	19
6.3	Configuration Using Dictionary . . . . .	20
6.4	Configuration from File . . . . .	21
6.4.1	YAML Configuration File . . . . .	21
<b>7</b>	<b>Best Practices for MLOps</b>	<b>23</b>
7.1	General Logging Best Practices . . . . .	23
7.2	MLOps-Specific Best Practices . . . . .	24
7.2.1	Log Model Training Progress . . . . .	24
7.2.2	Log Data Pipeline Steps . . . . .	25
7.2.3	Log Model Predictions with Timing . . . . .	25
7.2.4	Log Experiment Configuration . . . . .	26
7.3	Error Handling Best Practices . . . . .	26
<b>8</b>	<b>Complete Real-World Examples</b>	<b>28</b>
8.1	ML Training Script with Comprehensive Logging . . . . .	28
8.2	Data Pipeline with Logging . . . . .	32
8.3	Model Serving API with Logging . . . . .	37
<b>9</b>	<b>Advanced Logging Topics</b>	<b>40</b>
9.1	Logger Hierarchy . . . . .	40
9.2	Custom Log Levels . . . . .	40
9.3	Filters . . . . .	41
9.3.1	Context Filter Example . . . . .	41
9.4	Logging in Multiprocessing . . . . .	42
9.5	Structured Logging with JSON . . . . .	43
<b>10</b>	<b>Common Issues and Solutions</b>	<b>45</b>
10.1	Duplicate Log Messages . . . . .	45
10.2	Logs Not Appearing . . . . .	45
10.3	basicConfig() Not Working . . . . .	46
10.4	File Handler Not Writing . . . . .	46
10.5	Performance Issues . . . . .	47
<b>11</b>	<b>Quick Reference Guide</b>	<b>48</b>
11.1	Essential Commands Cheat Sheet . . . . .	48
11.2	Logging Levels Quick Reference . . . . .	50
11.3	Format String Placeholders . . . . .	50
11.4	Best Practices Checklist . . . . .	50
<b>12</b>	<b>Glossary</b>	<b>51</b>

<b>13 Conclusion</b>	<b>53</b>
13.1 Key Takeaways . . . . .	53
13.2 MLOps-Specific Recommendations . . . . .	54
13.3 Implementation Strategy . . . . .	54
13.4 Common Pitfalls to Avoid . . . . .	54
13.5 Next Steps . . . . .	55
13.6 Integration with Other Tools . . . . .	55
13.7 Final Recommendations . . . . .	56
13.8 Additional Resources . . . . .	56

## 1 Introduction to Python Logging

---

### 1.1 What is Logging?

**Logging** is the process of recording events, messages, and information during program execution. It provides a way to track what happens when software runs, which is essential for understanding application behavior, debugging issues, and monitoring systems in production.

#### 1.1.1 Why Logging Matters

Logging is critical for:

- **Debugging and Troubleshooting:** Identify and fix bugs efficiently
- **Monitoring:** Track application behavior in production
- **Auditing:** Maintain records for compliance and security
- **Performance Analysis:** Identify bottlenecks and optimization opportunities
- **Understanding Flow:** Trace execution path through complex systems

### 1.2 Why Use Logging in MLOps?

In Machine Learning Operations (MLOps), logging becomes even more critical:

#### 1. Model Training Monitoring:

- Track training epochs and iterations
- Record loss values and metrics
- Monitor convergence and performance
- Log hyperparameters for reproducibility

#### 2. Data Pipeline Tracking:

- Monitor data ingestion processes
- Log preprocessing steps
- Track data quality issues
- Record transformation operations

#### 3. Model Serving:

- Log prediction requests and responses
- Track API performance
- Monitor model confidence scores
- Record error rates

#### 4. Error Tracking:

- Identify failures quickly
- Debug production issues
- Track error patterns
- Generate alerts for critical errors

## 5. Performance Monitoring:

- Track inference times
- Monitor resource usage (CPU, GPU, memory)
- Identify performance degradation
- Optimize system bottlenecks

## 6. Reproducibility:

- Maintain audit trails for experiments
- Document model versions and configurations
- Track data versions used
- Enable experiment reproduction

### 1.3 Logging vs Print Statements

Many beginners use `print()` statements for debugging. However, in production code, the logging module is vastly superior:

print() Statements	logging Module
Always outputs to console	Flexible output destinations (file, console, remote)
No severity levels	Five severity levels (DEBUG to CRITICAL)
Difficult to filter messages	Easy filtering by level
No automatic timestamps	Automatic timestamps available
Cannot be disabled easily	Can be enabled/disabled by configuration
Not suitable for production	Production-ready and battle-tested
No structured format	Consistent, structured format
Limited context information	Rich context (filename, line number, function)
Performance impact on I/O	Optimized performance with lazy evaluation
No rotation capabilities	Automatic log file rotation

#### Warning

**Production Code Rule:** Never use `print()` statements in production code. Always use the logging module for better control, flexibility, and maintainability.

## 2 Core Logging Components

### 2.1 The Logging Architecture

The Python logging module follows a hierarchical architecture with four main components that work together to provide flexible and powerful logging capabilities.

1. **Loggers**: Entry point for logging messages in your application
2. **Handlers**: Determine where log records are sent (console, file, network)
3. **Formatters**: Define the layout and content of log messages
4. **Filters**: Provide fine-grained control over which log records are processed

### 2.2 Logger Objects

The **Logger** is the primary interface that applications use to write log messages.

#### 2.2.1 Creating a Logger

```
1 import logging
2
3 # Create a logger with module name (recommended best practice)
4 logger = logging.getLogger(__name__)
5
6 # Create a logger with custom name
7 logger = logging.getLogger('my_application')
8
9 # Get the root logger
10 root_logger = logging.getLogger()
```

#### Important Note

**Best Practice:** Always use `__name__` when creating loggers. This creates a hierarchical naming structure based on your module organization, making it easier to manage logging across large applications.

#### 2.2.2 Logger Methods

Loggers provide methods for different severity levels:

- `logger.debug(msg)` - Detailed diagnostic information for debugging
- `logger.info(msg)` - General informational messages about normal operation
- `logger.warning(msg)` - Warning messages indicating potential issues
- `logger.error(msg)` - Error messages for serious problems
- `logger.critical(msg)` - Critical error messages for severe failures
- `logger.exception(msg)` - Error with full traceback (use in except blocks)

### 2.3 Handler Objects

**Handlers** determine where log messages are sent. Multiple handlers can be attached to a single logger.

### 2.3.1 Common Handler Types

<b>StreamHandler</b>	Sends logs to console (stdout/stderr) - useful for development
<b>FileHandler</b>	Writes logs to a file - basic file logging
<b>RotatingFileHandler</b>	Rotates log files based on size - prevents files from growing too large
<b>TimedRotatingFileHandler</b>	Rotates log files based on time intervals - daily, weekly, etc.
<b>HTTPHandler</b>	Sends logs to HTTP server - for centralized logging
<b>SMTPHandler</b>	Sends logs via email - for critical alerts
<b>SysLogHandler</b>	Sends logs to Unix syslog daemon
<b>SocketHandler</b>	Sends logs over network socket
<b>QueueHandler</b>	Sends logs to queue for async processing

## 2.4 Formatter Objects

**Formatters** specify the layout and content of log messages, allowing you to customize how information is displayed.

### 2.4.1 Common Format Attributes

The following placeholders can be used in format strings:

- `%(asctime)s` - Human-readable timestamp of log record
- `%(name)s` - Name of the logger
- `%(levelname)s` - Logging level (DEBUG, INFO, WARNING, ERROR, CRITICAL)
- `%(message)s` - The actual log message
- `%(filename)s` - Source filename where log was called
- `%(lineno)d` - Line number in source code
- `%(funcName)s` - Function name where log was called
- `%(process)d` - Process ID
- `%(thread)d` - Thread ID
- `%(pathname)s` - Full pathname of source file
- `%(module)s` - Module name (filename without extension)
- `%(levelno)s` - Numeric logging level

## 2.5 Filter Objects

**Filters** provide fine-grained control over which log records are processed. They can be attached to both loggers and handlers.

Filters are useful for:

- Filtering by specific attributes
- Adding contextual information
- Implementing complex filtering logic
- Rate limiting log messages

## 3 Logging Levels

### 3.1 The Five Standard Levels

Python's logging module defines five severity levels in ascending order of importance. Each level has a numeric value:

Level Name	Numeric Value	Purpose
DEBUG	10	Detailed diagnostic information
INFO	20	Informational messages
WARNING	30	Warning messages
ERROR	40	Error messages
CRITICAL	50	Critical error messages

### 3.2 DEBUG Level (10)

**Purpose:** Most detailed information for diagnosing problems during development.

#### When to Use:

- Development and debugging phases
- Tracking variable values and state
- Understanding detailed program flow
- Loop iterations and function entry/exit
- Detailed API request/response information

#### DEBUG Level Examples in MLOps

```

1 logger.debug(f"Loading dataset from {data_path}")
2 logger.debug(f"Feature shape: {X.shape}, Target shape: {y.shape}")
3 logger.debug(f"Model parameters: {model.get_params()}")
4 logger.debug(f"Batch {batch_num}: input_shape={input_data.shape}")
5 logger.debug(f"GPU memory allocated: {torch.cuda.
    memory_allocated()}")

```

### 3.3 INFO Level (20)

**Purpose:** Confirmation that things are working as expected.

#### When to Use:

- Application startup and shutdown
- Successful completion of operations
- Milestone achievements
- Configuration information
- Normal operational messages

### INFO Level Examples in MLOps

```
1 logger.info("Model training started")
2 logger.info(f"Epoch {epoch}/{total_epochs}: loss={loss:.4f},
    accuracy={acc:.4f}")
3 logger.info("Model saved successfully to disk")
4 logger.info(f"API server started on port {port}")
5 logger.info(f"Loaded {len(dataset)} samples from database")
6 logger.info("Data preprocessing pipeline completed")
```

## 3.4 WARNING Level (30)

**Purpose:** Indication of unexpected events or potential problems that don't prevent the program from working.

### When to Use:

- Using deprecated features
- Resource constraints or limitations
- Recoverable errors
- Configuration issues that have fallbacks
- Performance concerns

### WARNING Level Examples in MLOps

```
1 logger.warning("Missing values detected in dataset")
2 logger.warning(f"GPU memory usage at {usage}%, close to limit")
3 logger.warning("Model confidence below threshold: 0.6")
4 logger.warning("Using default hyperparameters")
5 logger.warning("Training data smaller than recommended minimum")
6 logger.warning("Deprecated model architecture detected")
```

## 3.5 ERROR Level (40)

**Purpose:** Serious problems that prevented a specific function or operation from executing.

### When to Use:

- Failed operations
- Exception handling (non-critical)
- Resource unavailability
- Data validation failures
- API request failures

### ERROR Level Examples in MLOps

```

1 logger.error(f"Failed to load model from {model_path}")
2 logger.error("Database connection failed")
3 logger.error(f"Invalid input shape: expected {exp_shape}, got {actual_shape}")
4 logger.error("Prediction API returned error 500")
5 logger.error("Data quality check failed: too many outliers")
6 logger.error(f"Feature extraction failed for {feature_name}")

```

## 3.6 CRITICAL Level (50)

**Purpose:** Very serious errors indicating that the program may not be able to continue running.

### When to Use:

- System crashes or imminent crashes
- Data corruption
- Security breaches
- Unrecoverable errors
- Service unavailability

### CRITICAL Level Examples in MLOps

```

1 logger.critical("Out of memory - cannot continue training")
2 logger.critical("Model serving endpoint unreachable")
3 logger.critical("Critical security vulnerability detected")
4 logger.critical("Data pipeline completely failed")
5 logger.critical("Unable to connect to required external service")
6 logger.critical("System configuration corrupted")

```

## 3.7 Setting Logging Levels

You can set logging levels at both the logger and handler level:

```

1 import logging
2
3 logger = logging.getLogger(__name__)
4
5 # Set logger level (processes messages at this level and above)
6 logger.setLevel(logging.DEBUG)
7
8 # Set handler level (outputs messages at this level and above)
9 handler = logging.StreamHandler()
10 handler.setLevel(logging.WARNING)
11
12 logger.addHandler(handler)
13
14 # These will be processed but not output (handler level is WARNING)
15 logger.debug("Debug message")    # Not output
16 logger.info("Info message")     # Not output

```

```
17  
18 # These will be both processed and output  
19 logger.warning("Warning message") # Output  
20 logger.error("Error message") # Output
```

## Level Hierarchy

### How Levels Work:

- Logger set to INFO will process INFO, WARNING, ERROR, CRITICAL (but not DEBUG)
- Handler set to ERROR will only output ERROR and CRITICAL messages
- Messages below the set level are completely ignored
- Both logger AND handler must allow a level for it to be output

## 4 Handlers in Detail

### 4.1 StreamHandler (Console Handler)

The **StreamHandler** sends log output to streams like `sys.stdout` or `sys.stderr`. This is the most common handler for development.

#### 4.1.1 Basic Usage

```

1 import logging
2
3 logger = logging.getLogger(__name__)
4 logger.setLevel(logging.DEBUG)
5
6 # Create console handler
7 console_handler = logging.StreamHandler()
8 console_handler.setLevel(logging.INFO)
9
10 # Create formatter
11 formatter = logging.Formatter(
12     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
13 )
14 console_handler.setFormatter(formatter)
15
16 # Add handler to logger
17 logger.addHandler(console_handler)
18
19 # Use the logger
20 logger.info("This appears in console")
21 logger.debug("This won't appear (below handler level)")
```

#### 4.1.2 Directing to stderr

```

1 import sys
2 import logging
3
4 # Direct to stderr instead of stdout
5 error_handler = logging.StreamHandler(sys.stderr)
6 error_handler.setLevel(logging.ERROR)
7
8 logger.addHandler(error_handler)
```

### 4.2 FileHandler

The **FileHandler** writes log messages to a specified file on disk.

#### 4.2.1 Basic Usage

```

1 import logging
2
3 logger = logging.getLogger(__name__)
4 logger.setLevel(logging.DEBUG)
5
6 # Create file handler
7 file_handler = logging.FileHandler('application.log')
```

```

8 file_handler.setLevel(logging.ERROR)
9
10 formatter = logging.Formatter(
11     '%(asctime)s - %(levelname)s - %(message)s'
12 )
13 file_handler.setFormatter(formatter)
14
15 logger.addHandler(file_handler)
16
17 logger.error("This is written to application.log")
18 logger.info("This is not written (below handler level)")

```

#### 4.2.2 File Modes

- 'w' - Write mode: Overwrites existing file
- 'a' - Append mode: Appends to existing file (default)

```

1 # Overwrite mode - starts fresh each time
2 file_handler = logging.FileHandler('app.log', mode='w')
3
4 # Append mode - continues adding to existing file
5 file_handler = logging.FileHandler('app.log', mode='a')

```

### 4.3 RotatingFileHandler

The **RotatingFileHandler** automatically rotates log files when they reach a certain size. This is essential for production systems to prevent log files from consuming too much disk space.

#### 4.3.1 Usage Example

```

1 from logging.handlers import RotatingFileHandler
2 import logging
3
4 logger = logging.getLogger(__name__)
5 logger.setLevel(logging.DEBUG)
6
7 # Rotate after 10MB, keep 5 backup files
8 rotating_handler = RotatingFileHandler(
9     'app.log',
10    maxBytes=10*1024*1024, # 10 MB
11    backupCount=5
12 )
13
14 formatter = logging.Formatter(
15     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
16 )
17 rotating_handler.setFormatter(formatter)
18 logger.addHandler(rotating_handler)
19
20 logger.info("Logging with rotation enabled")

```

#### How File Rotation Works:

1. When `app.log` reaches 10MB, it's renamed to `app.log.1`

2. A new `app.log` file is created for current logging
3. On next rotation: `app.log.1` → `app.log.2`, new log → `app.log.1`
4. This continues until `backupCount` is reached
5. Oldest file is deleted when new rotation occurs beyond `backupCount`

#### Important Note

**Production Tip:** For production systems handling significant traffic, use `RotatingFileHandler` with appropriate size limits (10-50 MB) and backup counts (5-10 files) to balance between log retention and disk space.

## 4.4 TimedRotatingFileHandler

The `TimedRotatingFileHandler` rotates log files at specified time intervals rather than by file size.

### 4.4.1 Usage Example

```

1 from logging.handlers import TimedRotatingFileHandler
2 import logging
3
4 logger = logging.getLogger(__name__)
5 logger.setLevel(logging.INFO)
6
7 # Rotate daily at midnight, keep 30 days of logs
8 timed_handler = TimedRotatingFileHandler(
9     'app.log',
10    when='midnight',
11    interval=1,
12    backupCount=30
13 )
14
15 formatter = logging.Formatter(
16     '%(asctime)s - %(levelname)s - %(message)s'
17 )
18 timed_handler.setFormatter(formatter)
19 logger.addHandler(timed_handler)
20
21 logger.info("Logging with time-based rotation")

```

### 4.4.2 Common 'when' Values

Value	Type	Description
'S'	Seconds	Rotate every N seconds
'M'	Minutes	Rotate every N minutes
'H'	Hours	Rotate every N hours
'D'	Days	Rotate every N days
'midnight'	Special	Roll over at midnight
'W0'-'W6'	Weekday	Rotate on specific weekday (0=Monday)

### Time-Based Rotation Examples

```

1 # Rotate every hour
2 handler = TimedRotatingFileHandler('app.log', when='H',
3                                     interval=1)
4
5 # Rotate every 6 hours
6 handler = TimedRotatingFileHandler('app.log', when='H',
7                                     interval=6)
8
9 # Rotate daily at midnight
10 handler = TimedRotatingFileHandler('app.log', when='midnight')
11
12 # Rotate every Monday
13 handler = TimedRotatingFileHandler('app.log', when='W0')

```

## 4.5 Multiple Handlers

A single logger can have multiple handlers, each with different levels, formatters, and destinations. This is extremely useful for production systems.

### Complete Multi-Handler Setup

```

1 import logging
2 from logging.handlers import RotatingFileHandler
3
4 # Create logger
5 logger = logging.getLogger(__name__)
6 logger.setLevel(logging.DEBUG) # Logger processes all levels
7
8 # Console handler - INFO and above to console
9 console_handler = logging.StreamHandler()
10 console_handler.setLevel(logging.INFO)
11 console_formatter = logging.Formatter(
12     '%(levelname)s - %(message)s'
13 )
14 console_handler.setFormatter(console_formatter)
15
16 # File handler - ERROR and above to file
17 file_handler = RotatingFileHandler(
18     'errors.log',
19     maxBytes=5*1024*1024, # 5MB
20     backupCount=3
21 )
22 file_handler.setLevel(logging.ERROR)
23 file_formatter = logging.Formatter(
24     '%(asctime)s - %(name)s - %(levelname)s - ',
25     '[%(filename)s:%(lineno)d] - %(message)s'
26 )
27 file_handler.setFormatter(file_formatter)
28
29 # Debug file handler - all messages
30 debug_handler = RotatingFileHandler(
31     'debug.log',
32     maxBytes=10*1024*1024, # 10MB

```

```
33     backupCount=5
34 )
35 debug_handler.setLevel(logging.DEBUG)
36 debug_handler.setFormatter(file_formatter)
37
38 # Add all handlers to logger
39 logger.addHandler(console_handler)
40 logger.addHandler(file_handler)
41 logger.addHandler(debug_handler)
42
43 # Test different levels
44 logger.debug("Debug message")          # Only in debug.log
45 logger.info("Info message")           # Console + debug.log
46 logger.error("Error message")          # Console + errors.log +
    debug.log
```

This setup provides:

- Quick feedback in console (INFO and above)
- Separate error log for critical issues
- Complete debug log for thorough investigation

## 5 Formatters and Message Formatting

### 5.1 Creating Formatters

Formatters control the final output format of log messages. They can include various pieces of information about the log record.

#### 5.1.1 Basic Formatter

```

1 import logging
2
3 # Simple formatter with common fields
4 formatter = logging.Formatter(
5     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
6 )
7
8 # Apply to handler
9 handler = logging.StreamHandler()
10 handler.setFormatter(formatter)

```

#### 5.1.2 Custom Date Format

```

1 formatter = logging.Formatter(
2     fmt='%(asctime)s - %(levelname)s - %(message)s',
3     datefmt='%Y-%m-%d %H:%M:%S' # Custom date format
4 )
5
6 # Example output: 2024-12-17 14:30:45 - INFO - Message

```

## 5.2 Common Format Patterns

### 5.2.1 Simple Format (Development)

Best for quick development and debugging:

```
1 formatter = logging.Formatter('%(levelname)s - %(message)s')
```

Output example:

INFO - Model training started  
ERROR - Failed to load data

### 5.2.2 Standard Format (General Purpose)

Balanced format with essential information:

```

1 formatter = logging.Formatter(
2     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
3 )

```

Output example:

2024-12-17 14:30:45,123 - ml\_model - INFO - Model training started

### 5.2.3 Detailed Format (Production)

Includes file and line information for debugging:

```

1 formatter = logging.Formatter(
2     '%(asctime)s - %(name)s - %(levelname)s - ',
3     '[%(filename)s:%(lineno)d] - %(message)s',
4     datefmt='%Y-%m-%d %H:%M:%S',
5 )

```

Output example:

2024-12-17 14:30:45 - ml\_model - INFO - [train.py:42] - Model training started

### 5.2.4 MLOps Format (Comprehensive)

Includes process/thread info for distributed systems:

```

1 formatter = logging.Formatter(
2
3     '%(asctime)s - [PID:%(process)d TID:%(thread)d] - ',
4     '%(name)s - %(funcName)s - %(levelname)s - %(message)s',
5     datefmt='%Y-%m-%d %H:%M:%S',
6 )

```

Output example:

2024-12-17 14:30:45 - [PID:12345 TID:67890] - ml\_model -  
train\_model - INFO - Epoch 1 complete

## 5.3 Complete Format Attributes Reference

Attribute	Description
%(name)s	Logger name (typically module name)
%(levelname)s	Text logging level (DEBUG, INFO, etc.)
%(levelno)s	Numeric logging level (10, 20, etc.)
%(pathname)s	Full pathname of source file
%(filename)s	Filename portion of pathname
%(module)s	Module name (filename without .py)
%(lineno)d	Line number where log was called
%(funcName)s	Function name where log was called
%(created)f	Time when LogRecord was created (seconds since epoch)
%(asctime)s	Human-readable time
%(msecs)d	Millisecond portion of time
%(relativeCreated)d	Time in ms since logging module loaded
%(thread)d	Thread ID
%(threadName)s	Thread name
%(process)d	Process ID
%(message)s	The logged message

## 6 Logging Configuration

### 6.1 Basic Configuration with basicConfig()

The `basicConfig()` function provides a quick way to configure logging for simple applications.

#### 6.1.1 Simple Setup

```

1 import logging
2
3 # Basic configuration with default settings
4 logging.basicConfig(level=logging.INFO)
5
6 # Use root logger
7 logging.info("This is an info message")
8 logging.error("This is an error message")

```

#### 6.1.2 Detailed Configuration

```

1 import logging
2
3 logging.basicConfig(
4     level=logging.DEBUG,
5     format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
6     datefmt='%Y-%m-%d %H:%M:%S',
7     filename='app.log',
8     filemode='w' # 'w' to overwrite, 'a' to append
9 )
10
11 logging.info("Application started")
12 logging.debug("Debug information")

```

#### 6.1.3 Multiple Handlers with basicConfig()

```

1 import logging
2
3 logging.basicConfig(
4     level=logging.DEBUG,
5     format='%(asctime)s - %(levelname)s - %(message)s',
6     handlers=[
7         logging.FileHandler('debug.log'),
8         logging.StreamHandler()
9     ]
10 )
11
12 logging.info("Logs to both file and console")

```

#### Warning

**Important Limitation:** `basicConfig()` only works the first time it's called. Subsequent calls are ignored unless you use the `force=True` parameter (Python 3.8+).

```

1 # Python 3.8+ only
2 logging.basicConfig(level=logging.DEBUG, force=True)

```

## 6.2 Manual Configuration (Recommended)

For production applications, manual configuration provides more control and flexibility.

### 6.2.1 Complete Setup Example

```
 1 import logging
 2 from logging.handlers import RotatingFileHandler
 3
 4 def setup_logger(name, log_file, level=logging.INFO):
 5     """
 6         Function to setup logger with console and file handlers.
 7
 8     Args:
 9         name: Logger name
10         log_file: Path to log file
11         level: Logging level
12
13     Returns:
14         Configured logger object
15     """
16
17     # Create logger
18     logger = logging.getLogger(name)
19     logger.setLevel(level)
20
21     # Prevent duplicate handlers
22     if logger.handlers:
23         return logger
24
25     # Create formatters
26     console_formatter = logging.Formatter(
27         '%(levelname)s - %(message)s'
28     )
29     file_formatter = logging.Formatter(
30         '%(asctime)s - %(name)s - %(levelname)s - ',
31         '[%(filename)s:%(lineno)d] - %(message)s',
32         datefmt='%Y-%m-%d %H:%M:%S'
33     )
34
35     # Console handler
36     console_handler = logging.StreamHandler()
37     console_handler.setLevel(logging.INFO)
38     console_handler.setFormatter(console_formatter)
39
40     # File handler with rotation
41     file_handler = RotatingFileHandler(
42         log_file,
43         maxBytes=10*1024*1024, # 10MB
44         backupCount=5
45     )
46     file_handler.setLevel(logging.DEBUG)
47     file_handler.setFormatter(file_formatter)
48
49     # Add handlers to logger
50     logger.addHandler(console_handler)
51     logger.addHandler(file_handler)
52
```

```

53     return logger
54
55 # Usage
56 logger = setup_logger('ml_pipeline', 'pipeline.log')
57 logger.info("Logger configured successfully")
58 logger.debug("This goes to file only")

```

### 6.3 Configuration Using Dictionary

For complex applications, dictionary-based configuration provides a clean, declarative approach.

```

1 import logging
2 import logging.config
3
4 LOGGING_CONFIG = {
5     'version': 1,
6     'disable_existing_loggers': False,
7
8     'formatters': {
9         'standard': {
10             'format': '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
11         },
12         'detailed': {
13             'format': '%(asctime)s - %(name)s - %(levelname)s - '
14                         '[%(filename)s:%(lineno)d] - %(message)s'
15         }
16     },
17
18     'handlers': {
19         'console': {
20             'class': 'logging.StreamHandler',
21             'level': 'INFO',
22             'formatter': 'standard',
23             'stream': 'ext://sys.stdout'
24         },
25         'file': {
26             'class': 'logging.handlers.RotatingFileHandler',
27             'level': 'DEBUG',
28             'formatter': 'detailed',
29             'filename': 'app.log',
30             'maxBytes': 10485760, # 10MB
31             'backupCount': 5
32         },
33         'error_file': {
34             'class': 'logging.FileHandler',
35             'level': 'ERROR',
36             'formatter': 'detailed',
37             'filename': 'errors.log'
38         }
39     },
40
41     'loggers': {
42         '': { # Root logger
43             'handlers': ['console', 'file', 'error_file'],
44             'level': 'DEBUG',
45             'propagate': False
46         },

```

```

47         'ml_model': { # Specific logger
48             'handlers': ['console', 'file'],
49             'level': 'INFO',
50             'propagate': False
51         }
52     }
53 }
54
55 # Apply configuration
56 logging.config.dictConfig(LOGGING_CONFIG)
57
58 # Use loggers
59 root_logger = logging.getLogger()
60 ml_logger = logging.getLogger('ml_model')
61
62 root_logger.info("Root logger message")
63 ml_logger.info("ML model logger message")

```

## 6.4 Configuration from File

You can store configuration in external files for easy management.

### 6.4.1 YAML Configuration File

Create `logging_config.yaml`:

```

1 version: 1
2 disable_existing_loggers: False
3
4 formatters:
5     standard:
6         format: '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
7
8 handlers:
9     console:
10        class: logging.StreamHandler
11        level: INFO
12        formatter: standard
13        stream: ext://sys.stdout
14
15 file:
16        class: logging.handlers.RotatingFileHandler
17        level: DEBUG
18        formatter: standard
19        filename: app.log
20        maxBytes: 10485760
21        backupCount: 5
22
23 loggers:
24     '':
25         handlers: [console, file]
26         level: DEBUG
27         propagate: False

```

Load and apply:

```

1 import logging.config
2 import yaml

```

```
3
4 with open('logging_config.yaml', 'r') as f:
5     config = yaml.safe_load(f)
6     logging.config.dictConfig(config)
7
8 logger = logging.getLogger(__name__)
9 logger.info("Configuration loaded from YAML")
```

## 7 Best Practices for MLOps

### 7.1 General Logging Best Practices

#### 1. Use `__name__` for logger names

This creates a hierarchical logger structure based on your module organization:

```

1 # Good practice
2 logger = logging.getLogger(__name__)
3
4 # Avoid
5 logger = logging.getLogger('my_logger')
6

```

#### 2. Use appropriate logging levels

- DEBUG: Development and debugging only
- INFO: Production tracking and milestones
- WARNING: Recoverable issues and deprecations
- ERROR: Failures that affect functionality
- CRITICAL: System-level failures

#### 3. Use lazy formatting for performance

```

1 # Good - lazy evaluation (only formats if logged)
2 logger.info("Processing batch %d with size %d", batch_id, size)
3
4 # Bad - eager evaluation (always formats, even if not logged)
5 logger.info(f"Processing batch {batch_id} with size {size}")
6
7 # Also bad
8 logger.info("Processing batch " + str(batch_id) + " with size " +
    str(size))
9

```

#### 4. Use `logger.exception()` in except blocks

```

1 try:
2     model.predict(data)
3 except Exception as e:
4     logger.exception("Prediction failed")
5     # Automatically includes full traceback
6     raise # Re-raise after logging
7

```

#### 5. Don't log sensitive information

```

1 # Bad - security risk
2 logger.info(f"User password: {password}")
3 logger.info(f"API key: {api_key}")
4
5 # Good - no sensitive data
6 logger.info(f"User authentication attempt for: {username}")
7 logger.info("API authentication successful")
8

```

## 6. Configure logging once at startup

Set up logging configuration at application entry point, not in every module.

## 7. Use structured logging for complex data

```

1 # Use extra parameter for structured data
2 logger.info(
3     "Model evaluation complete",
4     extra={
5         'accuracy': 0.95,
6         'precision': 0.93,
7         'recall': 0.97,
8         'f1_score': 0.95
9     }
10 )
11

```

## 8. Avoid excessive logging

```

1 # Bad - logs every iteration (too verbose)
2 for i in range(10000):
3     logger.debug(f"Processing item {i}")
4
5 # Good - log periodically
6 for i in range(10000):
7     if i % 1000 == 0:
8         logger.info(f"Processed {i}/10000 items")
9

```

## 7.2 MLOps-Specific Best Practices

### 7.2.1 Log Model Training Progress

```

1 import logging
2
3 logger = logging.getLogger(__name__)
4
5 def train_model(model, train_loader, val_loader, epochs):
6     logger.info("=" * 60)
7     logger.info("STARTING MODEL TRAINING")
8     logger.info("=" * 60)
9     logger.info(f"Training samples: {len(train_loader.dataset)}")
10    logger.info(f"Validation samples: {len(val_loader.dataset)}")
11    logger.info(f"Epochs: {epochs}")
12
13    for epoch in range(epochs):
14        # Training phase
15        train_loss = train_epoch(model, train_loader)
16
17        # Validation phase
18        val_loss, val_acc = validate(model, val_loader)
19
20        # Log progress
21        logger.info(
22            f"Epoch {epoch+1}/{epochs}: "
23            f"train_loss={train_loss:.4f}, "
24            f"val_loss={val_loss:.4f}, "

```

```

25             f"val_acc={val_acc:.4f}"
26         )
27
28     # Log warnings if needed
29     if val_loss > train_loss * 1.5:
30         logger.warning("Possible overfitting detected")
31
32     logger.info("Training completed successfully")
33     logger.info(f"Final validation accuracy: {val_acc:.4f}")
34     logger.info("=" * 60)

```

## 7.2.2 Log Data Pipeline Steps

```

1 def process_data_pipeline(data_source):
2     logger.info(f"Loading data from {data_source}")
3     data = load_data(data_source)
4     logger.debug(f"Raw data shape: {data.shape}")
5
6     logger.info("Starting data preprocessing")
7     missing_count = data.isnull().sum().sum()
8     if missing_count > 0:
9         logger.warning(f"Found {missing_count} missing values")
10
11    cleaned_data = preprocess(data)
12    logger.debug(f"Cleaned data shape: {cleaned_data.shape}")
13
14    logger.info("Feature engineering started")
15    features = engineer_features(cleaned_data)
16    logger.info(f"Created {features.shape[1]} features")
17    logger.debug(f"Feature names: {list(features.columns)}")
18
19    logger.info("Data pipeline completed successfully")
20    logger.info(f"Final dataset shape: {features.shape}")
21
22    return features

```

## 7.2.3 Log Model Predictions with Timing

```

1 import time
2
3 def predict_with_logging(model, input_data, request_id=None):
4     if request_id is None:
5         request_id = str(time.time())
6
7     logger.info(f"[{request_id}] Prediction request received")
8     logger.debug(f"[{request_id}] Input shape: {input_data.shape}")
9
10    try:
11        start_time = time.time()
12        predictions = model.predict(input_data)
13        inference_time = time.time() - start_time
14
15        logger.info(
16            f"[{request_id}] Prediction completed in "
17            f"{inference_time:.3f}s"
18        )

```

```

19         logger.debug(
20             f"[{request_id}] Predictions: {predictions[:5]}...")
21     # Log first 5 only
22
23     return predictions
24
25 except Exception as e:
26     logger.exception(f"[{request_id}] Prediction failed")
27     raise

```

#### 7.2.4 Log Experiment Configuration

```

1 def log_experiment_config(config):
2     logger.info("=" * 70)
3     logger.info("EXPERIMENT CONFIGURATION")
4     logger.info("=" * 70)
5
6     for section, params in config.items():
7         logger.info(f"{section}:")
8         if isinstance(params, dict):
9             for key, value in params.items():
10                 logger.info(f"  {key}: {value}")
11         else:
12             logger.info(f"  {params}")
13
14     logger.info("=" * 70)
15
16 # Usage
17 config = {
18     'model': {
19         'type': 'RandomForest',
20         'n_estimators': 100,
21         'max_depth': 10
22     },
23     'training': {
24         'epochs': 50,
25         'batch_size': 32,
26         'learning_rate': 0.001
27     },
28     'data': {
29         'train_split': 0.8,
30         'random_state': 42
31     }
32 }
33
34 log_experiment_config(config)

```

### 7.3 Error Handling Best Practices

```

1 def load_and_process_data(file_path):
2     try:
3         logger.info(f"Attempting to load data from {file_path}")
4         data = pd.read_csv(file_path)
5
6         logger.info(f"Data loaded successfully: {data.shape}")
7         logger.debug(f"Columns: {list(data.columns)}")

```

```
8      # Data validation
9      if data.empty:
10         logger.error("Loaded data is empty")
11         raise ValueError("Empty dataset")
12
13     # Processing
14     logger.info("Starting data processing")
15     processed_data = process(data)
16
17     logger.info("Data processing completed successfully")
18     return processed_data
19
20
21 except FileNotFoundError:
22     logger.error(f"File not found: {file_path}")
23     logger.error("Please check the file path and try again")
24     raise
25
26 except pd.errors.EmptyDataError:
27     logger.error(f"Empty data file: {file_path}")
28     raise
29
30 except pd.errors.ParserError as e:
31     logger.error(f"Error parsing CSV file: {str(e)}")
32     logger.error("Check file format and encoding")
33     raise
34
35 except MemoryError:
36     logger.critical(
37         f"Out of memory while loading {file_path}"
38     )
39     logger.critical("Try loading data in chunks")
40     raise
41
42 except Exception as e:
43     logger.exception(
44         f"Unexpected error loading data from {file_path}"
45     )
46     raise
```

## 8 Complete Real-World Examples

### 8.1 ML Training Script with Comprehensive Logging

```

1 import logging
2 from logging.handlers import RotatingFileHandler
3 import time
4 import numpy as np
5 from sklearn.model_selection import train_test_split
6 from sklearn.ensemble import RandomForestClassifier
7 from sklearn.metrics import (
8     accuracy_score, precision_score,
9     recall_score, f1_score
10 )
11
12 # =====
13 # Logger Configuration
14 # =====
15
16 def setup_training_logger():
17     """Setup logger for ML training with file rotation"""
18     logger = logging.getLogger(__name__)
19     logger.setLevel(logging.DEBUG)
20
21     # Console handler - INFO and above
22     console_handler = logging.StreamHandler()
23     console_handler.setLevel(logging.INFO)
24     console_formatter = logging.Formatter(
25         '%(levelname)s - %(message)s'
26     )
27     console_handler.setFormatter(console_formatter)
28
29     # File handler - all messages with rotation
30     file_handler = RotatingFileHandler(
31         'training.log',
32         maxBytes=10*1024*1024, # 10MB
33         backupCount=5
34     )
35     file_handler.setLevel(logging.DEBUG)
36     file_formatter = logging.Formatter(
37         '%(asctime)s - %(name)s - %(levelname)s - ',
38         '[%(filename)s:%(lineno)d] - %(message)s',
39         datefmt='%Y-%m-%d %H:%M:%S'
40     )
41     file_handler.setFormatter(file_formatter)
42
43     # Add handlers
44     logger.addHandler(console_handler)
45     logger.addHandler(file_handler)
46
47     return logger
48
49 logger = setup_training_logger()
50
51 # =====
52 # Training Functions
53 # =====
54

```

```
55 def validate_data(X, y):
56     """Validate input data with logging"""
57     logger.debug("Validating input data")
58
59     if X is None or y is None:
60         logger.error("Input data is None")
61         raise ValueError("Invalid input data")
62
63     if len(X) != len(y):
64         logger.error(
65             f"Shape mismatch: X has {len(X)} samples, "
66             f"y has {len(y)} samples"
67         )
68         raise ValueError("X and y length mismatch")
69
70     if len(X) < 100:
71         logger.warning(
72             f"Small dataset: only {len(X)} samples"
73         )
74
75     logger.debug(f"Data validation passed: {len(X)} samples")
76
77 def train_model(X, y, config):
78     """
79     Train ML model with comprehensive logging
80
81     Args:
82         X: Feature matrix
83         y: Target vector
84         config: Training configuration dictionary
85
86     Returns:
87         Trained model
88     """
89
90     logger.info("=" * 70)
91     logger.info("STARTING MODEL TRAINING")
92     logger.info("=" * 70)
93
94     # Log configuration
95     logger.info("Training Configuration:")
96     for key, value in config.items():
97         logger.info(f" {key}: {value}")
98     logger.info("-" * 70)
99
100    try:
101        # Validate data
102        validate_data(X, y)
103
104        # Split data
105        logger.info("Splitting data into train and test sets")
106        X_train, X_test, y_train, y_test = train_test_split(
107            X,
108            test_size=config['test_size'],
109            random_state=config['random_state']
110        )
111
112        logger.info(f"Training set: {len(X_train)} samples")
```

```
113     logger.info(f"Test set: {len(X_test)} samples")
114     logger.debug(f"Training features shape: {X_train.shape}")
115     logger.debug(f"Test features shape: {X_test.shape}")
116
117     # Check class balance
118     unique, counts = np.unique(y_train, return_counts=True)
119     logger.debug(f"Class distribution: {dict(zip(unique, counts))}")
120
121     if max(counts) / min(counts) > 10:
122         logger.warning("Highly imbalanced dataset detected")
123
124     # Initialize model
125     logger.info("Initializing Random Forest model")
126     logger.debug(f"Model parameters: {config}")
127
128     model = RandomForestClassifier(
129         n_estimators=config['n_estimators'],
130         max_depth=config['max_depth'],
131         random_state=config['random_state'],
132         n_jobs=-1,    # Use all CPU cores
133         verbose=0
134     )
135
136     # Train model
137     logger.info("Starting model training...")
138     start_time = time.time()
139
140     model.fit(X_train, y_train)
141
142     training_time = time.time() - start_time
143     logger.info(
144         f"Training completed in {training_time:.2f} seconds"
145     )
146
147     # Evaluate on training set
148     logger.debug("Evaluating on training set")
149     y_train_pred = model.predict(X_train)
150     train_acc = accuracy_score(y_train, y_train_pred)
151     logger.debug(f"Training accuracy: {train_acc:.4f}")
152
153     # Evaluate on test set
154     logger.info("Evaluating model on test set")
155     start_eval = time.time()
156     y_pred = model.predict(X_test)
157     eval_time = time.time() - start_eval
158
159     logger.debug(f"Evaluation time: {eval_time:.3f}s")
160
161     # Calculate metrics
162     accuracy = accuracy_score(y_test, y_pred)
163     precision = precision_score(
164         y_test, y_pred, average='weighted', zero_division=0
165     )
166     recall = recall_score(
167         y_test, y_pred, average='weighted', zero_division=0
168     )
169     f1 = f1_score(
170         y_test, y_pred, average='weighted', zero_division=0
171     )
```

```
170     )
171
172     # Log results
173     logger.info("=" * 70)
174     logger.info("MODEL EVALUATION RESULTS")
175     logger.info("=" * 70)
176     logger.info(f"Accuracy: {accuracy:.4f}")
177     logger.info(f"Precision: {precision:.4f}")
178     logger.info(f"Recall: {recall:.4f}")
179     logger.info(f"F1-Score: {f1:.4f}")
180     logger.info("=" * 70)
181
182     # Feature importance
183     if hasattr(model, 'feature_importances_'):
184         importance = model.feature_importances_
185         logger.debug(f"Feature importances: {importance}")
186         top_features = np.argsort(importance)[-5:][::-1]
187         logger.info(
188             f"Top 5 important features: {top_features.tolist()}"
189         )
190
191     # Performance warnings
192     if accuracy < 0.7:
193         logger.warning(
194             f"Low accuracy: {accuracy:.4f}. "
195             "Consider feature engineering or hyperparameter
tuning"
196         )
197
198     if training_time > 300: # 5 minutes
199         logger.warning(
200             f"Long training time: {training_time:.2f}s. "
201             "Consider reducing model complexity"
202         )
203
204     logger.info("Model training pipeline completed successfully"
205     )
206
207     return model
208
209 except ValueError as e:
210     logger.error(f"Value error during training: {str(e)}")
211     raise
212
213 except MemoryError:
214     logger.critical(
215         "Out of memory during model training. "
216         "Try reducing dataset size or model complexity"
217     )
218     raise
219
220 except Exception as e:
221     logger.exception("Unexpected error during model training")
222     raise
223
224 # =====
225 # Main Execution
226 # =====
```

```

226
227 if __name__ == "__main__":
228     logger.info("Script started")
229     logger.info(f"Python logging version: {logging.__version__}")
230
231     # Configuration
232     config = {
233         'n_estimators': 100,
234         'max_depth': 10,
235         'test_size': 0.2,
236         'random_state': 42
237     }
238
239     try:
240         # Generate sample data
241         logger.info("Generating sample dataset")
242         np.random.seed(42)
243         X = np.random.rand(1000, 10)
244         y = np.random.randint(0, 2, 1000)
245         logger.info("Sample data generated successfully")
246
247         # Train model
248         model = train_model(X, y, config)
249
250         logger.info("Script completed successfully")
251
252     except KeyboardInterrupt:
253         logger.warning("Script interrupted by user")
254     except Exception as e:
255         logger.critical("Script failed with critical error")
256         logger.exception("Error details")
257         raise

```

## 8.2 Data Pipeline with Logging

```

1 import logging
2 import pandas as pd
3 import numpy as np
4 from datetime import datetime
5
6 logger = logging.getLogger(__name__)
7
8 class DataPipeline:
9     """
10     Data pipeline with comprehensive logging
11
12     Handles data loading, cleaning, and feature engineering
13     with detailed logging at each step.
14     """
15
16     def __init__(self, config):
17         self.config = config
18         logger.info("DataPipeline initialized")
19         logger.debug(f"Pipeline config: {config}")
20
21     def load_data(self, file_path):
22         """Load data with error handling and logging"""

```

```
23     try:
24         logger.info(f"Loading data from: {file_path}")
25         start_time = datetime.now()
26
27         data = pd.read_csv(file_path)
28
29         load_time = (datetime.now() - start_time).total_seconds()
30     ()
31         logger.info(f"Data loaded in {load_time:.2f} seconds")
32         logger.info(f"Dataset shape: {data.shape}")
33         logger.debug(f"Columns ({len(data.columns)}): {list(data
34 .columns)}")
35         logger.debug(f"Memory usage: {data.memory_usage(deep=
36 True).sum() / 1024**2:.2f} MB")
37
38     # Log data types
39     logger.debug("Data types:")
40     for col, dtype in data.dtypes.items():
41         logger.debug(f" {col}: {dtype}")
42
43     return data
44
45 except FileNotFoundError:
46     logger.error(f"File not found: {file_path}")
47     logger.error("Please verify the file path")
48     raise
49
50 except pd.errors.ParserError as e:
51     logger.error(f"Error parsing CSV: {str(e)}")
52     logger.error("Check file format, delimiter, and encoding")
53     raise
54
55 except pd.errors.EmptyDataError:
56     logger.error(f"Empty file: {file_path}")
57     raise
58
59 except Exception as e:
60     logger.exception("Unexpected error loading data")
61     raise
62
63 def clean_data(self, data):
64     """Clean data with detailed logging"""
65     logger.info("Starting data cleaning")
66
67     original_shape = data.shape
68     original_rows = original_shape[0]
69     logger.debug(f"Original shape: {original_shape}")
70
71     # Check missing values
72     missing_count = data.isnull().sum().sum()
73     if missing_count > 0:
74         missing_pct = (missing_count / data.size) * 100
75         logger.warning(
76             f"Found {missing_count} missing values "
77             f"({missing_pct:.2f}% of data)"
78         )
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
```

```
77         # Log missing values per column
78         missing_cols = data.isnull().sum()
79         missing_cols = missing_cols[missing_cols > 0]
80         for col, count in missing_cols.items():
81             col_pct = (count / len(data)) * 100
82             logger.debug(
83                 f" {col}: {count} missing ({col_pct:.1f}%)"
84             )
85
86         # Handle missing values
87         logger.info("Handling missing values")
88         numeric_cols = data.select_dtypes(
89             include=[np.number]
90         ).columns
91         data[numeric_cols] = data[numeric_cols].fillna(
92             data[numeric_cols].mean()
93         )
94         logger.info("Numeric missing values filled with mean")
95
96         categorical_cols = data.select_dtypes(
97             include=['object']
98         ).columns
99         data[categorical_cols] = data[categorical_cols].fillna(
100            data[categorical_cols].mode().iloc[0]
101        )
102         logger.info("Categorical missing values filled with mode")
103     ")
104 else:
105     logger.info("No missing values found")
106
107     # Remove duplicates
108     duplicates = data.duplicated().sum()
109     if duplicates > 0:
110         dup_pct = (duplicates / len(data)) * 100
111         logger.warning(
112             f"Found {duplicates} duplicate rows ({dup_pct:.2f}%)"
113         )
114         data = data.drop_duplicates()
115         logger.info(f"Removed {duplicates} duplicate rows")
116     else:
117         logger.info("No duplicates found")
118
119     # Remove outliers (optional)
120     numeric_cols = data.select_dtypes(include=[np.number]).columns
121     outliers_removed = 0
122     for col in numeric_cols:
123         Q1 = data[col].quantile(0.25)
124         Q3 = data[col].quantile(0.75)
125         IQR = Q3 - Q1
126         lower_bound = Q1 - 1.5 * IQR
127         upper_bound = Q3 + 1.5 * IQR
128
129         outliers = ((data[col] < lower_bound) |
130                     (data[col] > upper_bound)).sum()
131         if outliers > 0:
```

```
131             logger.debug(f" {col}: {outliers} outliers detected")
132         outliers_removed += outliers
133
134     if outliers_removed > 0:
135         logger.info(f"Total outliers detected: {outliers_removed}")
136
137     final_shape = data.shape
138     rows_removed = original_rows - final_shape[0]
139
140     logger.info("Data cleaning completed")
141     logger.info(f"Final shape: {final_shape}")
142     logger.info(f"Rows removed: {rows_removed}")
143
144     return data
145
146 def feature_engineering(self, data):
147     """Perform feature engineering with logging"""
148     logger.info("Starting feature engineering")
149
150     original_features = data.shape[1]
151     logger.debug(f"Original features: {original_features}")
152
153     try:
154         # Example: Create interaction features
155         logger.debug("Creating derived features")
156
157         # Add timestamp features if date column exists
158         date_cols = data.select_dtypes(
159             include=['datetime64'])
160         .columns
161         if len(date_cols) > 0:
162             logger.info(
163                 f"Processing {len(date_cols)} datetime columns"
164             )
165             for col in date_cols:
166                 data[f'{col}_year'] = data[col].dt.year
167                 data[f'{col}_month'] = data[col].dt.month
168                 data[f'{col}_day'] = data[col].dt.day
169                 logger.debug(
170                     f"Created time features from {col}"
171                 )
172
173         new_features = data.shape[1]
174         added_features = new_features - original_features
175
176         logger.info("Feature engineering completed")
177         logger.info(f"Total features: {new_features}")
178         logger.info(f"Added features: {added_features}")
179         logger.debug(f"New columns: {list(data.columns)}")
180
181         return data
182
183     except Exception as e:
184         logger.exception("Error during feature engineering")
185         raise
```

```
187     def run(self, file_path):
188         """Execute complete pipeline"""
189         logger.info("=" * 70)
190         logger.info("STARTING DATA PIPELINE")
191         logger.info("=" * 70)
192
193         pipeline_start = datetime.now()
194
195         try:
196             # Load data
197             data = self.load_data(file_path)
198
199             # Clean data
200             data = self.clean_data(data)
201
202             # Feature engineering
203             data = self.feature_engineering(data)
204
205             pipeline_time = (
206                 datetime.now() - pipeline_start
207             ).total_seconds()
208
209             logger.info("=" * 70)
210             logger.info("PIPELINE COMPLETED SUCCESSFULLY")
211             logger.info(
212                 f"Total pipeline time: {pipeline_time:.2f} seconds"
213             )
214             logger.info("=" * 70)
215
216             return data
217
218         except Exception as e:
219             logger.critical("Pipeline failed")
220             logger.exception("Pipeline error details")
221             raise
222
223 # Usage example
224 if __name__ == "__main__":
225     logging.basicConfig(
226         level=logging.DEBUG,
227         format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
228         ,
229         handlers=[
230             logging.FileHandler('pipeline.log'),
231             logging.StreamHandler()
232         ]
233     )
234
235     config = {'version': '1.0', 'mode': 'production'}
236     pipeline = DataPipeline(config)
237
238     # Run pipeline
239     try:
240         result = pipeline.run('data.csv')
241         logger.info("Pipeline execution successful")
242     except Exception:
243         logger.error("Pipeline execution failed")
```

### 8.3 Model Serving API with Logging

```

1 import logging
2 from flask import Flask, request, jsonify
3 import numpy as np
4 import time
5 import uuid
6
7 app = Flask(__name__)
8
9 # Configure logging
10 logging.basicConfig(
11     level=logging.INFO,
12     format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
13     handlers=[
14         logging.FileHandler('api.log'),
15         logging.StreamHandler()
16     ]
17 )
18
19 logger = logging.getLogger(__name__)
20
21 # Global model variable
22 model = None
23
24 @app.before_first_request
25 def setup():
26     """Initialize application"""
27     global model
28     logger.info("=" * 70)
29     logger.info("INITIALIZING MODEL SERVING API")
30     logger.info("=" * 70)
31
32     try:
33         logger.info("Loading model")
34         # model = load_model('model.pkl') # Placeholder
35         logger.info("Model loaded successfully")
36         logger.info("API initialization complete")
37     except Exception as e:
38         logger.critical("Failed to load model")
39         logger.exception("Model loading error")
40         raise
41
42 @app.route('/predict', methods=['POST'])
43 def predict():
44     """Prediction endpoint with comprehensive logging"""
45     request_id = str(uuid.uuid4())[:8]
46
47     logger.info(f"[{request_id}] Prediction request received")
48     logger.debug(f"[{request_id}] Request headers: {dict(request.headers)}")
49
50     try:
51         # Parse request
52         data = request.get_json()
53         logger.debug(f"[{request_id}] Request data keys: {list(data.keys())}")
54

```

```
55     # Validate request
56     if not data or 'features' not in data:
57         logger.warning(
58             f"[{request_id}] Invalid request: missing 'features'
59         )
60         return jsonify({'error': 'Missing features field'}), 400
61
62     features = np.array(data['features'])
63     logger.debug(f"[{request_id}] Features shape: {features.
64     shape}")
65
66     # Validate input shape
67     if features.ndim != 2:
68         logger.error(
69             f"[{request_id}] Invalid shape: {features.shape}. "
70             f"Expected 2D array"
71         )
72         return jsonify({
73             'error': f'Invalid feature shape: {features.shape}'
74         }), 400
75
76     # Check for invalid values
77     if np.isnan(features).any():
78         logger.error(f"[{request_id}] NaN values in input")
79         return jsonify({'error': 'NaN values in features'}), 400
80
81     if np.isinf(features).any():
82         logger.error(f"[{request_id}] Inf values in input")
83         return jsonify({'error': 'Inf values in features'}), 400
84
85     # Make prediction
86     logger.info(f"[{request_id}] Starting prediction")
87     start_time = time.time()
88
89     # prediction = model.predict(features) # Placeholder
90     prediction = np.random.rand(features.shape[0]) # Mock
91     prediction
92
93     inference_time = time.time() - start_time
94
95     logger.info(
96         f"[{request_id}] Prediction completed in "
97         f"{inference_time:.3f}s"
98     )
99     logger.debug(
100        f"[{request_id}] Prediction shape: {prediction.shape}"
101    )
102     logger.debug(
103        f"[{request_id}] Sample predictions: {prediction[:3]}"
104    )
105
106     # Log slow predictions
107     if inference_time > 1.0:
108         logger.warning(
109             f"[{request_id}] Slow prediction: {inference_time:.3
110             f}s"
111         )
```

```
109
110     response = {
111         'request_id': request_id,
112         'prediction': prediction.tolist(),
113         'inference_time': inference_time,
114         'status': 'success'
115     }
116
117     logger.info(f"[{request_id}] Request completed successfully")
118 )
119     return jsonify(response), 200
120
121 except ValueError as e:
122     logger.error(f"[{request_id}] Value error: {str(e)}")
123     return jsonify({
124         'error': 'Invalid input values',
125         'details': str(e)
126     }), 400
127
128 except Exception as e:
129     logger.exception(f"[{request_id}] Prediction failed")
130     return jsonify({
131         'error': 'Internal server error',
132         'request_id': request_id
133     }), 500
134
135 @app.route('/health', methods=['GET'])
136 def health():
137     """Health check endpoint"""
138     logger.debug("Health check requested")
139     return jsonify({
140         'status': 'healthy',
141         'timestamp': time.time()
142     }), 200
143
144 @app.route('/metrics', methods=['GET'])
145 def metrics():
146     """Metrics endpoint"""
147     logger.debug("Metrics requested")
148     # Return application metrics
149     return jsonify({
150         'requests_total': 0,    # Placeholder
151         'errors_total': 0,      # Placeholder
152         'avg_inference_time': 0.0  # Placeholder
153     }), 200
154
155 if __name__ == '__main__':
156     logger.info("Starting Flask application")
157     logger.info("Server configuration:")
158     logger.info("  Host: 0.0.0.0")
159     logger.info("  Port: 5000")
160     logger.info("  Debug: False")
161
162     app.run(host='0.0.0.0', port=5000, debug=False)
```

## 9 Advanced Logging Topics

### 9.1 Logger Hierarchy

Loggers follow a hierarchical naming structure using dots (.), similar to Python's module structure.

```

1 import logging
2
3 # Parent logger
4 parent_logger = logging.getLogger('myapp')
5
6 # Child loggers (automatically inherit from parent)
7 data_logger = logging.getLogger('myapp.data')
8 model_logger = logging.getLogger('myapp.model')
9 api_logger = logging.getLogger('myapp.api')
10
11 # Grandchild logger
12 preprocessing_logger = logging.getLogger('myapp.data.preprocessing')

```

#### Propagation Rules:

- Child loggers inherit settings from parent loggers
- Messages propagate up the hierarchy by default
- Set `propagate=False` to prevent propagation
- Root logger is at the top of all hierarchies

```

1 # Configure parent logger
2 parent = logging.getLogger('myapp')
3 parent.setLevel(logging.INFO)
4 handler = logging.StreamHandler()
5 parent.addHandler(handler)
6
7 # Child automatically inherits configuration
8 child = logging.getLogger('myapp.module')
9 child.info("This uses parent's handler") # Works!
10
11 # Disable propagation
12 child.propagate = False # Now won't use parent's handlers

```

### 9.2 Custom Log Levels

You can define custom logging levels for specialized needs:

```

1 import logging
2
3 # Define custom level between DEBUG and INFO
4 TRACE = 5
5 logging.addLevelName(TRACE, "TRACE")
6
7 def trace(self, message, *args, **kwargs):
8     """Add trace method to logger"""
9     if self.isEnabledFor(TRACE):
10         self._log(TRACE, message, args, **kwargs)
11

```

```

12 # Add method to Logger class
13 logging.Logger.trace = trace
14
15 # Usage
16 logger = logging.getLogger(__name__)
17 logger.setLevel(TRACE)
18
19 handler = logging.StreamHandler()
20 handler.setLevel(TRACE)
21 logger.addHandler(handler)
22
23 logger.trace("This is a trace message") # Works!
24 logger.debug("This is a debug message")

```

### 9.3 Filters

Filters provide fine-grained control over which log records are processed:

```

1 import logging
2
3 class LevelRangeFilter(logging.Filter):
4     """Filter to only allow specific level range"""
5
6     def __init__(self, min_level, max_level):
7         super().__init__()
8         self.min_level = min_level
9         self.max_level = max_level
10
11    def filter(self, record):
12        return self.min_level <= record.levelno <= self.max_level
13
14 # Usage
15 logger = logging.getLogger(__name__)
16 handler = logging.FileHandler('info_warnings.log')
17
18 # Only log INFO and WARNING (not DEBUG, ERROR, CRITICAL)
19 level_filter = LevelRangeFilter(logging.INFO, logging.WARNING)
20 handler.addFilter(level_filter)
21
22 logger.addHandler(handler)
23
24 logger.debug("Not logged")           # Below min
25 logger.info("Logged")               # In range
26 logger.warning("Logged")            # In range
27 logger.error("Not logged")          # Above max

```

#### 9.3.1 Context Filter Example

```

1 class ContextFilter(logging.Filter):
2     """Add contextual information to logs"""
3
4     def __init__(self, user_id=None):
5         super().__init__()
6         self.user_id = user_id
7
8     def filter(self, record):

```

```

9         record.user_id = self.user_id or 'anonymous'
10        return True
11
12 # Setup
13 logger = logging.getLogger(__name__)
14 handler = logging.StreamHandler()
15 formatter = logging.Formatter(
16     '%(asctime)s - [User:%(user_id)s] - %(levelname)s - %(message)s',
17 )
18 handler.setFormatter(formatter)
19
20 # Add filter
21 context_filter = ContextFilter(user_id='user123')
22 handler.addFilter(context_filter)
23
24 logger.addHandler(handler)
25
26 logger.info("Processing request")
27 # Output: 2024-12-17 14:30:45 - [User:user123] - INFO - Processing
    request

```

## 9.4 Logging in Multiprocessing

When using multiprocessing, special care is needed to avoid conflicts:

```

1 import logging
2 from logging.handlers import QueueHandler, QueueListener
3 from multiprocessing import Queue, Process
4 import time
5
6 def worker_process(queue, worker_id):
7     """Worker process with queue-based logging"""
8     # Configure worker to use queue
9     qh = QueueHandler(queue)
10    logger = logging.getLogger()
11    logger.addHandler(qh)
12    logger.setLevel(logging.INFO)
13
14    # Do work with logging
15    logger.info(f"Worker {worker_id} started")
16    time.sleep(1)
17    logger.info(f"Worker {worker_id} processing")
18    time.sleep(1)
19    logger.info(f"Worker {worker_id} completed")
20
21 if __name__ == '__main__':
22     # Create queue for log records
23     log_queue = Queue()
24
25     # Setup handlers that will process records
26     console_handler = logging.StreamHandler()
27     file_handler = logging.FileHandler('multiprocess.log')
28
29     formatter = logging.Formatter(
30         '%(asctime)s - [PID:%(process)d] - %(levelname)s - %(message)s',
31     )
32     console_handler.setFormatter(formatter)

```

```

33     file_handler.setFormatter(formatter)
34
35     # Create listener to process queue
36     listener = QueueListener(
37         log_queue,
38         console_handler,
39         file_handler,
40         respect_handler_level=True
41     )
42     listener.start()
43
44     # Create and start worker processes
45     processes = []
46     for i in range(4):
47         p = Process(target=worker_process, args=(log_queue, i))
48         p.start()
49         processes.append(p)
50
51     # Wait for all workers
52     for p in processes:
53         p.join()
54
55     # Stop listener
56     listener.stop()
57
58     print("All workers completed")

```

## 9.5 Structured Logging with JSON

For better log parsing and analysis, use JSON format:

```

1 import logging
2 import json
3 from datetime import datetime
4
5 class JsonFormatter(logging.Formatter):
6     """Format logs as JSON"""
7
8     def format(self, record):
9         log_data = {
10             'timestamp': datetime.utcnow().isoformat(),
11             'level': record.levelname,
12             'logger': record.name,
13             'message': record.getMessage(),
14             'module': record.module,
15             'function': record.funcName,
16             'line': record.lineno
17         }
18
19         # Add exception info if present
20         if record.exc_info:
21             log_data['exception'] = self.formatException(record.
22             exc_info)
23
24         # Add extra fields
25         if hasattr(record, 'user_id'):
26             log_data['user_id'] = record.user_id
27             if hasattr(record, 'request_id'):

```

```
27         log_data['request_id'] = record.request_id
28
29     return json.dumps(log_data)
30
31 # Usage
32 logger = logging.getLogger(__name__)
33 handler = logging.FileHandler('app.json.log')
34 handler.setFormatter(JsonFormatter())
35 logger.addHandler(handler)
36
37 logger.info("User logged in", extra={'user_id': 'user123'})
38 # Output: {"timestamp": "2024-12-17T14:30:45.123456",
39 #           "level": "INFO", "logger": "__main__",
40 #           "message": "User logged in", "user_id": "user123", ...}
```

## 10 Common Issues and Solutions

### 10.1 Duplicate Log Messages

**Problem:** Log messages appear multiple times in output.

**Cause:** Multiple handlers attached to logger or propagation causing duplicates.

**Solutions:**

```

1 # Solution 1: Check if handlers already exist
2 logger = logging.getLogger(__name__)
3 if not logger.handlers:
4     handler = logging.StreamHandler()
5     logger.addHandler(handler)
6
7 # Solution 2: Clear existing handlers
8 logger.handlers = []
9 logger.addHandler(new_handler)
10
11 # Solution 3: Disable propagation
12 logger.propagate = False
13
14 # Solution 4: Use hasHandlers() to check
15 if not logger.hasHandlers():
16     logger.addHandler(handler)
```

### 10.2 Logs Not Appearing

**Problem:** Log messages don't appear in expected output.

**Causes and Solutions:**

#### 1. Logger level too high

```

1 # Check and set logger level
2 logger = logging.getLogger(__name__)
3 print(f"Current level: {logger.level}")
4 logger.setLevel(logging.DEBUG)
5
```

#### 2. Handler level too high

```

1 # Set handler level appropriately
2 for handler in logger.handlers:
3     print(f"Handler level: {handler.level}")
4     handler.setLevel(logging.DEBUG)
5
```

#### 3. No handlers attached

```

1 # Verify handlers exist
2 print(f"Handlers: {logger.handlers}")
3 if not logger.handlers:
4     logger.addHandler(logging.StreamHandler())
5
```

#### 4. Root logger not configured

```

1 # Configure root logger
2 logging.basicConfig(level=logging.DEBUG)
3

```

### 10.3 basicConfig() Not Working

**Problem:** basicConfig() appears to have no effect.

**Cause:** basicConfig() only works once and only if root logger has no handlers.

**Solutions:**

```

1 # Solution 1: Use force=True (Python 3.8+)
2 logging.basicConfig(
3     level=logging.DEBUG,
4     format='%(asctime)s - %(levelname)s - %(message)s',
5     force=True    # Reconfigure even if already configured
6 )
7
8 # Solution 2: Manually reset root logger
9 root = logging.getLogger()
10 for handler in root.handlers[:]:
11     root.removeHandler(handler)
12 logging.basicConfig(level=logging.DEBUG)
13
14 # Solution 3: Check if already configured
15 if not logging.getLogger().hasHandlers():
16     logging.basicConfig(level=logging.DEBUG)

```

### 10.4 File Handler Not Writing

**Problem:** Log file is not created or not being written to.

**Solutions:**

#### 1. Check file permissions and path

```

1 import os
2
3 log_file = 'app.log'
4 log_dir = os.path.dirname(log_file) or '.'
5
6 # Check if directory exists
7 if not os.path.exists(log_dir):
8     os.makedirs(log_dir)
9
10 # Check write permissions
11 if not os.access(log_dir, os.W_OK):
12     print(f"No write permission for {log_dir}")
13

```

#### 2. Ensure proper levels set

```

1 logger.setLevel(logging.DEBUG)
2 file_handler.setLevel(logging.DEBUG)
3

```

### 3. Force flush after logging

```
1 for handler in logger.handlers:  
2     handler.flush()  
3
```

## 10.5 Performance Issues

**Problem:** Logging causes performance degradation.

**Solutions:**

### 1. Use lazy formatting

```
1 # Good - only formats if logged  
2 logger.debug("Value: %s", expensive_computation())  
3  
4 # Bad - always computes and formats  
5 logger.debug(f"Value: {expensive_computation()}")  
6
```

### 2. Increase logging level in production

```
1 # Development  
2 logger.setLevel(logging.DEBUG)  
3  
4 # Production  
5 logger.setLevel(logging.INFO)  
6
```

### 3. Use QueueHandler for async logging

```
1 from logging.handlers import QueueHandler  
2 import queue  
3  
4 log_queue = queue.Queue()  
5 queue_handler = QueueHandler(log_queue)  
6 logger.addHandler(queue_handler)  
7
```

## 11 Quick Reference Guide

### 11.1 Essential Commands Cheat Sheet

#### Basic Setup

```
1 import logging
2
3 # Quick setup
4 logging.basicConfig(level=logging.INFO)
5
6 # Create logger
7 logger = logging.getLogger(__name__)
8
9 # Set level
10 logger.setLevel(logging.DEBUG)
```

#### Logging Messages

```
1 # Five standard levels
2 logger.debug("Detailed diagnostic information")
3 logger.info("Informational message")
4 logger.warning("Warning message")
5 logger.error("Error message")
6 logger.critical("Critical error message")
7
8 # Exception logging (includes traceback)
9 try:
10     risky_operation()
11 except Exception:
12     logger.exception("Operation failed")
```

## Handlers and Formatters

```
1 # Create handlers
2 console_handler = logging.StreamHandler()
3 file_handler = logging.FileHandler('app.log')
4
5 # Rotating file handler
6 from logging.handlers import RotatingFileHandler
7 rotating_handler = RotatingFileHandler(
8     'app.log',
9     maxBytes=10*1024*1024,    # 10MB
10    backupCount=5
11 )
12
13 # Create formatter
14 formatter = logging.Formatter(
15     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
16 )
17
18 # Attach formatter to handler
19 handler.setFormatter(formatter)
20
21 # Add handler to logger
22 logger.addHandler(handler)
```

## Complete Setup Example

```
1 import logging
2
3 logger = logging.getLogger(__name__)
4 logger.setLevel(logging.DEBUG)
5
6 # Console handler
7 console = logging.StreamHandler()
8 console.setLevel(logging.INFO)
9 console.setFormatter(logging.Formatter('%(levelname)s - %(message)s'))
10
11 # File handler
12 file_h = logging.FileHandler('app.log')
13 file_h.setLevel(logging.DEBUG)
14 file_h.setFormatter(logging.Formatter(
15     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
16 ))
17
18 logger.addHandler(console)
19 logger.addHandler(file_h)
20
21 logger.info("Logger configured")
```

## 11.2 Logging Levels Quick Reference

Level	Numeric	Method	When to Use
DEBUG	10	<code>logger.debug()</code>	Detailed diagnostic info
INFO	20	<code>logger.info()</code>	General informational messages
WARNING	30	<code>logger.warning()</code>	Warning messages
ERROR	40	<code>logger.error()</code>	Error messages
CRITICAL	50	<code>logger.critical()</code>	Critical errors

## 11.3 Format String Placeholders

Placeholder	Description
<code>%(name)s</code>	Logger name
<code>%(levelname)s</code>	Log level name
<code>%(message)s</code>	Log message
<code>%(asctime)s</code>	Timestamp
<code>%(filename)s</code>	Source filename
<code>%(lineno)d</code>	Line number
<code>%(funcName)s</code>	Function name
<code>%(process)d</code>	Process ID
<code>%(thread)d</code>	Thread ID

## 11.4 Best Practices Checklist

- Use `logging.getLogger(__name__)` for logger names
- Configure logging once at application startup
- Use appropriate log levels (DEBUG in dev, INFO+ in prod)
- Use lazy formatting: `logger.info("User %s", user)`
- Use `logger.exception()` in except blocks
- Never log sensitive information (passwords, keys)
- Implement log rotation for production systems
- Add timestamps to all log messages
- Include context (request ID, user ID) when relevant
- Test logging configuration before deploying
- Monitor log file sizes regularly
- Use structured logging (JSON) for complex systems
- Set up centralized logging for distributed systems
- Document logging conventions for your team

## 12 Glossary

---

<b>Logger</b>	The main interface that applications use to log messages. Loggers are organized hierarchically by name.
<b>Handler</b>	Object responsible for dispatching log records to specific destinations (console, file, network, etc.).
<b>Formatter</b>	Specifies the layout of log messages, including which attributes to include and how to format them.
<b>Filter</b>	Provides fine-grained control over which log records are processed, beyond simple level filtering.
<b>Log Level</b>	Severity indicator for log messages (DEBUG=10, INFO=20, WARNING=30, ERROR=40, CRITICAL=50).
<b>Log Record</b>	Object containing all information about a single logging event (message, level, timestamp, location, etc.).
<b>Propagation</b>	Process by which log messages are passed from child loggers to parent loggers in the hierarchy.
<b>Root Logger</b>	The top-level logger in the hierarchy ( <code>logging.getLogger()</code> with no name), parent to all other loggers.
<b>StreamHandler</b>	Handler that sends log output to streams like <code>sys.stdout</code> or <code>sys.stderr</code> (console output).
<b>FileHandler</b>	Handler that writes log messages to a file on disk.
<b>RotatingFileHandler</b>	Handler that automatically rotates log files when they reach a specified size.
<b>TimedRotatingFileHandler</b>	Handler that rotates log files at specified time intervals (hourly, daily, etc.).
<b>basicConfig()</b>	Convenience function for simple logging configuration, works only once per program.
<b>Lazy Formatting</b>	Technique where string formatting is delayed until the message is actually logged, improving performance.
<b>Structured Logging</b>	Practice of logging data in a structured format (e.g., JSON) rather than plain text, making logs easier to parse and analyze.
<b>QueueHandler</b>	Handler that sends log records to a queue for asynchronous processing, useful in multiprocessing environments.
<b>QueueListener</b>	Receives log records from a queue and dispatches them to configured handlers.
<b>Logger Hierarchy</b>	Tree structure of loggers where child loggers inherit configuration from parent loggers.

**Exception Logging** Special logging that includes full traceback information, typically done with `logger.exception()`.

**Handler Level** Minimum severity level a handler will output, independent of the logger's level.

**Logger Level** Minimum severity level a logger will process before passing to handlers.

#### **Context Information**

Additional data attached to log records (user ID, request ID, session info) for better traceability.

## 13 Conclusion

---

Effective logging is a cornerstone of professional software development and absolutely essential for Machine Learning Operations (MLOps). The Python logging module provides a robust, flexible, and production-ready framework for tracking application behavior, debugging issues, and monitoring systems in production environments.

### 13.1 Key Takeaways

#### 1. Never Use `print()` in Production

Always use the logging module instead of `print()` statements. Logging provides:

- Severity levels for filtering
- Flexible output destinations
- Automatic timestamps and context
- Production-ready features

#### 2. Choose Appropriate Log Levels

Use the right level for each situation:

- DEBUG for development and detailed diagnostics
- INFO for production tracking and milestones
- WARNING for potential issues that don't stop execution
- ERROR for failures that affect functionality
- CRITICAL for system-level failures

#### 3. Configure Logging Properly

Set up logging at application startup with:

- Appropriate handlers (console, file, rotating)
- Clear, informative formatters
- Correct levels for each environment
- Log rotation to manage disk space

#### 4. Log at the Right Verbosity

Balance information and noise:

- Too much logging creates noise and performance issues
- Too little logging misses important events
- Use DEBUG liberally in development
- Be selective with INFO/WARNING in production

#### 5. Include Relevant Context

Make logs actionable by including:

- Request IDs for tracing requests
- User IDs for user-specific issues
- Timestamps for temporal analysis
- Function names and line numbers for debugging

## 6. Handle Errors Gracefully

Always use proper exception logging:

- Use `logger.exception()` in except blocks
- Include full tracebacks for debugging
- Log before re-raising exceptions
- Don't swallow exceptions silently

## 7. Monitor and Rotate Logs

In production systems:

- Implement log rotation to prevent disk full
- Monitor log file sizes regularly
- Archive old logs appropriately
- Set up alerts for critical errors

### 13.2 MLOps-Specific Recommendations

For machine learning operations, logging serves critical functions:

- **Model Training:** Track epochs, loss, metrics, and training time
- **Data Pipelines:** Log data loading, cleaning, and transformation steps
- **Model Serving:** Record prediction requests, responses, and inference times
- **Performance Monitoring:** Track resource usage (CPU, GPU, memory)
- **Error Tracking:** Identify and debug production failures quickly
- **Audit Trails:** Maintain compliance and reproducibility records
- **Experiment Tracking:** Document configurations and results

### 13.3 Implementation Strategy

When implementing logging in your projects:

1. **Start Simple:** Begin with `basicConfig()` for prototypes
2. **Expand Gradually:** Add handlers and formatters as needs grow
3. **Establish Standards:** Define logging conventions for your team
4. **Review Regularly:** Analyze logs to improve logging strategy
5. **Integrate Tools:** Connect logging with monitoring and alerting systems
6. **Test Configuration:** Verify logging works before deploying
7. **Document Practices:** Maintain clear documentation for team members

### 13.4 Common Pitfalls to Avoid

**Warning****Avoid These Common Mistakes:**

- Using `print()` instead of logging in production
- Logging sensitive information (passwords, API keys)
- Over-logging in tight loops (performance impact)
- Not rotating log files (disk space issues)
- Using string concatenation instead of lazy formatting
- Forgetting to set appropriate log levels
- Not including enough context in log messages
- Swallowing exceptions without logging
- Configuring logging multiple times
- Not testing logging configuration

### 13.5 Next Steps

To continue improving your logging practices:

1. **Practice:** Implement logging in personal projects
2. **Experiment:** Try different handlers and formatters
3. **Integrate:** Connect logging with monitoring tools (ELK, Splunk, Datadog)
4. **Learn Advanced Topics:** Explore async logging, custom handlers, filters
5. **Study Production Systems:** Analyze logging in open-source projects
6. **Optimize:** Profile and optimize logging performance
7. **Automate:** Set up automated log analysis and alerting

### 13.6 Integration with Other Tools

Python logging integrates well with:

- **MLflow:** Experiment tracking and model registry
- **TensorBoard:** Visualization and monitoring
- **ELK Stack:** Elasticsearch, Logstash, Kibana for log analysis
- **Splunk:** Enterprise log management
- **Datadog:** Cloud monitoring and analytics
- **Sentry:** Error tracking and monitoring
- **CloudWatch:** AWS log management
- **Prometheus/Grafana:** Metrics and dashboards

## 13.7 Final Recommendations

### Best Practices Summary

1. Use `logging.getLogger(__name__)` for all loggers
2. Configure logging once at application startup
3. Use appropriate levels: DEBUG < INFO < WARNING < ERROR < CRITICAL
4. Implement log rotation in production
5. Include context information (timestamps, IDs, locations)
6. Use lazy formatting for performance
7. Never log sensitive information
8. Use `logger.exception()` in except blocks
9. Test logging configuration thoroughly
10. Monitor and analyze logs regularly

## 13.8 Additional Resources

For further learning and reference:

- **Official Documentation:**

- Python Logging Module: <https://docs.python.org/3/library/logging.html>
- Logging Cookbook: <https://docs.python.org/3/howto/logging-cookbook.html>
- Logging HOWTO: <https://docs.python.org/3/howto/logging.html>

- **Tutorials and Guides:**

- Real Python Logging Guide: <https://realpython.com/python-logging/>
- Python Logging Best Practices: <https://www.loggly.com/ultimate-guide/python-logging-basics>

- **MLOps Integration:**

- MLflow Documentation: <https://www.mlflow.org/docs/latest/tracking.html>
- TensorBoard Logging: <https://www.tensorflow.org/tensorboard>

- **Log Management Tools:**

- ELK Stack: <https://www.elastic.co/elasticsearch>
- Splunk: <https://www.splunk.com>
- Datadog: <https://www.datadoghq.com>
- Sentry: <https://sentry.io/>

- **Books:**

- “Logging and Log Management” by Anton Chuvakin
- “Python Logging Essentials” by Chetan Giridhar

**Remember**

**Logging is not just about recording events—it's about making your application observable, debuggable, and maintainable.**

Good logging practices will:

- Save hours of debugging time
- Enable quick problem resolution in production
- Provide insights into system behavior
- Build confidence in your deployments
- Support compliance and auditing requirements
- Facilitate team collaboration

Invest time in proper logging setup—it pays dividends throughout the application lifecycle!

---

*End of Python Logging Module Complete Reference Guide*

*"The most effective debugging tool is still careful thought,  
coupled with judiciously placed print statements."*

— Brian Kernighan