

# MLOps Project

## with DVC, Git & MLflow

Complete Integration Guide

A Comprehensive Guide to Building Production MLOps Pipelines with DVC Version Control, Git Integration, and MLflow Experiment Tracking via DagShub

Sujil S

sujil9480@gmail.com

December 25, 2025

## Contents

---

<b>1</b>	<b>Introduction to MLOps with DagShub</b>	<b>2</b>
1.1	What is This Project About? . . . . .	2
1.2	The MLOps Stack . . . . .	2
1.3	What Makes This Different? . . . . .	2
1.4	Key Differences from Previous Workflow . . . . .	3
<b>2</b>	<b>Initial Project Setup</b>	<b>4</b>
2.1	Step 1: Create Python Environment . . . . .	4
2.2	Step 2: Activate Environment . . . . .	4
2.3	Step 3: Create Source Directory . . . . .	4
2.4	Step 4: Create Requirements File . . . . .	4
2.5	Step 5: Install Requirements . . . . .	4
<b>3</b>	<b>Building Pipeline Components</b>	<b>5</b>
3.1	Component Overview . . . . .	5
3.2	Data Ingestion Module . . . . .	5
3.3	Data Pre-Processing Module . . . . .	7
3.4	Feature Engineering Module . . . . .	10
3.5	Model Building Module . . . . .	13
<b>4</b>	<b>MLflow Integration with DagShub</b>	<b>16</b>
4.1	What is DagShub? . . . . .	16
4.2	Connecting to DagShub . . . . .	16
4.2.1	Step 1: Create DagShub Account . . . . .	16
4.2.2	Step 2: Obtain Tracking URI . . . . .	16
4.3	Model Evaluation with MLflow . . . . .	16
4.4	Understanding the MLflow Integration . . . . .	20
4.4.1	Key Components . . . . .	20
<b>5</b>	<b>Configuration Files</b>	<b>21</b>
5.1	Step 7: Create .gitignore . . . . .	21
5.2	Step 8: Initialize DVC and Git . . . . .	21
5.3	Step 9: Create DVC Pipeline . . . . .	22
5.4	Step 10: Create Parameters File . . . . .	23
<b>6</b>	<b>Initial Git Commit and DVC Setup</b>	<b>24</b>
6.1	Step 11: Check Git Status . . . . .	24
6.2	Step 12: Check DVC Status . . . . .	24
6.3	Step 13: Make Initial Git Commit . . . . .	25
6.4	Step 14: Configure DVC Remote Storage . . . . .	25
6.4.1	Create Local Storage Directory . . . . .	25
6.4.2	Add as DVC Remote . . . . .	25
6.4.3	Check DVC Configuration . . . . .	26
6.4.4	Add S3 to .gitignore . . . . .	26
<b>7</b>	<b>Running the Pipeline and First Experiment</b>	<b>27</b>
7.1	Step 15: Execute the Pipeline . . . . .	27
7.2	Pipeline Execution Output . . . . .	27
7.3	Understanding the Output . . . . .	28
7.4	Step 16: Understanding DVC Cache . . . . .	29

<b>8 Running Multiple Experiments</b>	<b>30</b>
8.1 Step 17: Experimentation Phase . . . . .	30
8.2 Experiment Workflow . . . . .	30
8.3 Running Experiment Variations . . . . .	30
8.3.1 Experiment 1: Initial Configuration . . . . .	30
8.3.2 Experiment 2: Fewer Trees . . . . .	30
8.3.3 Experiment 3: More Features . . . . .	31
8.3.4 Experiment 4: Different Split . . . . .	31
8.3.5 Experiment 5: Balanced Configuration . . . . .	31
<b>9 Step 18: Analyzing Experiments in MLflow</b>	<b>33</b>
9.1 Accessing the MLflow UI . . . . .	33
9.2 Experiment Results Summary . . . . .	33
9.3 Detailed Run Analysis . . . . .	33
9.3.1 Run 1: able-ray-853 . . . . .	33
9.3.2 Run 2: handsome-asp-382 . . . . .	33
9.3.3 Run 5: receptive-crow-130 . . . . .	34
<b>10 Step 19: Deciding What "Best" Means</b>	<b>35</b>
10.1 Defining Success Criteria . . . . .	35
10.2 Analysis for Spam Classification . . . . .	35
10.3 Ranking the Runs . . . . .	35
<b>11 Step 20: Re-creating the Winning Run</b>	<b>36</b>
11.1 Why Re-create? . . . . .	36
11.2 Obtaining Winning Parameters . . . . .	36
11.3 Step 21: Update params.yaml . . . . .	36
11.4 Step 22: Run Pipeline with Winning Parameters . . . . .	36
11.5 Step 23: Verify Metrics (Sanity Check) . . . . .	37
<b>12 Step 24: The Critical Commit</b>	<b>39</b>
12.1 Why This Commit Matters . . . . .	39
12.2 Check Git Status . . . . .	39
12.3 Stage and Commit . . . . .	40
12.4 What This Commit Means . . . . .	40
<b>13 Step 25: Push Data and Model to DVC Remote</b>	<b>41</b>
13.1 Why Push to DVC Remote? . . . . .	41
13.2 Execute DVC Push . . . . .	41
13.3 What Gets Pushed . . . . .	41
13.4 Verification . . . . .	41
<b>14 Step 26: Linking Git Commits to MLflow Runs</b>	<b>43</b>
14.1 Why Tag MLflow Runs? . . . . .	43
14.2 Create Tagging Script . . . . .	43
14.3 Understanding the Script . . . . .	44
14.3.1 Key Components . . . . .	44
14.4 Finding the Run ID . . . . .	44
14.5 Step 27: Execute Tagging Script . . . . .	44
14.6 Step 28: Verify Tags in MLflow UI . . . . .	45

<b>15 Step 29: Multi-Model Experimentation</b>	<b>46</b>
15.1 Scenario: Team with Multiple Models . . . . .	46
15.2 Using Different Experiment Names . . . . .	46
15.2.1 RandomForest Experiments . . . . .	46
15.2.2 XGBoost Experiments . . . . .	46
15.2.3 Neural Network Experiments . . . . .	46
15.3 Benefits of Separate Experiments . . . . .	46
15.4 Viewing All Experiments . . . . .	46
<b>16 Reproducing the Project on a New Machine</b>	<b>48</b>
16.1 The Reproducibility Challenge . . . . .	48
16.2 Important Note About Local DVC Remote . . . . .	48
16.3 Reproduction Workflow . . . . .	49
16.4 Step-by-Step Reproduction . . . . .	49
16.4.1 Step 1: Clone Repository . . . . .	49
16.4.2 Step 2: Create and Activate Environment . . . . .	50
16.4.3 Step 3: Install Requirements . . . . .	50
16.4.4 Step 4: Configure DVC Remote Access . . . . .	51
16.4.5 Step 5: Pull Data and Models . . . . .	52
16.4.6 Step 6: Verify Pipeline Status . . . . .	52
16.4.7 Step 7 (Optional): Re-run Pipeline . . . . .	53
16.5 Verification Checklist . . . . .	53
16.6 Manual Verification . . . . .	54
16.6.1 Verify Metrics . . . . .	54
16.6.2 Verify Parameters . . . . .	54
16.6.3 Test Model Loading . . . . .	54
16.7 Why This Matters in MLOps . . . . .	55
16.8 Troubleshooting Reproduction Issues . . . . .	55
16.8.1 Issue 1: "Remote not found" . . . . .	55
16.8.2 Issue 2: "Permission denied" . . . . .	56
16.8.3 Issue 3: "Checksum mismatch" . . . . .	56
16.8.4 Issue 4: "Python package conflicts" . . . . .	56
<b>17 Complete Mental Model</b>	<b>57</b>
17.1 The Three-System Architecture . . . . .	57
17.2 Understanding Each Component . . . . .	57
17.2.1 Git: Code and Metadata . . . . .	57
17.2.2 DVC: Data and Models . . . . .	58
17.2.3 MLflow: Experiment Tracking . . . . .	58
17.3 Data Flow Summary . . . . .	59
17.4 Complete Workflow Visualization . . . . .	59
17.5 Key Concepts Review . . . . .	60
<b>18 MLOps Best Practices</b>	<b>61</b>
18.1 Experiment Management . . . . .	61
18.1.1 Naming Conventions . . . . .	61
18.1.2 Commit Message Best Practices . . . . .	61
18.1.3 Parameter Documentation . . . . .	61
18.2 DVC Remote Management . . . . .	61
18.2.1 Production Setup Recommendations . . . . .	62
18.2.2 Migrating to AWS S3 . . . . .	62
18.3 MLflow Best Practices . . . . .	63

18.3.1	What to Log . . . . .	63
18.3.2	Comprehensive Tagging Strategy . . . . .	63
18.4	Workflow Best Practices . . . . .	64
18.4.1	Before Starting Experiments . . . . .	64
18.4.2	During Experimentation . . . . .	64
18.4.3	After Choosing Best Model . . . . .	65
18.5	Team Collaboration Guidelines . . . . .	65
18.5.1	Communication Best Practices . . . . .	65
18.5.2	Avoiding Conflicts . . . . .	66
18.6	Code Quality Best Practices . . . . .	66
18.6.1	Logging Standards . . . . .	66
18.6.2	Error Handling . . . . .	66
18.6.3	Type Hints and Documentation . . . . .	67
<b>19</b>	<b>Comprehensive Troubleshooting Guide</b>	<b>68</b>
19.1	DVC Issues . . . . .	68
19.1.1	Issue 1: Failed to Push Data . . . . .	68
19.1.2	Issue 2: DVC Pull Fails . . . . .	68
19.1.3	Issue 3: Pipeline Won't Re-run . . . . .	69
19.1.4	Issue 4: Checksum Mismatch . . . . .	69
19.2	MLflow Issues . . . . .	70
19.2.1	Issue 5: MLflow Not Logging Runs . . . . .	70
19.2.2	Issue 6: Cannot Find Run ID . . . . .	71
19.3	Git Issues . . . . .	71
19.3.1	Issue 7: dvc.lock Merge Conflicts . . . . .	71
19.3.2	Issue 8: Large Files Committed to Git . . . . .	72
19.4	Environment Issues . . . . .	72
19.4.1	Issue 9: Import Errors . . . . .	72
19.4.2	Issue 10: NLTK Data Not Found . . . . .	73
19.5	Performance Issues . . . . .	73
19.5.1	Issue 11: Slow DVC Operations . . . . .	73
19.5.2	Issue 12: Large Cache Size . . . . .	73
<b>20</b>	<b>Quick Reference Guide</b>	<b>75</b>
20.1	Essential Commands . . . . .	75
20.2	Complete Workflow Summary . . . . .	76
20.3	Project Structure Reference . . . . .	77
20.4	Key Files Explained . . . . .	78
20.5	Common Command Combinations . . . . .	79
<b>21</b>	<b>Comparison: DVC-Only vs DVC+MLflow+Git Workflow</b>	<b>80</b>
21.1	Architecture Comparison . . . . .	80
21.2	Feature-by-Feature Comparison . . . . .	80
21.3	Workflow Process Comparison . . . . .	80
21.3.1	Experimentation Phase . . . . .	80
21.3.2	Model Freezing Phase . . . . .	81
21.4	Team Collaboration Comparison . . . . .	81
21.5	Use Case Recommendations . . . . .	81
21.5.1	When to Use DVC-Only . . . . .	81
21.5.2	When to Use DVC + MLflow + Git . . . . .	81
21.6	Real-World Scenario Comparison . . . . .	82
21.6.1	Scenario: Finding Best Model from 20+ Experiments . . . . .	82

21.6.2 Scenario: Reproducing 6-Month-Old Model . . . . .	83
21.7 Cost-Benefit Analysis . . . . .	83
21.8 Migration Path . . . . .	83
21.8.1 Upgrading from DVC-Only to DVC+MLflow . . . . .	84
<b>22 Conclusion and Next Steps</b> . . . . .	<b>85</b>
22.1 What You've Accomplished . . . . .	85
22.2 Key Takeaways . . . . .	85
22.2.1 The Three-System Architecture . . . . .	85
22.2.2 Critical Best Practices . . . . .	86
22.3 Common Pitfalls to Avoid . . . . .	86
22.4 Next Steps and Extensions . . . . .	87
22.4.1 Immediate Next Steps . . . . .	87
22.4.2 Advanced Extensions . . . . .	88
22.5 Learning Resources . . . . .	89
22.5.1 Official Documentation . . . . .	89
22.5.2 Community and Support . . . . .	89
22.5.3 Advanced Topics to Explore . . . . .	89
22.6 Final Thoughts . . . . .	89
22.7 Project Checklist . . . . .	91
22.8 Get In Touch . . . . .	91

## 1 Introduction to MLOps with DagShub

### 1.1 What is This Project About?

This comprehensive guide demonstrates a complete MLOps workflow that integrates three powerful tools:

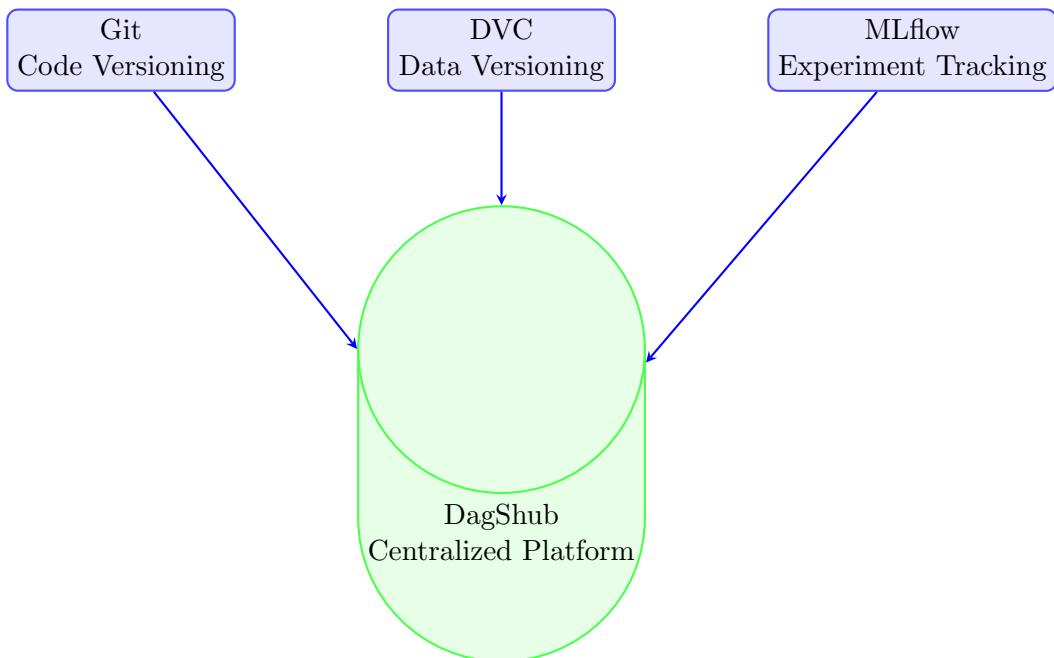
1. **DVC (Data Version Control)**: For versioning data, models, and pipeline artifacts
2. **Git**: For versioning code and metadata
3. **MLflow via DagShub**: For experiment tracking and model registry

#### Project Goals

By the end of this guide, you will:

- Build a complete ML pipeline with proper logging
- Version control data and models with DVC
- Track experiments with MLflow on DagShub
- Compare multiple model runs systematically
- Freeze and deploy the best performing model
- Enable team collaboration and reproducibility

### 1.2 The MLOps Stack



### 1.3 What Makes This Different?

Unlike the previous DVC-only workflow, this project adds:

- **Centralized Experiment Tracking**: MLflow UI hosted on DagShub

- **Team Collaboration:** Everyone sees all experiments
- **Model Comparison:** Visual comparison of multiple runs
- **Experiment Tagging:** Link Git commits to MLflow runs
- **Model Registry:** Track model lineage and deployment status

#### 1.4 Key Differences from Previous Workflow

DVC-Only Workflow	DVC + MLflow Workflow
Local experiment tracking	Centralized MLflow tracking
Manual metric comparison	Visual metric comparison
Experiments private to user	All experiments visible to team
Limited experiment history	Full experiment history preserved
No model registry	MLflow model registry

## 2 Initial Project Setup

### 2.1 Step 1: Create Python Environment

```
# Create virtual environment  
python -m venv venv
```

### 2.2 Step 2: Activate Environment

```
# Windows  
venv\Scripts\activate.bat  
  
# Linux/Mac  
source venv/bin/activate
```

### 2.3 Step 3: Create Source Directory

```
mkdir src
```

### 2.4 Step 4: Create Requirements File

#### requirements.txt

```
dvc  
numpy  
pandas  
matplotlib  
wordcloud  
nltk  
scikit-learn  
xgboost  
pyyaml  
mlflow  
dagshub
```

#### Important Note

##### New Dependencies:

- **mlflow**: Experiment tracking library
- **dagshub**: DagShub integration for MLflow
- **xgboost**: Additional model option (if needed)

### 2.5 Step 5: Install Requirements

```
pip install -r requirements.txt
```

## 3 Building Pipeline Components

### 3.1 Component Overview

The ML pipeline consists of five stages:

1. **Data Ingestion:** Load and split data
2. **Data Pre-Processing:** Clean and transform text
3. **Feature Engineering:** Apply TF-IDF
4. **Model Building:** Train RandomForest
5. **Model Evaluation:** Calculate metrics and log to MLflow

### 3.2 Data Ingestion Module

src/Data\_Ingestion.py

```
1 import pandas as pd
2 import os
3 from sklearn.model_selection import train_test_split
4 import logging
5 import yaml
6
7 # Ensure the "logs" directory exists
8 log_dir = 'logs'
9 os.makedirs(log_dir, exist_ok=True)
10
11 # Logging configuration
12 logger = logging.getLogger('Data_Ingestion')
13 logger.setLevel('DEBUG')
14
15 console_handler = logging.StreamHandler()
16 console_handler.setLevel('DEBUG')
17
18 log_file_path = os.path.join(log_dir, 'Data_Ingestion.log')
19 file_handler = logging.FileHandler(log_file_path)
20 file_handler.setLevel('DEBUG')
21
22 formatter = logging.Formatter(
23     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
24 )
25 console_handler.setFormatter(formatter)
26 file_handler.setFormatter(formatter)
27
28 logger.addHandler(console_handler)
29 logger.addHandler(file_handler)
30
31
32 def load_params(params_path: str) -> dict:
33     """Load parameters from a YAML file."""
34     try:
35         with open(params_path, 'r') as file:
36             params = yaml.safe_load(file)
37             logger.debug('Parameters retrieved from %s',
params_path)
```

```
38     return params
39 except FileNotFoundError:
40     logger.error('File not found: %s', params_path)
41     raise
42 except yaml.YAMLError as e:
43     logger.error('YAML error: %s', e)
44     raise
45 except Exception as e:
46     logger.error('Unexpected error: %s', e)
47     raise
48
49
50 def load_data(data_url: str) -> pd.DataFrame:
51     """Load data from a CSV file."""
52     try:
53         df = pd.read_csv(data_url)
54         logger.debug('Data loaded from %s', data_url)
55         return df
56     except pd.errors.ParserError as e:
57         logger.error('Failed to parse the CSV file: %s', e)
58         raise
59     except Exception as e:
60         logger.error('Unexpected error while loading data: %s', e)
61         raise
62
63
64 def preprocess_data(df: pd.DataFrame) -> pd.DataFrame:
65     """Preprocess the data."""
66     try:
67         df.drop(columns=['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],
68                 inplace=True)
69         df.rename(columns={'v1': 'target', 'v2': 'text'},
70                   inplace=True)
71         logger.debug('Data preprocessing completed')
72         return df
73     except KeyError as e:
74         logger.error('Missing column in dataframe: %s', e)
75         raise
76     except Exception as e:
77         logger.error('Unexpected error during preprocessing: %s', e)
78         raise
79
80
81 def save_data(train_data: pd.DataFrame,
82               test_data: pd.DataFrame,
83               data_path: str) -> None:
84     """Save the train and test datasets."""
85     try:
86         raw_data_path = os.path.join(data_path, 'raw')
87         os.makedirs(raw_data_path, exist_ok=True)
88
89         train_data.to_csv(
90             os.path.join(raw_data_path, "train.csv"),
```

```

1      index=False
2  )
3  test_data.to_csv(
4      os.path.join(raw_data_path, "test.csv"),
5      index=False
6  )
7  logger.debug('Train and test data saved to %s',
8          raw_data_path)
9  except Exception as e:
10     logger.error('Error while saving data: %s', e)
11     raise
12
13
14 def main():
15     try:
16         params = load_params(params_path='params.yaml')
17         test_size = params['data_ingestion']['test_size']
18
19         data_path = 'https://raw.githubusercontent.com/' \
20                     'Error-Makes-Clever/MLOPS-Data-Versioning-' \
21                     '\
22                     'using-DVC-Project-1/refs/heads/main/' \
23                     'Experiments/spam.csv'
24
25         df = load_data(data_url=data_path)
26         final_df = preprocess_data(df)
27
28         train_data, test_data = train_test_split(
29             final_df, test_size=test_size, random_state=2
30         )
31
32         save_data(train_data, test_data, data_path='./data')
33     except Exception as e:
34         logger.error('Failed to complete data ingestion: %s', e)
35     print(f"Error: {e}")
36
37
38 if __name__ == '__main__':
39     main()

```

### 3.3 Data Pre-Processing Module

#### src/Data\_Pre\_Processing.py - Part 1

```

1 import os
2 import logging
3 import pandas as pd
4 from sklearn.preprocessing import LabelEncoder
5 from nltk.stem.porter import PorterStemmer
6 from nltk.corpus import stopwords
7 import string
8 import nltk
9

```

```
10 nltk.download('stopwords')
11 nltk.download('punkt')
12
13 # Ensure the "logs" directory exists
14 log_dir = 'logs'
15 os.makedirs(log_dir, exist_ok=True)
16
17 # Setting up logger
18 logger = logging.getLogger('Data_Pre_Processing')
19 logger.setLevel('DEBUG')
20
21 console_handler = logging.StreamHandler()
22 console_handler.setLevel('DEBUG')
23
24 log_file_path = os.path.join(log_dir, 'Data_Pre_Processing.log')
25 file_handler = logging.FileHandler(log_file_path)
26 file_handler.setLevel('DEBUG')
27
28 formatter = logging.Formatter(
29     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
30 )
31 console_handler.setFormatter(formatter)
32 file_handler.setFormatter(formatter)
33
34 logger.addHandler(console_handler)
35 logger.addHandler(file_handler)
36
37
38 def transform_text(text):
39     """
40         Transform text: lowercase, tokenize, remove stopwords,
41         punctuation, and stem.
42     """
43     ps = PorterStemmer()
44
45     # Convert to lowercase
46     text = text.lower()
47
48     # Tokenize
49     text = nltk.word_tokenize(text)
50
51     # Remove non-alphanumeric tokens
52     text = [word for word in text if word.isalnum()]
53
54     # Remove stopwords and punctuation
55     text = [word for word in text
56             if word not in stopwords.words('english')
57             and word not in string.punctuation]
58
59     # Stem the words
60     text = [ps.stem(word) for word in text]
61
62     # Join back into string
63     return " ".join(text)
64
```

```
65
66 def preprocess_df(df, text_column='text',
67                     target_column='target'):
68     """
69     Preprocess DataFrame: encode target, remove duplicates,
70     transform text.
71     """
72     try:
73         logger.debug('Starting preprocessing for DataFrame')
74
75         # Encode the target column
76         encoder = LabelEncoder()
77         df[target_column] = encoder.fit_transform(df[
78             target_column])
79         logger.debug('Target column encoded')
80
81         # Remove duplicate rows
82         df = df.drop_duplicates(keep='first')
83         logger.debug('Duplicates removed')
84
85         # Apply text transformation
86         df.loc[:, text_column] = df[text_column].apply(
87             transform_text
88         )
89         logger.debug('Text column transformed')
90
91     return df
92 except KeyError as e:
93     logger.error('Column not found: %s', e)
94     raise
95 except Exception as e:
96     logger.error('Error during text normalization: %s', e)
97     raise
98
99 def main(text_column='text', target_column='target'):
100     """
101     Main function: load raw data, preprocess, save processed.
102     """
103     try:
104         # Load data from data/raw
105         train_data = pd.read_csv('./data/raw/train.csv')
106         test_data = pd.read_csv('./data/raw/test.csv')
107         logger.debug('Data loaded properly')
108
109         # Transform the data
110         train_processed = preprocess_df(train_data,
111                                         text_column,
112                                         target_column)
113         test_processed = preprocess_df(test_data,
114                                         text_column,
115                                         target_column)
116
117         # Store in data/interim
118         data_path = os.path.join("./data", "interim")
119         os.makedirs(data_path, exist_ok=True)
```

```

120     train_processed.to_csv(
121         os.path.join(data_path, "train_processed.csv"),
122         index=False
123     )
124     test_processed.to_csv(
125         os.path.join(data_path, "test_processed.csv"),
126         index=False
127     )
128
129     logger.debug('Processed data saved to %s', data_path)
130 except FileNotFoundError as e:
131     logger.error('File not found: %s', e)
132 except pd.errors.EmptyDataError as e:
133     logger.error('No data: %s', e)
134 except Exception as e:
135     logger.error('Failed to complete data transformation: %s', e)
136     print(f"Error: {e}")
137
138
139
140 if __name__ == '__main__':
141     main()

```

### 3.4 Feature Engineering Module

src/Feature\_Engineering.py

```

1 import pandas as pd
2 import os
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 import logging
5 import yaml
6
7 # Ensure the "logs" directory exists
8 log_dir = 'logs'
9 os.makedirs(log_dir, exist_ok=True)
10
11 # Logging configuration
12 logger = logging.getLogger('Feature_Engineering')
13 logger.setLevel('DEBUG')
14
15 console_handler = logging.StreamHandler()
16 console_handler.setLevel('DEBUG')
17
18 log_file_path = os.path.join(log_dir, 'Feature_Engineering.log')
19 file_handler = logging.FileHandler(log_file_path)
20 file_handler.setLevel('DEBUG')
21
22 formatter = logging.Formatter(
23     '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
24 )
25 console_handler.setFormatter(formatter)

```

```
26     file_handler.setFormatter(formatter)
27
28     logger.addHandler(console_handler)
29     logger.addHandler(file_handler)
30
31
32     def load_params(params_path: str) -> dict:
33         """Load parameters from a YAML file."""
34         try:
35             with open(params_path, 'r') as file:
36                 params = yaml.safe_load(file)
37                 logger.debug('Parameters retrieved from %s',
38 params_path)
39                 return params
40         except FileNotFoundError:
41             logger.error('File not found: %s', params_path)
42             raise
43         except yaml.YAMLError as e:
44             logger.error('YAML error: %s', e)
45             raise
46         except Exception as e:
47             logger.error('Unexpected error: %s', e)
48             raise
49
50     def load_data(file_path: str) -> pd.DataFrame:
51         """Load data from a CSV file."""
52         try:
53             df = pd.read_csv(file_path)
54             df.fillna(' ', inplace=True)
55             logger.debug('Data loaded and NaNs filled from %s',
56 file_path)
57             return df
58         except pd.errors.ParserError as e:
59             logger.error('Failed to parse the CSV file: %s', e)
60             raise
61         except Exception as e:
62             logger.error('Unexpected error while loading data: %s',
63 e)
64             raise
65
66     def apply_tfidf(train_data: pd.DataFrame,
67                     test_data: pd.DataFrame,
68                     max_features: int) -> tuple:
69         """Apply TF-IDF to the data."""
70         try:
71             vectorizer = TfidfVectorizer(max_features=max_features)
72
73             X_train = train_data['text'].values
74             y_train = train_data['target'].values
75             X_test = test_data['text'].values
76             y_test = test_data['target'].values
77
78             X_train_bow = vectorizer.fit_transform(X_train)
79             X_test_bow = vectorizer.transform(X_test)
```

```
80
81     train_df = pd.DataFrame(X_train_bow.toarray())
82     train_df['label'] = y_train
83
84     test_df = pd.DataFrame(X_test_bow.toarray())
85     test_df['label'] = y_test
86
87     logger.debug('TF-IDF applied and data transformed')
88     return train_df, test_df
89 except Exception as e:
90     logger.error('Error during TF-IDF transformation: %s',
91     e)
92     raise
93
94 def save_data(df: pd.DataFrame, file_path: str) -> None:
95     """Save the dataframe to a CSV file."""
96     try:
97         os.makedirs(os.path.dirname(file_path), exist_ok=True)
98         df.to_csv(file_path, index=False)
99         logger.debug('Data saved to %s', file_path)
100    except Exception as e:
101        logger.error('Unexpected error while saving data: %s',
102        e)
103        raise
104
105 def main():
106     try:
107         params = load_params(params_path='params.yaml')
108         max_features = params['feature_engineering'][
109             'max_features']
110
111         train_data = load_data('./data/interim/train_processed.
112         csv')
113         test_data = load_data('./data/interim/test_processed.
114         csv')
115
116         train_df, test_df = apply_tfidf(train_data, test_data,
117                                         max_features)
118
119         save_data(train_df,
120                   os.path.join("./data", "processed",
121                               "train_tfidf.csv"))
122         save_data(test_df,
123                   os.path.join("./data", "processed",
124                               "test_tfidf.csv"))
125     except Exception as e:
126         logger.error('Failed to complete feature engineering: %
127 s', e)
128         print(f"Error: {e}")
129
130
131 if __name__ == '__main__':
132     main()
```

### 3.5 Model Building Module

src/Model\_Building\_Rf.py

```
 1 import os
 2 import numpy as np
 3 import pandas as pd
 4 import pickle
 5 import logging
 6 from sklearn.ensemble import RandomForestClassifier
 7 import yaml
 8
 9 # Ensure the "logs" directory exists
10 log_dir = 'logs'
11 os.makedirs(log_dir, exist_ok=True)
12
13 # Logging configuration
14 logger = logging.getLogger('Model_Building')
15 logger.setLevel('DEBUG')
16
17 console_handler = logging.StreamHandler()
18 console_handler.setLevel('DEBUG')
19
20 log_file_path = os.path.join(log_dir, 'Model_Building.log')
21 file_handler = logging.FileHandler(log_file_path)
22 file_handler.setLevel('DEBUG')
23
24 formatter = logging.Formatter(
25     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
26 )
27 console_handler.setFormatter(formatter)
28 file_handler.setFormatter(formatter)
29
30 logger.addHandler(console_handler)
31 logger.addHandler(file_handler)
32
33
34 def load_params(params_path: str) -> dict:
35     """Load parameters from a YAML file."""
36     try:
37         with open(params_path, 'r') as file:
38             params = yaml.safe_load(file)
39             logger.debug('Parameters retrieved from %s',
40                         params_path)
41         return params
42     except FileNotFoundError:
43         logger.error('File not found: %s', params_path)
44         raise
45     except yaml.YAMLError as e:
46         logger.error('YAML error: %s', e)
47         raise
48     except Exception as e:
49         logger.error('Unexpected error: %s', e)
50         raise
51
52 def load_data(file_path: str) -> pd.DataFrame:
```

```
53     """Load data from a CSV file."""
54     try:
55         df = pd.read_csv(file_path)
56         logger.debug('Data loaded from %s with shape %s',
57                      file_path, df.shape)
58         return df
59     except pd.errors.ParserError as e:
60         logger.error('Failed to parse the CSV file: %s', e)
61         raise
62     except FileNotFoundError as e:
63         logger.error('File not found: %s', e)
64         raise
65     except Exception as e:
66         logger.error('Unexpected error while loading data: %s',
67                      e)
68         raise
69
70 def train_model(X_train: np.ndarray,
71                  y_train: np.ndarray,
72                  params: dict) -> RandomForestClassifier:
73     """Train the RandomForest model."""
74     try:
75         if X_train.shape[0] != y_train.shape[0]:
76             raise ValueError(
77                 "Number of samples in X_train and y_train must
78                 match")
79
80         logger.debug('Initializing RandomForest with params: %s
81         ',
82                     params)
83         clf = RandomForestClassifier(
84             n_estimators=params['n_estimators'],
85             random_state=params['random_state'],
86             max_depth=params['max_depth']
87         )
88
89         logger.debug('Model training started with %d samples',
90                     X_train.shape[0])
91         clf.fit(X_train, y_train)
92         logger.debug('Model training completed')
93
94         return clf
95     except ValueError as e:
96         logger.error('ValueError during model training: %s', e)
97         raise
98     except Exception as e:
99         logger.error('Error during model training: %s', e)
100        raise
101
102 def save_model(model, file_path: str) -> None:
103     """Save the trained model to a file."""
104     try:
105         os.makedirs(os.path.dirname(file_path), exist_ok=True)
```

```
106
107     with open(file_path, 'wb') as file:
108         pickle.dump(model, file)
109     logger.debug('Model saved to %s', file_path)
110 except FileNotFoundError as e:
111     logger.error('File path not found: %s', e)
112     raise
113 except Exception as e:
114     logger.error('Error while saving model: %s', e)
115     raise
116
117
118 def main():
119     try:
120         params = load_params('params.yaml')[‘model_building’]
121
122         train_data = load_data('./data/processed/train_tfidf.
123 csv')
124         X_train = train_data.iloc[:, :-1].values
125         y_train = train_data.iloc[:, -1].values
126
127         clf = train_model(X_train, y_train, params)
128
129         model_save_path = ‘models/model_rf.pkl’
130         save_model(clf, model_save_path)
131
132     except Exception as e:
133         logger.error('Failed to complete model building: %s', e
134     )
135
136 if __name__ == '__main__':
137     main()
```

## 4 MLflow Integration with DagShub

### 4.1 What is DagShub?

DagShub is a platform for data scientists and ML engineers that provides:

- **Hosted MLflow:** No need to run your own MLflow server
- **Git Integration:** Direct connection to GitHub repositories
- **DVC Support:** Native DVC integration
- **Team Collaboration:** Shared experiment tracking
- **Free Tier:** Generous free plan for small teams

### 4.2 Connecting to DagShub

#### 4.2.1 Step 1: Create DagShub Account

1. Go to <https://dagshub.com>
2. Sign up with GitHub account
3. Create a new repository or connect existing one

#### 4.2.2 Step 2: Obtain Tracking URI

After connecting your repository to DagShub:

```
# Your tracking URI will look like:  
https://dagshub.com/username/repo-name.mlflow
```

### 4.3 Model Evaluation with MLflow

#### src/Model\_Evaluation\_Rf.py - Complete Implementation

```
1 import os  
2 import numpy as np  
3 import pandas as pd  
4 import pickle  
5 import json  
6 from sklearn.metrics import (accuracy_score, precision_score,  
7                               recall_score, roc_auc_score)  
8 import logging  
9 import mlflow  
10 import yaml  
11 import dagshub  
12  
13 # Initialize DagShub connection  
14 dagshub.init(  
15     repo_owner='Error-Makes-Clever',  
16     repo_name='MLOPS-Dagshub-DVC-Git-Project',  
17     mlflow=True  
18 )  
19  
20 # Set MLflow tracking URI
```

```
21 mlflow.set_tracking_uri(
22     "https://dagshub.com/Error-Makes-Clever/"
23     "MLOPS-Dagshub-DVC-Git-Project.mlflow"
24 )
25
26 # Set experiment name
27 mlflow.set_experiment("MLOPS-Project-with-Random-Forest")
28
29 # Ensure the "logs" directory exists
30 log_dir = 'logs'
31 os.makedirs(log_dir, exist_ok=True)
32
33 # Logging configuration
34 logger = logging.getLogger('Model_Evaluation')
35 logger.setLevel('DEBUG')
36
37 console_handler = logging.StreamHandler()
38 console_handler.setLevel('DEBUG')
39
40 log_file_path = os.path.join(log_dir, 'Model_Evaluation.log')
41 file_handler = logging.FileHandler(log_file_path)
42 file_handler.setLevel('DEBUG')
43
44 formatter = logging.Formatter(
45     '%(asctime)s - %(name)s - %(levelname)s - %(message)s',
46 )
47 console_handler.setFormatter(formatter)
48 file_handler.setFormatter(formatter)
49
50 logger.addHandler(console_handler)
51 logger.addHandler(file_handler)
52
53
54 def load_params(params_path: str) -> dict:
55     """Load parameters from a YAML file."""
56     try:
57         with open(params_path, 'r') as file:
58             params = yaml.safe_load(file)
59             logger.debug('Parameters retrieved from %s', params_path)
60         return params
61     except FileNotFoundError:
62         logger.error('File not found: %s', params_path)
63         raise
64     except yaml.YAMLError as e:
65         logger.error('YAML error: %s', e)
66         raise
67     except Exception as e:
68         logger.error('Unexpected error: %s', e)
69         raise
70
71
72 def load_model(file_path: str):
73     """Load the trained model from a file."""
74     try:
75         with open(file_path, 'rb') as file:
```

```
    model = pickle.load(file)
    logger.debug('Model loaded from %s', file_path)
    return model
except FileNotFoundError:
    logger.error('File not found: %s', file_path)
    raise
except Exception as e:
    logger.error('Unexpected error while loading model: %s', e)
    raise

def load_data(file_path: str) -> pd.DataFrame:
    """Load data from a CSV file."""
    try:
        df = pd.read_csv(file_path)
        logger.debug('Data loaded from %s', file_path)
        return df
    except pd.errors.ParserError as e:
        logger.error('Failed to parse the CSV file: %s', e)
        raise
    except Exception as e:
        logger.error('Unexpected error while loading data: %s', e)
        raise

def evaluate_model(clf, X_test: np.ndarray,
                   y_test: np.ndarray) -> dict:
    """Evaluate the model and return metrics."""
    try:
        y_pred = clf.predict(X_test)
        y_pred_proba = clf.predict_proba(X_test)[:, 1]

        accuracy = accuracy_score(y_test, y_pred)
        precision = precision_score(y_test, y_pred)
        recall = recall_score(y_test, y_pred)
        auc = roc_auc_score(y_test, y_pred_proba)

        metrics_dict = {
            'accuracy': accuracy,
            'precision': precision,
            'recall': recall,
            'auc': auc
        }
        logger.debug('Model evaluation metrics calculated')
        return metrics_dict
    except Exception as e:
        logger.error('Error during model evaluation: %s', e)
        raise

def save_metrics(metrics: dict, file_path: str) -> None:
    """Save the evaluation metrics to a JSON file."""
    try:
        os.makedirs(os.path.dirname(file_path), exist_ok=True)
```

```

130
131     with open(file_path, 'w') as file:
132         json.dump(metrics, file, indent=4)
133     logger.debug('Metrics saved to %s', file_path)
134 except Exception as e:
135     logger.error('Error while saving metrics: %s', e)
136     raise
137
138
139 def main():
140     try:
141         params = load_params('params.yaml')
142
143         clf = load_model('./models/model_rf.pkl')
144         test_data = load_data('./data/processed/test_tfidf.csv')
145     )
146
147     X_test = test_data.iloc[:, :-1].values
148     y_test = test_data.iloc[:, -1].values
149
150     metrics = evaluate_model(clf, X_test, y_test)
151
152     # MLflow tracking
153     with mlflow.start_run():
154         # Log metrics
155         mlflow.log_metric('accuracy', metrics['accuracy'])
156         mlflow.log_metric('precision', metrics['precision'])
157
158         # Log parameters
159         mlflow.log_param('n_estimators',
160                           params['model_building'][,
161                           'n_estimators'])
162         mlflow.log_param('random_state',
163                           params['model_building'][,
164                           'random_state'])
165         mlflow.log_param('max_depth',
166                           params['model_building'][, 'max_depth'])
167
168         mlflow.log_param('max_features',
169                           params['feature_engineering'][,
170                           'max_features'])
171         mlflow.log_param('test_size',
172                           params['data_ingestion'][, 'test_size'])
173
174     save_metrics(metrics, 'reports/metrics.json')
175 except Exception as e:
176     logger.error('Failed to complete model evaluation: %s',
177     e)
178     print(f"Error: {e}")
179
180
181 if __name__ == '__main__':

```

178        **main()**

## 4.4 Understanding the MLflow Integration

### 4.4.1 Key Components

1. **dagshub.init()**: Establishes connection to DagShub

```
1 dagshub.init(  
2     repo_owner='username',  
3     repo_name='project-name',  
4     mlflow=True  
5 )  
6
```

2. **mlflow.set\_tracking\_uri()**: Points MLflow to DagShub

```
1 mlflow.set_tracking_uri(  
2     "https://dagshub.com/username/project.mlflow"  
3 )  
4
```

3. **mlflow.set\_experiment()**: Groups related runs

```
1 mlflow.set_experiment("MLOPS-Project-with-Random-Forest")  
2
```

4. **mlflow.start\_run()**: Creates a new experiment run

```
1 with mlflow.start_run():  
2     mlflow.log_metric('accuracy', 0.95)  
3     mlflow.log_param('n_estimators', 100)  
4
```

## MLflow Concepts

**Experiment:** A collection of related runs (e.g., all RandomForest experiments)

**Run:** A single execution of your ML code with specific parameters

**Metrics:** Performance measurements (accuracy, precision, recall)

**Parameters:** Hyperparameters used in the run (n\_estimators, max\_depth)

**Artifacts:** Files generated during the run (models, plots, reports)

## 5 Configuration Files

### 5.1 Step 7: Create .gitignore

```
.gitignore

# Virtual environment
venv/
env/

# Data directories (DVC will track these)
data/
models/
reports/

# Logs
logs/

# Python cache
__pycache__/
*.pyc
*.pyo

# IDE
.vscode/
.idea/

# OS
.DS_Store
Thumbs.db
```

#### Important Note

**Important:** The `data/`, `models/`, and `reports/` directories are excluded from Git because DVC will handle versioning for these large files.

### 5.2 Step 8: Initialize DVC and Git

```
# Initialize Git
git init

# Initialize DVC
dvc init
```

#### What happens:

- Git creates `.git/` directory
- DVC creates `.dvc/` directory
- DVC modifies `.gitignore` to exclude cache

### 5.3 Step 9: Create DVC Pipeline

#### dvc.yaml

```
1 stages:
2   data_ingestion:
3     cmd: python src/Data_Ingestion.py
4     deps:
5       - src/Data_Ingestion.py
6     params:
7       - data_ingestion.test_size
8     outs:
9       - data/raw
10
11   data_preprocessing:
12     cmd: python src/Data_Pre_Processing.py
13     deps:
14       - data/raw
15       - src/Data_Pre_Processing.py
16     outs:
17       - data/interim
18
19   feature_engineering:
20     cmd: python src/Feature_Engineering.py
21     deps:
22       - data/interim
23       - src/Feature_Engineering.py
24     params:
25       - feature_engineering.max_features
26     outs:
27       - data/processed
28
29   model_building:
30     cmd: python src/Model_Building_Rf.py
31     deps:
32       - data/processed
33       - src/Model_Building_Rf.py
34     params:
35       - model_building.n_estimators
36       - model_building.random_state
37       - model_building.max_depth
38     outs:
39       - models/model_rf.pkl
40
41   model_evaluation:
42     cmd: python src/Model_Evaluation_Rf.py
43     deps:
44       - models/model_rf.pkl
45       - src/Model_Evaluation_Rf.py
46     metrics:
47       - reports/metrics.json
```

## 5.4 Step 10: Create Parameters File

params.yaml

```
1 data_ingestion:  
2   test_size: 0.25  
3  
4 feature_engineering:  
5   max_features: 35  
6  
7 model_building:  
8   n_estimators: 22  
9   random_state: 2  
10  max_depth: 11
```

## 6 Initial Git Commit and DVC Setup

---

### 6.1 Step 11: Check Git Status

```
git status
```

**Expected Output:**

```
On branch main

No commits yet

Changes to be committed:
  (use "git rm --cached <file>..." to unstage)
    new file:   .dvc/.gitignore
    new file:   .dvc/config
    new file:   .dvcignore

Untracked files:
  (use "git add <file>..." to include in what will be committed)
    .gitignore
    dvc.yaml
    params.yaml
    requirements.txt
    src/
```

### 6.2 Step 12: Check DVC Status

```
dvc status
```

**Expected Output:**

```
data_ingestion:
  changed deps:
    modified:           src\Data_Ingestion.py
    new:                params.yaml
  changed outs:
    deleted:           data\raw

data_preprocessing:
  changed deps:
    deleted:           data\raw
    modified:          src\Data_Pre_Processing.py
  changed outs:
    deleted:           data\interim

feature_engineering:
  changed deps:
    deleted:           data\interim
    modified:          src\Feature_Engineering.py
    new:                params.yaml
  changed outs:
    deleted:           data\processed
```

```

model_building:
  changed deps:
    deleted:          data\processed
    modified:         src\Model_Building_Rf.py
    new:              params.yaml
  changed outs:
    deleted:         models\model_rf.pkl

model_evaluation:
  changed deps:
    deleted:         models\model_rf.pkl
    modified:        src\Model_Evaluation_Rf.py
  changed outs:
    deleted:        reports\metrics.json

```

### Important Note

This output is normal! DVC is detecting that the pipeline has never been run. All outputs are listed as "deleted" because they don't exist yet.

## 6.3 Step 13: Make Initial Git Commit

```

# Stage all files
git add .

# Commit
git commit -m "Initial commit: set up ML pipeline components"

# Push to remote (after creating GitHub repo)
git remote add origin https://github.com/username/project.git
git push -u origin main

```

## 6.4 Step 14: Configure DVC Remote Storage

### 6.4.1 Create Local Storage Directory

```

# Create S3 directory for local storage
mkdir S3

```

### 6.4.2 Add as DVC Remote

```

# Add local directory as DVC remote
dvc remote add -d dvc_origin S3

# Verify remote configuration
dvc remote list

```

#### Expected Output:

```
dvc_origin      S3      (default)
```

#### 6.4.3 Check DVC Configuration

```
# View .dvc/config file
cat .dvc/config
```

Content:

```
[core]
remote = dvc_origin

['remote "dvc_origin"]'
url = S3
```

#### 6.4.4 Add S3 to .gitignore

```
# Add to .gitignore
echo "S3/" >> .gitignore
```

#### Warning

**Important:** The S3 directory should NOT be tracked by Git. It will contain all your data versions and can become very large.

## 7 Running the Pipeline and First Experiment

### 7.1 Step 15: Execute the Pipeline

```
dvc repro
```

### 7.2 Pipeline Execution Output

#### Complete Pipeline Output

```
Running stage 'data_ingestion':  
> python src/Data_Ingestion.py  
2025-12-20 18:28:06,435 - Data_Ingestion - DEBUG - Parameters  
    retrieved from params.yaml  
2025-12-20 18:28:07,002 - Data_Ingestion - DEBUG - Data loaded  
    from https://raw.githubusercontent.com/...  
2025-12-20 18:28:07,004 - Data_Ingestion - DEBUG - Data  
    preprocessing completed  
2025-12-20 18:28:07,034 - Data_Ingestion - DEBUG - Train and  
    test data saved to ./data/raw  
Updating lock file 'dvc.lock'  
  
Running stage 'data_preprocessing':  
> python src/Data_Pre_Processing.py  
[nltk_data] Downloading package stopwords...  
[nltk_data]   Package stopwords is already up-to-date!  
[nltk_data] Downloading package punkt...  
[nltk_data]   Package punkt is already up-to-date!  
2025-12-20 18:28:16,011 - Data_Pre_Processing - DEBUG - Data  
    loaded properly  
2025-12-20 18:28:16,012 - Data_Pre_Processing - DEBUG - Starting  
    preprocessing for DataFrame  
2025-12-20 18:28:16,015 - Data_Pre_Processing - DEBUG - Target  
    column encoded  
2025-12-20 18:28:16,023 - Data_Pre_Processing - DEBUG -  
    Duplicates removed  
2025-12-20 18:28:40,913 - Data_Pre_Processing - DEBUG - Text  
    column transformed  
Updating lock file 'dvc.lock'  
  
Running stage 'feature_engineering':  
> python src/Feature_Engineering.py  
2025-12-20 18:28:56,433 - Feature_Engineering - DEBUG -  
    Parameters retrieved from params.yaml  
2025-12-20 18:28:56,452 - Feature_Engineering - DEBUG - Data  
    loaded and NaNs filled  
2025-12-20 18:28:56,607 - Feature_Engineering - DEBUG - TF-IDF  
    applied and data transformed  
Updating lock file 'dvc.lock'  
  
Running stage 'model_building':
```

```
> python src/Model_Building_Rf.py
2025-12-20 18:29:04,072 - Model_Building - DEBUG - Parameters
    retrieved from params.yaml
2025-12-20 18:29:04,101 - Model_Building - DEBUG - Data loaded
    with shape (4152, 26)
2025-12-20 18:29:04,102 - Model_Building - DEBUG - Initializing
    RandomForest model
2025-12-20 18:29:04,103 - Model_Building - DEBUG - Model training
    started with 4152 samples
2025-12-20 18:29:04,211 - Model_Building - DEBUG - Model training
    completed
2025-12-20 18:29:04,214 - Model_Building - DEBUG - Model saved
    to models/model_rf.pkl
Updating lock file 'dvc.lock'

Running stage 'model_evaluation':
> python src/Model_Evaluation_Rf.py
Accessing as Error-Makes-Clever
Initialized MLflow to track repo
    "Error-Makes-Clever/MLOPS-Dagshub-DVC-Git-Project"
Repository Error-Makes-Clever/MLOPS-Dagshub-DVC-Git-Project
    initialized!
2025-12-20 18:29:20,460 - Model_Evaluation - DEBUG - Parameters
    retrieved from params.yaml
2025-12-20 18:29:20,685 - Model_Evaluation - DEBUG - Model loaded
2025-12-20 18:29:20,692 - Model_Evaluation - DEBUG - Data loaded
2025-12-20 18:29:20,714 - Model_Evaluation - DEBUG - Model
    evaluation metrics calculated
View run handsome-asp-382 at:
    https://dagshub.com/.../runs/beaa07ff...
View experiment at:
    https://dagshub.com/.../experiments/0
2025-12-20 18:29:25,787 - Model_Evaluation - DEBUG - Metrics
    saved to reports/metrics.json
Updating lock file 'dvc.lock'

To track the changes with git, run:
    git add dvc.lock

To enable auto staging, run:
    dvc config core.autostage true

Use 'dvc push' to send your updates to remote storage.
```

### 7.3 Understanding the Output

1. **Stage Execution:** Each stage runs sequentially
2. **dvc.lock Updates:** After each stage, DVC updates the lock file
3. **MLflow Integration:** Model evaluation connects to DagShub

4. **Run Link:** DagShub provides a direct link to view the run

## 7.4 Step 16: Understanding DVC Cache

After running the pipeline, check your project structure:

```
project/
+-- .dvc/
|   +-- cache/           # Contains cached artifacts
|   +-- config
|   +-- .gitignore
+-- S3/                  # Remote storage (local directory)
+-- data/
|   +-- raw/
|   +-- interim/
|   +-- processed/
+-- models/
|   +-- model_rf.pkl
+-- reports/
|   +-- metrics.json
+-- logs/
|   +-- *.log files
+-- dvc.lock             # Generated lock file
```

### What's in Each Directory?

#### .dvc/cache/:

- Stores all versions of tracked data locally
- Content-addressable storage (files named by hash)
- Acts as local backup

#### S3/:

- Remote storage location
- Mirrors .dvc/cache/ structure
- Enables team sharing

#### dvc.lock:

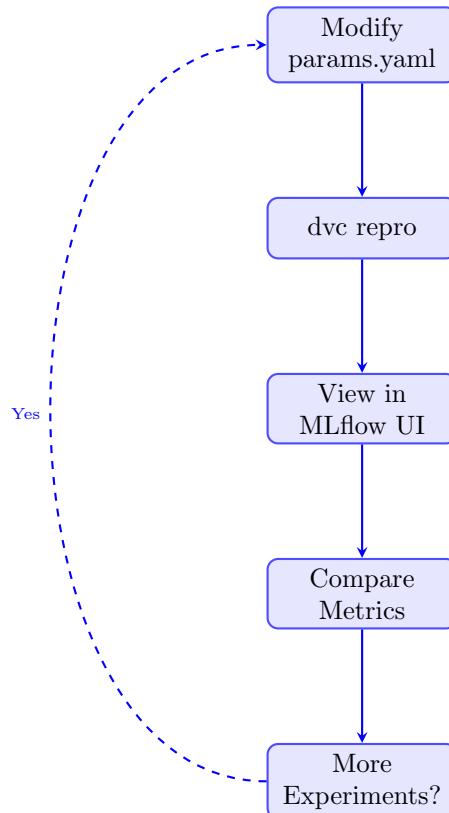
- Records exact file versions (MD5 hashes)
- Ensures reproducibility
- Tracked by Git

## 8 Running Multiple Experiments

### 8.1 Step 17: Experimentation Phase

Now we'll run multiple experiments with different hyperparameters to find the best model.

### 8.2 Experiment Workflow



### 8.3 Running Experiment Variations

#### 8.3.1 Experiment 1: Initial Configuration

```
1 # params.yaml - Run 1
2 data_ingestion:
3     test_size: 0.25
4
5 feature_engineering:
6     max_features: 35
7
8 model_building:
9     n_estimators: 22
10    random_state: 2
11    max_depth: 11
```

dvc repro

#### 8.3.2 Experiment 2: Fewer Trees

```
1 # params.yaml - Run 2
2 data_ingestion:
3   test_size: 0.25
4
5 feature_engineering:
6   max_features: 35
7
8 model_building:
9   n_estimators: 18      # Changed
10  random_state: 2
11  max_depth: 5         # Changed
```

```
dvc repro
```

### 8.3.3 Experiment 3: More Features

```
1 # params.yaml - Run 3
2 data_ingestion:
3   test_size: 0.25
4
5 feature_engineering:
6   max_features: 50      # Changed
7
8 model_building:
9   n_estimators: 22
10  random_state: 2
11  max_depth: 11
```

```
dvc repro
```

### 8.3.4 Experiment 4: Different Split

```
1 # params.yaml - Run 4
2 data_ingestion:
3   test_size: 0.20      # Changed
4
5 feature_engineering:
6   max_features: 35
7
8 model_building:
9   n_estimators: 30      # Changed
10  random_state: 2
11  max_depth: 15         # Changed
```

```
dvc repro
```

### 8.3.5 Experiment 5: Balanced Configuration

```
1 # params.yaml - Run 5
2 data_ingestion:
3   test_size: 0.25
4
```

```
5 feature_engineering:  
6   max_features: 40      # Changed  
7  
8 model_building:  
9   n_estimators: 25      # Changed  
10  random_state: 2  
11  max_depth: 12         # Changed
```

```
dvc repro
```

## 9 Step 18: Analyzing Experiments in MLflow

### 9.1 Accessing the MLflow UI

Navigate to your DagShub MLflow tracking URI:

```
https://dagshub.com/Error-Makes-Clever/  
MLOPS-Dagshub-DVC-Git-Project.mlflow
```

### 9.2 Experiment Results Summary

Run	AUC	Accuracy	Precision	Recall
Run 1 (able-ray-853)	0.9187	0.9378	0.8302	0.6947
Run 2 (handsome-asp-382)	0.8656	0.8919	0.8776	0.2756
Run 3 (placid-gnu-625)	0.8846	0.8922	0.8704	0.2765
Run 4 (wise-elk-465)	0.8994	0.9254	0.7933	0.6263
Run 5 (receptive-crow-130)	0.9178	0.9334	0.8414	0.6421

### 9.3 Detailed Run Analysis

#### 9.3.1 Run 1: able-ray-853

##### Run 1 Metrics

###### Performance:

- AUC: 0.9187
- Accuracy: 0.9378
- Precision: 0.8302
- Recall: 0.6947

###### Parameters:

- n\_estimators: 22
- max\_depth: 11
- max\_features: 35
- test\_size: 0.25

**Analysis:** Best overall performance with highest AUC and accuracy.

#### 9.3.2 Run 2: handsome-asp-382

##### Run 2 Metrics

###### Performance:

- AUC: 0.8656
- Accuracy: 0.8919
- Precision: 0.8776

- Recall: 0.2756

**Parameters:**

- n\_estimators: 18
- max\_depth: 5
- max\_features: 35
- test\_size: 0.25

**Analysis:** High precision but very low recall. Model is too conservative.

### 9.3.3 Run 5: receptive-crow-130

#### Run 5 Metrics

**Performance:**

- AUC: 0.9178
- Accuracy: 0.9334
- Precision: 0.8414
- Recall: 0.6421

**Parameters:**

- n\_estimators: 25
- max\_depth: 12
- max\_features: 40
- test\_size: 0.25

**Analysis:** Second-best performance, close to Run 1.

## 10 Step 19: Deciding What "Best" Means

### 10.1 Defining Success Criteria

Before committing anything, define one primary metric based on your use case:

Metric Priority	Use Case
Recall	Spam detection - catch all spam (few false negatives)
Precision	Medical diagnosis - avoid false positives
Accuracy	General classification with balanced classes
AUC	Ranking, threshold tuning, imbalanced data
F1-Score	Balance between precision and recall

### 10.2 Analysis for Spam Classification

For spam detection, we typically want:

1. High AUC (good separation between classes)
2. High accuracy (overall correctness)
3. Balanced precision and recall

### 10.3 Ranking the Runs

Run	AUC	Accuracy	Recommendation
Run 1	0.9187	0.9378	Winner ✓
Run 5	0.9178	0.9334	Runner-up
Run 4	0.8994	0.9254	Good
Run 3	0.8846	0.8922	Acceptable
Run 2	0.8656	0.8919	Poor recall

#### Important Note

**Winner: Run 1 (able-ray-853)**

**Justification:**

- Highest AUC (0.9187) - best class separation
- Highest accuracy (0.9378) - best overall performance
- Balanced precision (0.8302) and recall (0.6947)
- Reliable and reproducible

## 11 Step 20: Re-creating the Winning Run

### 11.1 Why Re-create?

#### Warning

##### Critical MLOps Principle:

Earlier runs were **exploratory**. Now you are making a **decision**.

##### Why re-create:

- Makes the decision explicit
- Creates a clean Git commit representing the approved model
- Enables future rollback to this exact state
- Documents which experiment was promoted to production

**Don't skip this step!** It's the difference between "I ran some experiments" and "This is our production model."

### 11.2 Obtaining Winning Parameters

Go to MLflow UI → Click on Run 1 (able-ray-853) → View Parameters:

Parameter	Value
test_size	0.25
n_estimators	22
random_state	2
max_depth	11
max_features	35

### 11.3 Step 21: Update params.yaml

#### params.yaml - Winning Configuration

```

1 data_ingestion:
2   test_size: 0.25
3
4 feature_engineering:
5   max_features: 35
6
7 model_building:
8   n_estimators: 22
9   random_state: 2
10  max_depth: 11

```

### 11.4 Step 22: Run Pipeline with Winning Parameters

```
dvc repro
```

Expected Output:

```
Stage 'data_ingestion' didn't change, skipping
Stage 'data_preprocessing' didn't change, skipping
Stage 'feature_engineering' didn't change, skipping
Stage 'model_building' is cached - skipping run,
    checking out outputs
Updating lock file 'dvc.lock'
```

```
Stage 'model_evaluation' is cached - skipping run,
    checking out outputs
Updating lock file 'dvc.lock'
```

To track the changes with git, run:  
`git add dvc.lock`

To enable auto staging, run:  
`dvc config core.autostage true`

Use 'dvc push' to send your updates to remote storage.

### Why No New MLflow Run?

**DVC detected that:**

- Code did not change
- Parameters did not change
- Data did not change

**Result:** DVC used cached outputs. No new execution = no new MLflow run.

**This is correct behavior!** We're verifying the winning configuration, not running a new experiment.

## 11.5 Step 23: Verify Metrics (Sanity Check)

```
# Check metrics.json
cat reports/metrics.json
```

**Expected Content:**

```
{
    "accuracy": 0.9378200438917337,
    "precision": 0.8301886792452831,
    "recall": 0.6947368421052632,
    "auc": 0.9186759379331932
}
```

### Important Note

**Verification:**

- Metrics match (or are very close to) Run 1
- No accidental changes

- Configuration is reproducible

If metrics differ wildly → check random seed and data consistency.

## 12 Step 24: The Critical Commit

### 12.1 Why This Commit Matters

#### The Freezing Commit

This is the **most important commit** in your MLOps workflow.

##### What this commit represents:

- ✓ This model is **approved**
- ✓ This state is **reproducible**
- ✓ This is **rollback-safe**
- ✓ This is **deployable**

##### What gets committed:

- Code (src/)
- Pipeline definition (dvc.yaml)
- Parameters (params.yaml)
- Lock file (dvc.lock) - contains exact data/model versions

##### What does NOT get committed:

- Data files (tracked by DVC)
- Model files (tracked by DVC)
- Logs

### 12.2 Check Git Status

```
git status
```

#### Expected Output:

On branch main

Your branch is up to date with 'origin/main'.

#### Changes not staged for commit:

```
(use "git add <file>..." to update what will be committed)
(use "git restore <file>..." to discard changes)
    modified:   .dvc/config
    modified:   dvc.yaml
    modified:   requirements.txt
    modified:   src/Model_Evaluation_Rf.py
```

#### Untracked files:

```
(use "git add <file>..." to include in what will be committed)
    dvc.lock
    logs/
```

no changes added to commit

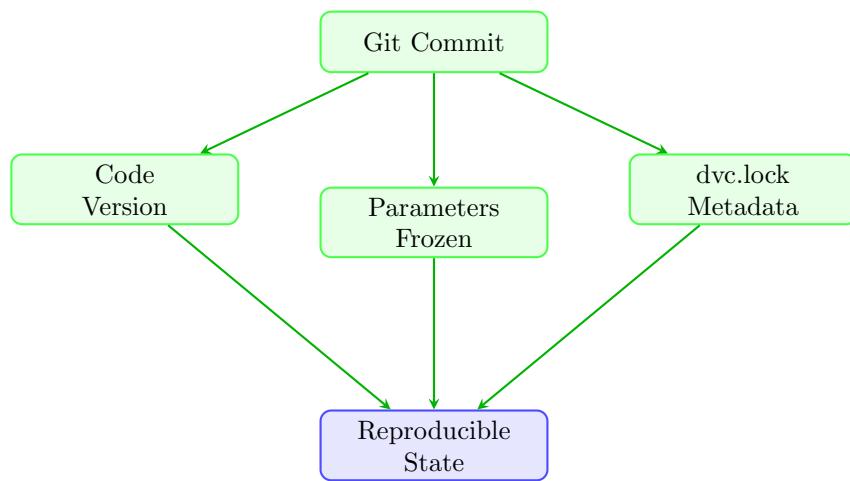
### 12.3 Stage and Commit

```
# Stage all changes
git add .

# Commit with descriptive message
git commit -m "Freeze best RF model (AUC=0.918, accuracy=0.937)"

# Push to remote
git push origin main
```

### 12.4 What This Commit Means



## 13 Step 25: Push Data and Model to DVC Remote

### 13.1 Why Push to DVC Remote?

#### Warning

**Git vs DVC Storage:**

**What Git has:**

- Code
- dvc.lock (metadata with hashes)
- Parameters

**What Git does NOT have:**

- Actual data files
- Actual model files

**Problem:** Without `dvc push`, teammates cannot reproduce your work!

### 13.2 Execute DVC Push

```
dvc push
```

**Expected Output:**

```
Collecting                         | 8.00 [00:00, 2.50entry/s]
Pushing                           | 8.00 [00:02, 3.20file/s]
8 files pushed
```

### 13.3 What Gets Pushed

Local Location	DVC Remote (S3)
data/raw/	Uploaded to S3/
data/interim/	Uploaded to S3/
data/processed/	Uploaded to S3/
models/model_rf.pkl	Uploaded to S3/
reports/metrics.json	Uploaded to S3/

### 13.4 Verification

```
# Check S3 directory
ls S3/

# Should see hash-named directories
```

#### Complete State Now

**GitHub has:**

- Code (src/)

- Pipeline (dvc.yaml)
- Parameters (params.yaml)
- Metadata (dvc.lock)

**DVC Remote (S3) has:**

- All data versions
- All model versions
- All reports

**Result:** Complete reproducibility ✓

## 14 Step 26: Linking Git Commits to MLflow Runs

### 14.1 Why Tag MLflow Runs?

#### The Problem

##### Situation:

- You have 5 MLflow runs
- You committed Run 1 as the winner
- Weeks later: "Which MLflow run is our production model?"

**Solution:** Tag the MLflow run with the Git commit hash!

##### Benefit:

- Clear link between MLflow experiment and Git commit
- Easy to identify production model
- Audit trail for compliance

### 14.2 Create Tagging Script

#### Adding tag\_to\_runs.py

```
 1 import subprocess
 2 import mlflow
 3 from mlflow.tracking import MlflowClient
 4 import dagshub
 5
 6 # Authenticate with DagShub
 7 dagshub.init(
 8     repo_owner="Error-Makes-Clever",
 9     repo_name="MLOPS-Dagshub-DVC-Git-Project",
10     mlflow=True
11 )
12
13 # Set tracking URI
14 mlflow.set_tracking_uri(
15     "https://dagshub.com/Error-Makes-Clever/"
16     "MLOPS-Dagshub-DVC-Git-Project.mlflow"
17 )
18
19 # Initialize MLflow client
20 client = MlflowClient()
21
22 # Run ID from MLflow UI (Run 1)
23 run_id = "eaa2537460244123bc672bee9c22da05"
24
25 # Get current Git commit hash
26 commit = subprocess.check_output(
27     ["git", "rev-parse", "HEAD"]
28 ).decode().strip()
29
```

```

30 # Add tag to MLflow run
31 client.set_tag(run_id, "git_commit", commit)
32
33 # Optional: Add more tags
34 client.set_tag(run_id, "status", "production")
35 client.set_tag(run_id, "approved_by", "data_science_team")
36 client.set_tag(run_id, "deployment_date", "2025-12-20")
37
38 print(f"Tags successfully added to run {run_id}")
39 print(f"    Git commit: {commit}")

```

## 14.3 Understanding the Script

### 14.3.1 Key Components

#### 1. Get Git Commit Hash:

```

1 commit = subprocess.check_output(
2     ["git", "rev-parse", "HEAD"]
3 ).decode().strip()
4

```

#### 2. Set Tag in MLflow:

```

1 client.set_tag(run_id, "git_commit", commit)
2

```

#### 3. Additional Metadata Tags:

```

1 client.set_tag(run_id, "status", "production")
2 client.set_tag(run_id, "approved_by", "team_name")
3

```

## 14.4 Finding the Run ID

### Option 1: From MLflow UI:

1. Go to DagShub MLflow interface
2. Click on Run 1 (able-ray-853)
3. Copy the Run ID from URL or run details

### Option 2: From Console Output:

View run at: .../runs/eaa2537460244123bc672bee9c22da05  
 ^^^^^^^^^^^^^^^^^^  
 This is the Run ID

## 14.5 Step 27: Execute Tagging Script

```
python Adding_tag_to_runs.py
```

### Expected Output:

Tags successfully added to run eaa2537460244123bc672bee9c22da05  
 Git commit: a3f8e2c9d1b4e5f6g7h8i9j0k1l2m3n4

## 14.6 Step 28: Verify Tags in MLflow UI

1. Go to MLflow UI
2. Click on Run 1 (able-ray-853)
3. Scroll to "Tags" section
4. Verify tags appear:

Tag Key	Tag Value
git_commit	a3f8e2c9d1b4e5f6...
status	production
approved_by	data_science_team
deployment_date	2025-12-20

### Important Note

**Best Practice:** Always tag production runs with:

- Git commit hash
- Deployment status
- Approval information
- Date/time of promotion

This creates a complete audit trail for compliance and debugging.

## 15 Step 29: Multi-Model Experimentation

### 15.1 Scenario: Team with Multiple Models

#### Real-World Team Setup

##### Scenario:

- Data Scientist A: Working on RandomForest
- Data Scientist B: Working on XGBoost
- Data Scientist C: Working on Neural Networks

**Goal:** Everyone shares the same infrastructure but tracks separate experiments.

### 15.2 Using Different Experiment Names

#### 15.2.1 RandomForest Experiments

```
1 # src/Model_Evaluation_Rf.py
2 mlflow.set_experiment("MLOPS-Project-with-Random-Forest")
```

#### 15.2.2 XGBoost Experiments

```
1 # src/Model_Evaluation_Xgb.py
2 mlflow.set_experiment("MLOPS-Project-with-XGBoost")
```

#### 15.2.3 Neural Network Experiments

```
1 # src/Model_Evaluation_NN.py
2 mlflow.set_experiment("MLOPS-Project-with-Neural-Networks")
```

### 15.3 Benefits of Separate Experiments

1. **Organization:** Experiments grouped by model type
2. **Clear Ownership:** Each team member manages their experiment
3. **Easy Comparison:** Compare within model family first
4. **Scalability:** Add new model types without conflicts

### 15.4 Viewing All Experiments

In the DagShub MLflow UI:

Experiment	Runs	Owner
MLOPS-Project-with-Random-Forest	5	Data Scientist A
MLOPS-Project-with-XGBoost	8	Data Scientist B
MLOPS-Project-with-Neural-Networks	12	Data Scientist C

**Important Note****Key Advantage:**

Everyone in your team can:

- See all experiments
- Compare across models
- Learn from each other's work
- Avoid duplicate efforts

This is impossible with local-only experiment tracking!

## 16 Reproducing the Project on a New Machine

### 16.1 The Reproducibility Challenge

#### Warning

##### Common Scenario:

A new team member joins or you need to reproduce the model on a production server.

##### Requirements:

- Exact code version
- Exact data version
- Exact model version
- Same environment

**Solution:** Git + DVC + proper setup

### 16.2 Important Note About Local DVC Remote

#### Critical: Local S3 Directory Limitation

In this project, the DVC remote is configured as a **local directory**:

```
S3/ (located at: S:\MLOPS-Dagshub-DVC-Git-Project\S3)
```

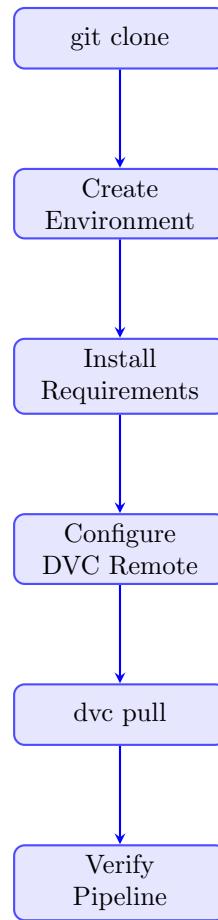
**Problem:** This is a local path, NOT accessible from other machines!

##### Options for Team Sharing:

1. Place S3 folder on **shared network drive**
2. Use **NAS / file server**
3. Map the **same path on all machines** (e.g., Z:\Shared\S3)
4. **Replace with real cloud remote** (AWS S3, GCS, Azure Blob)

**Important:** If S3 directory is NOT accessible, `dvc pull` will fail!

### 16.3 Reproduction Workflow



### 16.4 Step-by-Step Reproduction

#### 16.4.1 Step 1: Clone Repository

```
# Clone Git repository
git clone https://github.com/Error-Makes-Clever/
    MLOPS-Dagshub-DVC-Git-Project.git

cd MLOPS-Dagshub-DVC-Git-Project
```

**What this does:**

- Downloads all code files
- Downloads dvc.yaml (pipeline definition)
- Downloads dvc.lock (file version metadata)
- Downloads params.yaml (hyperparameters)
- Does NOT download data/models (DVC handles this)

#### 16.4.2 Step 2: Create and Activate Environment

```
# Create virtual environment
python -m venv venv

# Activate (Windows)
venv\Scripts\activate

# Activate (Linux/Mac)
source venv/bin/activate
```

Verify activation:

```
# Check Python path (should point to venv)
which python    # Linux/Mac
where python    # Windows
```

#### 16.4.3 Step 3: Install Requirements

```
pip install -r requirements.txt
```

Expected packages:

- dvc
- numpy, pandas
- matplotlib, wordcloud
- nltk
- scikit-learn
- xgboost
- pyyaml
- mlflow
- dagshub

Verify installation:

```
# Check installed packages
pip list

# Verify DVC
dvc version

# Verify MLflow
mlflow --version
```

#### 16.4.4 Step 4: Configure DVC Remote Access

```
# Check current remote configuration
dvc remote list
```

**Expected Output:**

```
dvc_origin    S3
```

#### Important Note

##### Understanding the Remote Path:

The remote path in the repository is relative (S3), which means it expects the S3 folder to be in the same parent directory as the project.

##### Original setup:

```
S:\MLOPS-Dagshub-DVC-Git-Project\
    +-- S3/                  (DVC remote)
    +-- project/             (Git repository)
```

**For reproduction:** You need to ensure the S3 directory is accessible at the same relative location OR update the remote URL.

##### If path needs updating:

#### Update DVC Remote Path

```
1 # Option 1: Use shared network drive (Windows)
2 dvc remote modify dvc_origin url Z:\Shared\MLOPS\S3
3
4 # Option 2: Use UNC path (Windows)
5 dvc remote modify dvc_origin url \\server\share\S3
6
7 # Option 3: Use mounted path (Linux/Mac)
8 dvc remote modify dvc_origin url /mnt/shared/S3
9
10 # Option 4: Use absolute local path
11 dvc remote modify dvc_origin url /path/to/shared/S3
12
13 # Verify updated configuration
14 dvc remote list -v
```

##### Check .dvc/config after modification:

```
cat .dvc/config
```

**Expected content:**

```
[core]
remote = dvc_origin

['remote "dvc_origin"]'
url = /path/to/shared/S3
```

#### 16.4.5 Step 5: Pull Data and Models

```
dvc pull
```

**Expected Output:**

```
Collecting                | 8.00 [00:00, 10.0entry/s]
Fetching
Building workspace index | 8.00 [00:00, 100entry/s]
Comparing indexes         | 8.00 [00:00, 200entry/s]
Applying changes          | 8.00 [00:01, 5.00file/s]
8 files fetched
```

**What gets downloaded:**

- data/raw/train.csv
- data/raw/test.csv
- data/interim/train\_processed.csv
- data/interim/test\_processed.csv
- data/processed/train\_tfidf.csv
- data/processed/test\_tfidf.csv
- models/model\_rf.pkl
- reports/metrics.json

#### How DVC Pull Works

**Process:**

1. DVC reads `dvc.lock` to get file hashes
2. Checks if files exist in local cache (`.dvc/cache/`)
3. If missing, downloads from remote storage (S3)
4. Copies files from cache to workspace
5. Verifies checksums match

**Result:** Exact same data/models as the original machine!

#### 16.4.6 Step 6: Verify Pipeline Status

```
dvc status
```

**Expected Output:**

Data and pipelines are up to date.

**If you see changes:**

```
data_ingestion:
    changed deps:
        modified: src/Data_Ingestion.py
```

**Possible causes:**

- Line ending differences (Windows vs Linux)
- Incomplete pull
- Local modifications

**Solution:**

```
# Discard local changes
git checkout .

# Re-pull data
dvc pull

# Check status again
dvc status
```

**16.4.7 Step 7 (Optional): Re-run Pipeline**

```
dvc repro
```

**Expected Result:**

```
Stage 'data_ingestion' didn't change, skipping
Stage 'data_preprocessing' didn't change, skipping
Stage 'feature_engineering' didn't change, skipping
Stage 'model_building' is cached - skipping run
Stage 'model_evaluation' is cached - skipping run
```

Data and pipelines are up to date.

**What this means:**

- All stages are cached (no re-execution)
- No new MLflow runs created
- Guaranteed reproducibility ✓
- Pipeline is working correctly ✓

**16.5 Verification Checklist**

Item	Status	Command
Code matches	✓	git log --oneline
Data matches	✓	dvc status
Model exists	✓	ls models/
Metrics match	✓	cat reports/metrics.json
Environment ready	✓	pip list
Pipeline works	✓	dvc repro

## 16.6 Manual Verification

### 16.6.1 Verify Metrics

```
cat reports/metrics.json
```

Expected content:

```
{  
    "accuracy": 0.9378200438917337,  
    "precision": 0.8301886792452831,  
    "recall": 0.6947368421052632,  
    "auc": 0.9186759379331932  
}
```

### 16.6.2 Verify Parameters

```
cat params.yaml
```

Expected content:

```
data_ingestion:  
    test_size: 0.25  
  
feature_engineering:  
    max_features: 35  
  
model_building:  
    n_estimators: 22  
    random_state: 2  
    max_depth: 11
```

### 16.6.3 Test Model Loading

```
test_model.py
```

```
1 import pickle  
2 import pandas as pd  
3  
4 # Load model  
5 with open('models/model_rf.pkl', 'rb') as f:  
6     model = pickle.load(f)  
7  
8 print(f"Model type: {type(model)}")  
9 print(f"Number of estimators: {model.n_estimators}")  
10 print(f"Max depth: {model.max_depth}")  
11  
12 # Load test data  
13 test_data = pd.read_csv('data/processed/test_tfidf.csv')  
14 X_test = test_data.iloc[:, :-1].values  
15 y_test = test_data.iloc[:, -1].values  
16  
17 # Make predictions  
18 predictions = model.predict(X_test)  
19 print(f"Predictions shape: {predictions.shape}")
```

```
20 print(f"First 10 predictions: {predictions[:10]}")  
21  
22 print("\nModel loaded and working correctly!")
```

```
python test_model.py
```

**Expected output:**

```
Model type: <class 'sklearn.ensemble._forest.RandomForestClassifier'>  
Number of estimators: 22  
Max depth: 11  
Predictions shape: (1384,)  
First 10 predictions: [0 0 0 1 0 0 0 0 0 0]
```

Model loaded and working correctly!

## 16.7 Why This Matters in MLOps

### Reproducibility Benefits

**Without dvc pull:**

- Data is missing
- Models cannot be loaded
- Results cannot be verified
- Team collaboration breaks
- Production deployment fails

**With dvc pull:**

- Exact data restored
- Approved model retrieved
- Everyone works on same version
- Perfect reproducibility
- Confident deployment

## 16.8 Troubleshooting Reproduction Issues

### 16.8.1 Issue 1: "Remote not found"

**Error:**

```
ERROR: failed to pull data from the cloud -  
[Errno 2] No such file or directory: 'S3'
```

**Solution:**

1. Confirm S3 directory location with team

2. Update remote URL:

```
dvc remote modify dvc_origin url <correct-path>
```

3. Verify access permissions

4. Try pulling again

#### 16.8.2 Issue 2: "Permission denied"

Error:

```
ERROR: failed to pull data - [Errno 13] Permission denied
```

Solution:

- Check read permissions on shared directory
- On Linux/Mac: `ls -la /path/to/S3`
- On Windows: Check folder properties → Security tab
- Contact IT for network drive access

#### 16.8.3 Issue 3: "Checksum mismatch"

Error:

```
ERROR: checksum mismatch for 'data/raw/train.csv'
```

Solution:

```
# Remove corrupted cache
rm -rf .dvc/cache

# Re-pull data
dvc pull
```

#### 16.8.4 Issue 4: "Python package conflicts"

Error:

```
ImportError: cannot import name 'X' from 'Y'
```

Solution:

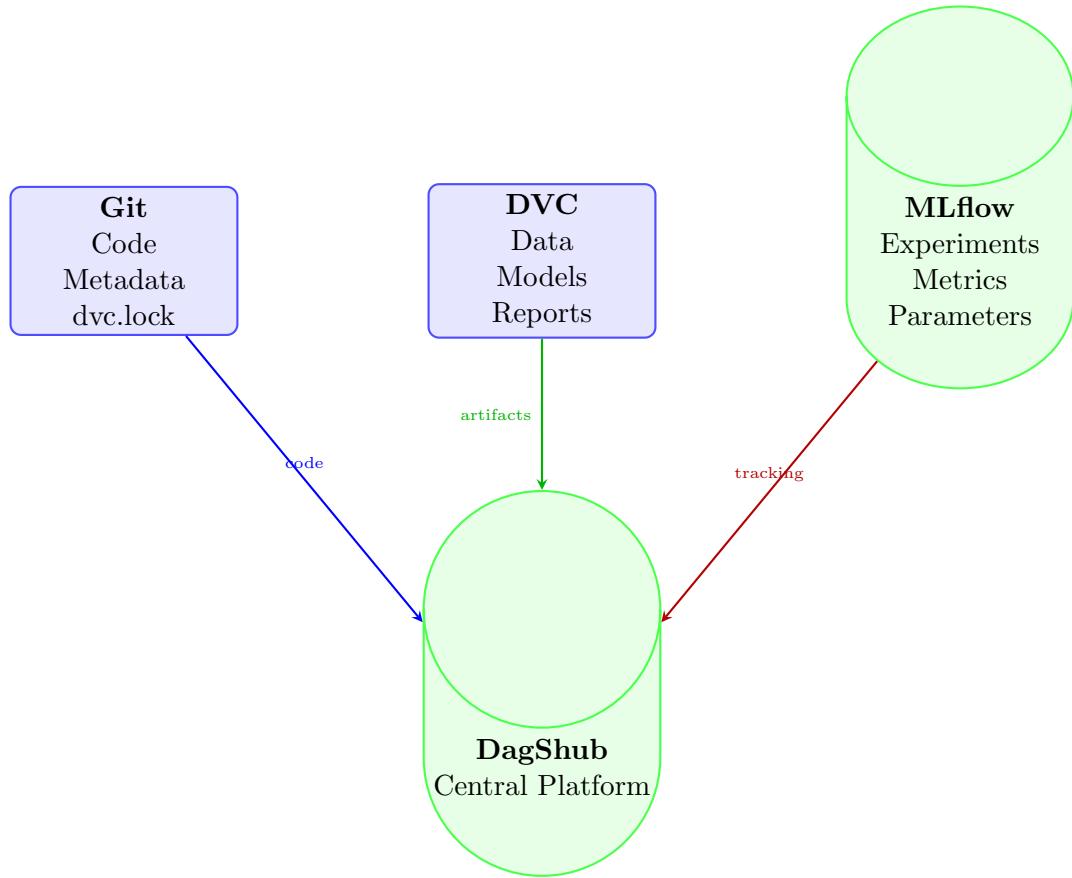
```
# Recreate environment
deactivate
rm -rf venv
python -m venv venv
source venv/bin/activate # or venv\Scripts\activate

# Reinstall exact versions
pip install -r requirements.txt

# Verify
pip list
```

## 17 Complete Mental Model

### 17.1 The Three-System Architecture



### 17.2 Understanding Each Component

#### 17.2.1 Git: Code and Metadata

##### Git Responsibilities

###### What Git tracks:

- Source code (`src/`)
- Pipeline definition (`dvc.yaml`)
- Parameters (`params.yaml`)
- Lock file (`dvc.lock`) - contains hashes
- Configuration files
- Documentation (`README.md`)

###### What Git does NOT track:

- Data files (too large)
- Model files (binary)

- Logs
- Cache directories

**Key Commands:**

- `git add .` - Stage changes
- `git commit -m "message"` - Save snapshot
- `git push` - Upload to GitHub
- `git checkout <hash>` - Switch versions

### 17.2.2 DVC: Data and Models

**DVC Responsibilities****What DVC tracks:**

- All data versions (raw, processed)
- All model versions
- Reports and metrics files
- Pipeline outputs

**How DVC works:**

- Stores actual files in remote storage (S3)
- Stores metadata (hashes) in Git (`dvc.lock`)
- Local cache (`.dvc/cache/`) for speed
- Content-addressable storage

**Key Commands:**

- `dvc add <file>` - Track file
- `dvc push` - Upload to remote
- `dvc pull` - Download from remote
- `dvc repro` - Run pipeline

### 17.2.3 MLflow: Experiment Tracking

**MLflow Responsibilities****What MLflow tracks:**

- All experiment runs
- Metrics (accuracy, precision, etc.)

- Parameters (hyperparameters)
- Artifacts (optional)
- Tags and metadata

**Benefits:**

- Visual comparison of experiments
- Centralized tracking (team-wide)
- Experiment history preserved
- Easy metric comparison

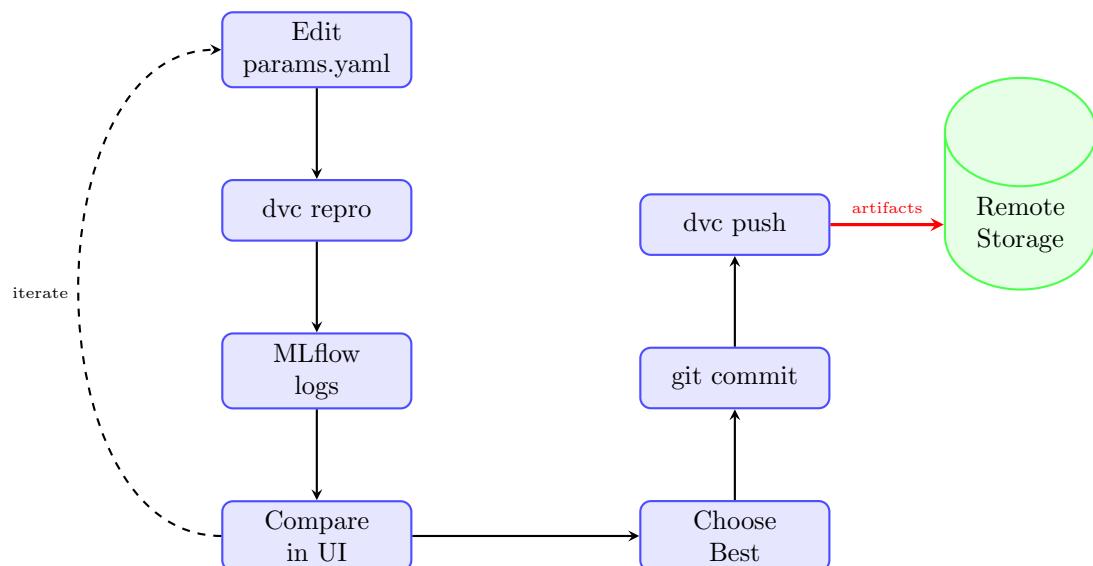
**Key Functions:**

- `mlflow.set_experiment()` - Group runs
- `mlflow.start_run()` - Begin tracking
- `mlflow.log_metric()` - Log metric
- `mlflow.log_param()` - Log parameter

### 17.3 Data Flow Summary

Command	What Happens
<code>git commit</code>	Saves code, params.yaml, dvc.lock to Git
<code>git push</code>	Uploads Git commits to GitHub
<code>dvc push</code>	Uploads data/models to DVC remote (S3)
<code>dvc pull</code>	Downloads data/models from DVC remote
<code>dvc repro</code>	Runs pipeline, updates dvc.lock
<code>mlflow.log_*</code>	Logs metrics/params to MLflow (DagShub)
<code>git checkout</code>	Restores code to specific commit

### 17.4 Complete Workflow Visualization



## 17.5 Key Concepts Review

Concept	Meaning
<b>Experiment</b>	Collection of related runs (e.g., all RandomForest tests)
<b>Run</b>	Single execution with specific parameters
<b>dvc.lock</b>	File with exact versions (hashes) of all data/models
<b>params.yaml</b>	Configuration file with hyperparameters
<b>DVC Remote</b>	Storage location for data/models (S3, GCS, local)
<b>Freezing</b>	Committing approved model configuration to Git
<b>Tagging</b>	Linking MLflow runs to Git commits
<b>Reproducibility</b>	Ability to recreate exact results on any machine

## 18 MLOps Best Practices

### 18.1 Experiment Management

#### 18.1.1 Naming Conventions

##### Good Experiment Names

- ✓ "MLOPS-Project-with-Random-Forest"
- ✓ "MLOPS-Project-with-XGBoost"
- ✓ "MLOPS-Project-with-BERT-Embeddings"
- ✓ "Spam-Detection-RandomForest-v2"
  
- ✗ "experiment1"
- ✗ "test"
- ✗ "my\_experiment"
- ✗ "asdf"

#### 18.1.2 Commit Message Best Practices

##### Good Commit Messages

- ✓ Freeze best RF model (AUC=0.918, accuracy=0.937)
- ✓ Experiment: increase max\_depth to 15, acc +0.02
- ✓ Fix: correct data leakage in preprocessing
- ✓ Add: XGBoost baseline (AUC=0.901)
- ✓ Refactor: modularize feature engineering
  
- ✗ Update
- ✗ Fix bug
- ✗ Changes
- ✗ WIP
- ✗ temp

#### 18.1.3 Parameter Documentation

```

1 # params.yaml with comments
2 data_ingestion:
3     test_size: 0.25      # Standard 75/25 split
4
5 feature_engineering:
6     max_features: 35    # Optimal from grid search
7
8 model_building:
9     n_estimators: 22    # Best balance: speed vs accuracy
10    random_state: 2     # For reproducibility
11    max_depth: 11       # Prevents overfitting

```

### 18.2 DVC Remote Management

### 18.2.1 Production Setup Recommendations

#### Warning

##### For Production: Use Cloud Storage

Why local directories don't scale:

- Not accessible across team
- No disaster recovery
- Access control issues
- Cannot use in CI/CD pipelines
- No versioning/backup

Recommended Solutions:

- AWS S3 (most popular, cost-effective)
- Google Cloud Storage (GCS)
- Azure Blob Storage
- MinIO (self-hosted S3-compatible)

### 18.2.2 Migrating to AWS S3

#### AWS S3 Setup

```
1 # Install AWS support
2 pip install dvc[s3] awscli
3
4 # Configure AWS credentials
5 aws configure
6 # Enter: Access Key ID
7 # Enter: Secret Access Key
8 # Enter: Region (e.g., us-east-1)
9 # Enter: Output format (json)
10
11 # Remove local remote
12 dvc remote remove dvc_origin
13
14 # Add S3 remote
15 dvc remote add -d s3remote s3://my-bucket/dvc-storage
16
17 # Verify configuration
18 dvc remote list
19
20 # Push data to S3
21 dvc push
22
23 # Commit updated config
24 git add .dvc/config
25 git commit -m "Migrate to AWS S3 remote storage"
```

```
26 git push
```

## 18.3 MLflow Best Practices

### 18.3.1 What to Log

Always Log	Consider Logging
Primary metrics (accuracy, AUC)	Training time
All hyperparameters	Memory usage
Data split ratios	Number of features
Random seeds	Data version/hash
Model type	Hardware info (GPU/CPU)
Feature names	Dataset size

### 18.3.2 Comprehensive Tagging Strategy

#### Complete Tagging Example

```

1 from mlflow.tracking import MlflowClient
2 import subprocess
3
4 client = MlflowClient()
5 run_id = "your-run-id"
6
7 # Tag with Git info
8 commit_hash = subprocess.check_output(
9     ["git", "rev-parse", "HEAD"])
10).decode().strip()
11 branch = subprocess.check_output(
12     ["git", "rev-parse", "--abbrev-ref", "HEAD"])
13).decode().strip()
14
15 client.set_tag(run_id, "git_commit", commit_hash)
16 client.set_tag(run_id, "git_branch", branch)
17
18 # Tag with deployment info
19 client.set_tag(run_id, "status", "production")
20 client.set_tag(run_id, "environment", "production")
21
22 # Tag with approval metadata
23 client.set_tag(run_id, "approved_by", "john_doe")
24 client.set_tag(run_id, "approval_date", "2025-12-20")
25 client.set_tag(run_id, "model_version", "v2.1.0")
26
27 # Tag with business context
28 client.set_tag(run_id, "use_case", "spam_detection")
29 client.set_tag(run_id, "department", "email_security")
30 client.set_tag(run_id, "priority", "high")
31
32 print(f"All tags added to run {run_id}")

```

## 18.4 Workflow Best Practices

### 18.4.1 Before Starting Experiments

#### 1. Define success criteria:

- What metric matters most?
- What's acceptable performance?
- What's the baseline?

#### 2. Set baseline:

- Run simplest model first
- Document baseline performance
- Use as comparison point

#### 3. Document assumptions:

- Data assumptions
- Business constraints
- Resource limitations

#### 4. Plan experiments:

- Create experiment matrix
- Prioritize high-impact experiments
- Estimate time/resources

### 18.4.2 During Experimentation

#### 1. Change one variable at a time:

- Isolate effects
- Understand what works
- Build intuition

#### 2. Log everything:

- All parameters
- All metrics
- Observations/notes
- Unexpected results

#### 3. Keep detailed notes:

- Why you tried something
- What you observed
- Hypotheses for next steps

#### 4. Compare systematically:

- Use MLflow UI
- Create comparison tables
- Visualize trends

### 18.4.3 After Choosing Best Model

#### 1. Verify reproducibility:

- Run pipeline again
- Confirm metrics match
- Check random seed

#### 2. Document decision:

- Why this model?
- What alternatives considered?
- Trade-offs made

#### 3. Freeze configuration:

- Git commit with clear message
- Include metrics in message
- Tag as milestone

#### 4. Push artifacts:

- `dvc push` to remote
- Verify upload successful
- Test pull on different machine

#### 5. Tag MLflow run:

- Link to Git commit
- Add production status
- Include approval info

#### 6. Update documentation:

- README with latest model info
- Wiki with experiment summary
- Deployment guide

## 18.5 Team Collaboration Guidelines

### 18.5.1 Communication Best Practices

- **Share experiment plans:** Before starting work
- **Review each other's runs:** Learn together
- **Discuss surprises:** Unexpected results
- **Document learnings:** Wiki or shared docs
- **Regular sync meetings:** Weekly experiment reviews
- **Clear ownership:** Who owns which experiment
- **Celebrate wins:** Share successful experiments

### 18.5.2 Avoiding Conflicts

- **Use separate experiments:** One per person/model type
- **Coordinate on shared data:** Don't modify raw data
- **Review before freezing:** Get team approval
- **Clear ownership:** Who decides what to freeze
- **Communicate changes:** Notify team of updates
- **Use branches:** Feature branches for big changes

## 18.6 Code Quality Best Practices

### 18.6.1 Logging Standards

```

1 import logging
2
3 # Configure logger
4 logger = logging.getLogger(__name__)
5 logger.setLevel('DEBUG')
6
7 # Good logging practice
8 logger.debug('Loading data from %s', file_path)
9 logger.info('Training started with %d samples', n_samples)
10 logger.warning('Missing %d features, using defaults', n_missing)
11 logger.error('Failed to load model: %s', error_msg)
12
13 # Bad logging practice
14 logger.info('Loading data') # Not informative
15 print('Training started') # Should use logger

```

### 18.6.2 Error Handling

```

1 def load_data(file_path: str) -> pd.DataFrame:
2     """Load data with comprehensive error handling."""
3     try:
4         # Attempt operation
5         df = pd.read_csv(file_path)
6         logger.info(f'Loaded {len(df)} rows from {file_path}')
7         return df
8
9     except FileNotFoundError:
10         logger.error(f'File not found: {file_path}')
11         raise
12
13     except pd.errors.ParserError as e:
14         logger.error(f'Failed to parse CSV: {e}')
15         raise
16
17     except Exception as e:
18         logger.error(f'Unexpected error: {e}')
19         raise

```

### 18.6.3 Type Hints and Documentation

```
 1 from typing import Tuple
 2 import numpy as np
 3
 4 def train_model(
 5     X_train: np.ndarray,
 6     y_train: np.ndarray,
 7     params: dict
 8 ) -> RandomForestClassifier:
 9     """
10         Train a RandomForest classifier.
11
12     Args:
13         X_train: Training features, shape (n_samples, n_features)
14         y_train: Training labels, shape (n_samples,)
15         params: Dictionary containing:
16             - n_estimators: Number of trees
17             - max_depth: Maximum tree depth
18             - random_state: Random seed
19
20     Returns:
21         Trained RandomForestClassifier instance
22
23     Raises:
24         ValueError: If X_train and y_train shapes don't match
25
26     Example:
27         >>> params = {'n_estimators': 100, 'max_depth': 10}
28         >>> model = train_model(X_train, y_train, params)
29         """
30     # Implementation
31     pass
```

## 19 Comprehensive Troubleshooting Guide

### 19.1 DVC Issues

#### 19.1.1 Issue 1: Failed to Push Data

Error Message:

```
ERROR: failed to push data to the cloud -  
[Errno 2] No such file or directory: 'S3'
```

Causes:

1. Remote storage path doesn't exist
2. Incorrect remote URL configuration
3. Permission issues
4. Network connectivity problems

Solutions:

```
# Step 1: Verify remote configuration  
dvc remote list -v  
  
# Step 2: Check if remote path exists  
# For local directory:  
ls /path/to/S3  
  
# For AWS S3:  
aws s3 ls s3://bucket-name/  
  
# Step 3: Update remote URL if needed  
dvc remote modify dvc_origin url <correct-path>  
  
# Step 4: Verify permissions  
# Check read/write access to remote location  
  
# Step 5: Try verbose push  
dvc push -v
```

#### 19.1.2 Issue 2: DVC Pull Fails

Error Message:

```
ERROR: failed to pull data - file not found in remote
```

Causes:

1. Data was never pushed to remote
2. Remote configuration changed
3. Cache corruption

Solutions:

```
# Solution 1: Push from original machine
# (On machine with data)
dvc push

# Solution 2: Check cloud status
dvc status -c

# Solution 3: Clear and re-pull cache
rm -rf .dvc/cache
dvc fetch
dvc checkout

# Solution 4: Verify remote accessibility
dvc remote list -v
```

### 19.1.3 Issue 3: Pipeline Won't Re-run

**Symptom:**

Stage 'model\_building' didn't change, skipping  
Data and pipelines are up to date.

**Cause:** DVC detects no changes (this is actually correct behavior!)

**Solutions:**

```
# Force re-run entire pipeline
dvc repro --force

# Force re-run specific stage only
dvc repro --force model_building

# Force stage and all downstream stages
dvc repro --force --downstream model_building

# Remove specific stage output to trigger re-run
dvc remove model_building.dvc
dvc repro
```

### 19.1.4 Issue 4: Checksum Mismatch

**Error Message:**

ERROR: checksum mismatch for 'data/raw/train.csv'

**Cause:** File modified outside DVC or cache corruption

**Solutions:**

```
# Solution 1: Restore from remote
dvc checkout data/raw/train.csv

# Solution 2: Clear cache and re-pull
rm -rf .dvc/cache
dvc pull
```

```
# Solution 3: If file was intentionally modified
dvc add data/raw/train.csv
git add data/raw/train.csv.dvc
git commit -m "Update data"
```

## 19.2 MLflow Issues

### 19.2.1 Issue 5: MLflow Not Logging Runs

**Symptom:** Script completes but no runs appear in DagShub

**Diagnosis and Solutions:**

1. Check DagShub Initialization:

```
1 import dagshub
2
3 # Must be at top of script
4 dagshub.init(
5     repo_owner='username',
6     repo_name='project-name',
7     mlflow=True
8 )
9
```

2. Verify Tracking URI:

```
1 import mlflow
2
3 # Check tracking URI
4 print(mlflow.get_tracking_uri())
5
6 # Should output:
7 # https://dagshub.com/username/project.mlflow
8
9 # If incorrect, set it:
10 mlflow.set_tracking_uri(
11     "https://dagshub.com/username/project.mlflow"
12 )
13
```

3. Check Internet Connection:

```
# Test DagShub accessibility
curl https://dagshub.com

# Check if blocked by firewall
ping dagshub.com
```

4. Verify MLflow Start Run:

```
1 # Ensure you're using context manager
2 with mlflow.start_run():
3     mlflow.log_metric('accuracy', 0.95)
4     mlflow.log_param('n_estimators', 100)
5
```

```

6 # Run gets finalized when exiting context
7

```

### 19.2.2 Issue 6: Cannot Find Run ID

**Problem:** Need run ID for tagging but don't know it

**Solution:**

```

1 import mlflow
2 from mlflow.tracking import MlflowClient
3
4 client = MlflowClient()
5
6 # Get experiment by name
7 experiment = mlflow.get_experiment_by_name(
8     "MLOPS-Project-with-Random-Forest"
9 )
10
11 # List all runs in experiment
12 runs = client.search_runs(
13     experiment_ids=[experiment.experiment_id],
14     order_by=["metrics.accuracy DESC"] # Sort by accuracy
15 )
16
17 # Print run IDs with metrics
18 for run in runs:
19     print(f"Run ID: {run.info.run_id}")
20     print(f"Metrics: {run.data.metrics}")
21     print(f"Params: {run.data.params}")
22     print("---")

```

## 19.3 Git Issues

### 19.3.1 Issue 7: dvc.lock Merge Conflicts

**Symptom:**

CONFLICT (content): Merge conflict in dvc.lock

**Cause:** Multiple people modified pipeline simultaneously

**Solutions:**

```

# Option 1: Accept their version
git checkout --theirs dvc.lock
dvc repro
git add dvc.lock
git commit -m "Resolve dvc.lock conflict"

# Option 2: Accept your version
git checkout --ours dvc.lock
dvc repro
git add dvc.lock
git commit -m "Resolve dvc.lock conflict"

# Option 3: Re-run pipeline (recommended)

```

```
git checkout --theirs dvc.lock
dvc repro --force
git add dvc.lock
git commit -m "Resolve conflict by re-running pipeline"
```

### 19.3.2 Issue 8: Large Files Committed to Git

**Symptom:**

```
remote: error: File data/train.csv is 150 MB;
this exceeds GitHub's file size limit
```

**Cause:** Data files not in .gitignore, tracked by Git instead of DVC

**Solutions:**

```
# Step 1: Remove from Git tracking
git rm --cached data/train.csv

# Step 2: Add to .gitignore
echo "data/" >> .gitignore

# Step 3: Track with DVC
dvc add data/train.csv

# Step 4: Commit
git add data/train.csv.dvc .gitignore
git commit -m "Move data tracking from Git to DVC"

# If already pushed to Git:
# Use git filter-branch or BFG Repo-Cleaner
git filter-branch --tree-filter 'rm -f data/train.csv' HEAD
```

## 19.4 Environment Issues

### 19.4.1 Issue 9: Import Errors

**Error Message:**

```
ImportError: No module named 'pandas'
ModuleNotFoundError: No module named 'sklearn'
```

**Solutions:**

```
# Verify virtual environment is activated
which python    # Should show venv path
where python    # Windows

# Reinstall requirements
pip install -r requirements.txt

# Check installed packages
pip list | grep pandas

# If specific package missing
```

```

pip install pandas scikit-learn

# If version conflicts
pip install --upgrade pip
pip install -r requirements.txt --force-reinstall

```

#### 19.4.2 Issue 10: NLTK Data Not Found

**Error Message:**

LookupError: Resource 'stopwords' not found

**Solution:**

```

1 import nltk
2
3 # Download required NLTK data
4 nltk.download('stopwords')
5 nltk.download('punkt')
6 nltk.download('punkt_tab')
7
8 # Verify download
9 from nltk.corpus import stopwords
10 print(len(stopwords.words('english'))) # Should print ~150
11
12 # Alternative: Download all NLTK data
13 nltk.download('all') # Warning: Large download!

```

### 19.5 Performance Issues

#### 19.5.1 Issue 11: Slow DVC Operations

**Symptom:** Very slow push/pull operations

**Solutions:**

```

# Use parallel jobs for faster transfers
dvc push -j 4    # Use 4 parallel jobs
dvc pull -j 8    # Use 8 parallel jobs

# Use hardlinks/symlinks for faster cache
dvc config cache.type hardlink,symlink

# Move cache to faster storage (SSD)
dvc cache dir /path/to/ssd/cache

# For AWS S3, tune buffer size
dvc remote modify myremote buffer_size 10485760

```

#### 19.5.2 Issue 12: Large Cache Size

**Symptom:** .dvc/cache/ consuming too much disk space

**Solutions:**

```
# Remove unused cache files
dvc gc --workspace    # Keep only workspace files
dvc gc --cloud         # Keep only files in remote

# Check cache size
du -sh .dvc/cache

# Move cache to external drive
dvc cache dir /mnt/external/cache

# Use shared cache for team
dvc config cache.dir /shared/team/cache
dvc config cache.shared group
```

## 20 Quick Reference Guide

### 20.1 Essential Commands

#### DVC Commands

```
dvc init                      # Initialize DVC
dvc remote add -d name url    # Add default remote
dvc remote list                 # List remotes
dvc remote list -v              # List with URLs
dvc add <file>                 # Track file
dvc push                        # Upload to remote
dvc pull                        # Download from remote
dvc fetch                        # Download without checkout
dvc checkout                     # Update workspace
dvc repro                        # Run pipeline
dvc repro --force               # Force re-run
dvc status                       # Check status
dvc status -c                   # Compare with cloud
dvc diff                         # Show differences
dvc dag                          # Visualize pipeline
dvc gc                           # Garbage collect cache
```

#### Git Commands

```
git init                        # Initialize repository
git add .                        # Stage all changes
git add <file>                  # Stage specific file
git commit -m "message"         # Commit changes
git push origin main            # Push to remote
git pull origin main            # Pull from remote
git status                       # Check status
git log --oneline               # View commit history
git checkout <commit>           # Switch to commit
git checkout main                # Return to main
git branch                       # List branches
git diff                         # Show differences
```

## MLflow Commands

```
# In Python code:
import mlflow

mlflow.set_tracking_uri(url)      # Set tracking server
mlflow.set_experiment(name)       # Set experiment
mlflow.start_run()                # Start tracking run
mlflow.log_metric(key, value)     # Log metric
mlflow.log_param(key, value)      # Log parameter
mlflow.log_artifact(path)         # Log file
mlflow.end_run()                  # End run

# Using context manager (recommended):
with mlflow.start_run():
    mlflow.log_metric('accuracy', 0.95)
    mlflow.log_param('n_estimators', 100)
```

## 20.2 Complete Workflow Summary

## Phase 1: Experimentation

```
# 1. Modify parameters
vim params.yaml

# 2. Run pipeline
dvc repro

# 3. Check MLflow UI for results
# https://dagshub.com/username/project.mlflow

# 4. Repeat steps 1-3 for multiple experiments

# 5. Compare all runs in MLflow UI
# Choose best performing model
```

## Phase 2: Freezing Best Model

```
# 6. Set params.yaml to winning configuration
# (Copy parameters from MLflow UI)

# 7. Verify reproducibility
dvc repro
cat reports/metrics.json

# 8. Freeze the decision
git add .
git commit -m "Freeze best model (AUC=0.918, acc=0.937)"
git push origin main

# 9. Push data and models
```

```
dvc push

# 10. Tag MLflow run (optional)
python tag_run.py
```

### Phase 3: Reproduction on New Machine

```
# 11. Clone repository
git clone https://github.com/username/project.git
cd project

# 12. Setup environment
python -m venv venv
source venv/bin/activate # or venv\Scripts\activate
pip install -r requirements.txt

# 13. Configure DVC remote (if needed)
dvc remote modify dvc_origin url <shared-path>

# 14. Pull data and models
dvc pull

# 15. Verify everything is ready
dvc status          # Should say "up to date"
dvc repro           # Should use cache (no re-run)
```

## 20.3 Project Structure Reference

```
MLOPS-Dagshub-DVC-Git-Project/
|
+-- .dvc/                      # DVC configuration
|   +-- cache/                  # Local data cache
|   +-- config                  # Remote storage config
|
+-- .git/                       # Git repository
|
+-- S3/                         # Local DVC remote (gitignored)
|
+-- src/                        # Source code
|   +-- Data_Ingestion.py
|   +-- Data_Pre_Processing.py
|   +-- Feature_Engineering.py
|   +-- Model_Building_Rf.py
|   +-- Model_Evaluation_Rf.py
|
+-- data/                        # Data files (DVC tracked)
|   +-- raw/
|   +-- interim/
|   +-- processed/
```

```
+-- models/                                # Model files (DVC tracked)
|   +- model_rf.pkl
|
+-- reports/                               # Reports (DVC tracked)
|   +- metrics.json
|
+-- logs/                                  # Log files (gitignored)
|   +- *.log
|
+-- dvc.yaml                               # Pipeline definition
+-- dvc.lock                                # Pipeline lock file
--- params.yaml                            # Hyperparameters
--- requirements.txt                      # Python dependencies
--- .gitignore                             # Git ignore rules
--- .dvcignore                            # DVC ignore rules
--- README.md                              # Documentation
```

## 20.4 Key Files Explained

File	Purpose
dvc.yaml	Defines pipeline stages, dependencies, outputs, parameters
dvc.lock	Records exact file versions (MD5 hashes) for reproducibility
params.yaml	Centralized hyperparameter configuration
.gitignore	Tells Git what NOT to track (data/, models/, logs/)
.dvc/config	DVC remote storage configuration
.dvc/cache/	Local cache storing all data versions
requirements.txt	Python package dependencies with versions

## 20.5 Common Command Combinations

### Frequently Used Sequences

```
# Complete experiment cycle
dvc repro && dvc push && git add . && \
git commit -m "Experiment X" && git push

# Quick status check
dvc status && git status

# Full pipeline re-run with push
dvc repro --force && dvc push && \
git add dvc.lock && git commit -m "Re-run pipeline" && git push

# Sync with team
git pull && dvc pull

# Clean up experiments and cache
dvc gc --workspace && dvc gc --cloud

# View pipeline and check status
dvc dag && dvc status
```

## 21 Comparison: DVC-Only vs DVC+MLflow+Git Workflow

### 21.1 Architecture Comparison

DVC-Only Workflow	DVC + MLflow + Git Workflow
<b>Code Versioning</b>	
Git for code versioning	Git for code versioning
<b>Data &amp; Model Versioning</b>	
DVC for data/model versioning	DVC for data/model versioning
<b>Experiment Tracking</b>	
Local metrics.json file	Centralized MLflow server (DagShub)
Manual experiment comparison	Visual comparison in MLflow UI
Private to individual user	Shared across entire team
Limited experiment history	Complete experiment history
No run tagging	Git commit tagging support
<b>Model Registry</b>	
Manual model management	MLflow model registry
No deployment tracking	Deployment status tracking

### 21.2 Feature-by-Feature Comparison

Feature	DVC-Only	DVC+MLflow
Code versioning	✓	✓
Data versioning	✓	✓
Model versioning	✓	✓
Pipeline automation	✓	✓
Parameter tracking	✓	✓
Reproducibility	✓	✓
Team collaboration	Limited	✓
Visual metric comparison	✗	✓
Centralized tracking	✗	✓
Experiment history	Limited	✓
Run tagging	✗	✓
Model registry	✗	✓
Real-time monitoring	✗	✓
Cross-model comparison	Manual	Automatic
Deployment tracking	✗	✓

### 21.3 Workflow Process Comparison

#### 21.3.1 Experimentation Phase

DVC-Only	DVC + MLflow
1. Modify params.yaml	1. Modify params.yaml
2. Run <code>dvc repro</code>	2. Run <code>dvc repro</code>
3. Check metrics.json manually	3. View results in MLflow UI
4. Copy metrics to spreadsheet	4. Compare visually in UI
5. Repeat for each experiment	5. All runs automatically tracked
6. Manually compare results	6. Sort/filter by any metric
7. Choose best from spreadsheet	7. Click to view best run

### 21.3.2 Model Freezing Phase

DVC-Only	DVC + MLflow
<ol style="list-style-type: none"> <li>1. Identify best params manually</li> <li>2. Update params.yaml</li> <li>3. Run <code>dvc repro</code></li> <li>4. <code>git commit</code></li> <li>5. <code>dvc push</code></li> <li>6. No linking to experiments</li> </ol>	<ol style="list-style-type: none"> <li>1. Click best run in MLflow UI</li> <li>2. Copy parameters from UI</li> <li>3. Update params.yaml</li> <li>4. Run <code>dvc repro</code></li> <li>5. <code>git commit</code></li> <li>6. <code>dvc push</code></li> <li>7. Tag MLflow run with Git commit</li> <li>8. Mark as "production" in MLflow</li> </ol>

## 21.4 Team Collaboration Comparison

Aspect	DVC-Only	DVC + MLflow
<b>Experiment Visibility</b>	Private/Local	Centralized/Shared
<b>Metric Sharing</b>	Manual (email/Slack)	Automatic (MLflow UI)
<b>Best Model Decision</b>	Individual/unclear	Transparent/documentated
<b>Experiment History</b>	Lost after cleanup	Permanently preserved
<b>Cross-team Learning</b>	Difficult	Easy
<b>Onboarding New Members</b>	Manual explanation	Self-service in UI
<b>Avoiding Duplicate Work</b>	Hard to coordinate	Clear visibility

## 21.5 Use Case Recommendations

### 21.5.1 When to Use DVC-Only

- **Solo projects:** Individual research or learning
- **Simple pipelines:** Few experiments, straightforward workflows
- **Limited resources:** No budget for hosted services
- **Privacy constraints:** Cannot use external services
- **Prototype phase:** Early exploration, not production-ready
- **Academic research:** Focus on reproducibility over collaboration

### 21.5.2 When to Use DVC + MLflow + Git

- **Team projects:** Multiple data scientists collaborating
- **Production systems:** Models deployed to production
- **Complex experiments:** Many hyperparameter variations
- **Model comparison:** Testing multiple model architectures
- **Compliance needs:** Audit trails required
- **Continuous improvement:** Ongoing model iterations
- **Knowledge sharing:** Cross-functional teams
- **Model monitoring:** Track production model performance

## 21.6 Real-World Scenario Comparison

### 21.6.1 Scenario: Finding Best Model from 20+ Experiments

#### DVC-Only Approach

**Process:**

1. Run 20 experiments over several days
2. Copy metrics from each metrics.json to Excel
3. Manually create comparison charts
4. Search through Git commits to find parameters
5. Risk: Easy to lose track of which experiment was which
6. Risk: metrics.json overwritten between experiments

**Time:** 30+ minutes of manual work**Error prone:** Yes, easy to lose experiment data

#### DVC + MLflow Approach

**Process:**

1. Run 20 experiments (automatically logged)
2. Open MLflow UI
3. Click "Compare" and select all runs
4. Sort by any metric (e.g., AUC)
5. View parallel coordinates plot
6. Click best run to see all details

**Time:** 2 minutes**Error prone:** No, all data preserved automatically

### 21.6.2 Scenario: Reproducing 6-Month-Old Model

#### DVC-Only Approach

##### Challenges:

- Find Git commit (if tagged)
- No experiment metadata preserved
- No link to original metrics
- Manual verification needed
- Risk: Was this really the production model?

**Confidence:** Moderate

#### DVC + MLflow Approach

##### Process:

1. Search MLflow for "status:production"
2. Find run with "deployment\_date: 2024-06"
3. See Git commit hash in tags
4. View exact metrics from that run
5. `git checkout <commit_hash>`
6. `dvc checkout`
7. Verified reproduction

**Confidence:** High

## 21.7 Cost-Benefit Analysis

Factor	DVC-Only	DVC + MLflow
<b>Setup Time</b>	30 minutes	45-60 minutes
<b>Learning Curve</b>	Moderate	Steeper initially
<b>Per-Experiment Overhead</b>	None	Minimal (auto-logged)
<b>Experiment Comparison Time</b>	15-30 minutes	2-5 minutes
<b>Team Onboarding</b>	Manual explanation	Self-service UI
<b>Maintenance</b>	Low	Low (hosted on DagShub)
<b>Cost (DagShub Free Tier)</b>	\$0	\$0
<b>Scalability</b>	Limited	Excellent
<b>ROI for 1-person project</b>	Good	Moderate
<b>ROI for 3+ person team</b>	Poor	Excellent

## 21.8 Migration Path

### 21.8.1 Upgrading from DVC-Only to DVC+MLflow

If you already have a DVC-only project, migration is straightforward:

1. **Install additional dependencies:**

```
pip install mlflow dagshub
```

2. **Add MLflow logging to evaluation script:**

```
1 # At top of Model_Evaluation.py
2 import mlflow
3 import dagshub
4
5 dagshub.init(repo_owner='...', repo_name='...', mlflow=True)
6 mlflow.set_tracking_uri("...")
7 mlflow.set_experiment("...")
8
9 # In main function, add:
10 with mlflow.start_run():
11     mlflow.log_metric('accuracy', accuracy)
12     mlflow.log_param('n_estimators', params['n_estimators'])
13     # ... log all metrics and parameters
14
```

3. **No changes needed to:**

- DVC pipeline (dvc.yaml)
- Data versioning
- Git workflow
- Other pipeline stages

4. **Run pipeline:**

```
dvc repro
```

5. **Verify in MLflow UI:** Check that run appears

6. **Commit changes:**

```
git add .
git commit -m "Add MLflow experiment tracking"
git push
```

**Result:** All future experiments automatically tracked in MLflow, while maintaining all DVC benefits!

## 22 Conclusion and Next Steps

### 22.1 What You've Accomplished

By completing this guide, you now have a production-ready MLOps pipeline that:

#### Core Achievements

##### Version Control:

- ✓ Code versioned with Git
- ✓ Data versioned with DVC
- ✓ Models versioned with DVC
- ✓ Complete reproducibility guaranteed

##### Experiment Management:

- ✓ Centralized tracking with MLflow
- ✓ Visual metric comparison
- ✓ Complete experiment history
- ✓ Production model clearly identified

##### Team Collaboration:

- ✓ Shared experiment visibility
- ✓ Clear communication channel (DagShub)
- ✓ Avoid duplicate work
- ✓ Knowledge sharing enabled

##### Production Readiness:

- ✓ Automated pipeline execution
- ✓ Deployment tracking capability
- ✓ Audit trail for compliance
- ✓ Rollback capability to any version

### 22.2 Key Takeaways

#### 22.2.1 The Three-System Architecture

##### Remember the Triangle

**Git** → Code + Metadata

**DVC** → Data + Models

**MLflow** → Experiments + Metrics

Each system has a specific purpose. Together, they provide complete MLOps infrastructure.

### 22.2.2 Critical Best Practices

#### 1. Experiment freely, freeze deliberately:

- Run as many experiments as needed
- Only commit when you have a decision
- Tag production runs explicitly

#### 2. Always push artifacts:

- `git push` for code
- `dvc push` for data/models
- Both are necessary for team work

#### 3. Link everything together:

- Git commits contain code + metadata
- DVC tracks exact data versions
- MLflow runs tagged with Git commits
- Complete traceability

#### 4. Document your decisions:

- Clear commit messages
- Experiment notes in MLflow
- README updates
- Why you chose this model

## 22.3 Common Pitfalls to Avoid

### Warning

#### Watch Out For:

##### 1. Forgetting to push DVC artifacts:

- Always run `dvc push` after committing
- Team cannot reproduce without artifacts
- Set up reminder or automation

##### 2. Not linking MLflow runs to Git commits:

- Lost connection between code and experiments
- Hard to reproduce later
- Tag runs immediately

##### 3. Committing data files to Git:

- Verify `.gitignore` before committing

- Data should only be in DVC
  - Check `git status` carefully
4. **Modifying params.yaml without re-running:**
- Changes not reflected until `dvc repro`
  - `dvc.lock` becomes out of sync
  - Always repro after param changes
5. **Losing local DVC cache:**
- Cache is critical for speed
  - Back up cache directory
  - Or ensure remote is always accessible

## 22.4 Next Steps and Extensions

### 22.4.1 Immediate Next Steps

#### 1. Add more models:

- XGBoost classifier
- Logistic Regression baseline
- Neural network (LSTM/BERT)
- Compare in MLflow UI

#### 2. Improve monitoring:

```

1 # Log additional metrics
2 mlflow.log_metric('train_time', train_time)
3 mlflow.log_metric('inference_time', inference_time)
4 mlflow.log_metric('model_size_mb', model_size)
5

```

#### 3. Add data validation:

- Check data quality in ingestion
- Validate schema consistency
- Alert on data drift

#### 4. Enhance visualizations:

```

1 # Log confusion matrix
2 import matplotlib.pyplot as plt
3 from sklearn.metrics import ConfusionMatrixDisplay
4
5 cm = confusion_matrix(y_test, y_pred)
6 disp = ConfusionMatrixDisplay(cm)
7 disp.plot()
8 plt.savefig('confusion_matrix.png')
9 mlflow.log_artifact('confusion_matrix.png')
10

```

### 22.4.2 Advanced Extensions

#### 1. CI/CD Integration:

- GitHub Actions for automated testing
- Automatic pipeline runs on push
- Model validation gates
- Automated deployment to staging

#### 2. Model Registry:

```

1 # Register model in MLflow
2 mlflow.sklearn.log_model(
3     model,
4     "spam_classifier",
5     registered_model_name="SpamDetector"
6 )
7
8 # Transition to production
9 from mlflow.tracking import MlflowClient
10 client = MlflowClient()
11 client.transition_model_version_stage(
12     name="SpamDetector",
13     version=1,
14     stage="Production"
15 )
16

```

#### 3. Hyperparameter Optimization:

```

1 import optuna
2
3 def objective(trial):
4     n_estimators = trial.suggest_int('n_estimators', 10, 100)
5     max_depth = trial.suggest_int('max_depth', 5, 20)
6
7     # Train and evaluate
8     model = train_model(n_estimators, max_depth)
9     score = evaluate_model(model)
10
11     # Log to MLflow
12     with mlflow.start_run(nested=True):
13         mlflow.log_params(trial.params)
14         mlflow.log_metric('accuracy', score)
15
16     return score
17
18 study = optuna.create_study(direction='maximize')
19 study.optimize(objective, n_trials=50)
20

```

#### 4. Production Deployment:

- Deploy model as REST API (FastAPI)
- Container with Docker
- Kubernetes orchestration

- Model serving with MLflow

## 5. Monitoring and Alerting:

- Track prediction distribution
- Monitor model performance drift
- Alert on accuracy degradation
- Trigger retraining pipeline

## 22.5 Learning Resources

### 22.5.1 Official Documentation

- DVC: <https://dvc.org/doc>
- MLflow: <https://mlflow.org/docs/latest/index.html>
- DagShub: <https://dagshub.com/docs>
- Git: <https://git-scm.com/doc>

### 22.5.2 Community and Support

- DVC Discord: Active community for questions
- MLflow Slack: Official Slack workspace
- DagShub Community: Forum and discussions
- Stack Overflow: Tagged questions for all tools

### 22.5.3 Advanced Topics to Explore

#### 1. MLOps Maturity Models:

- Google's MLOps maturity framework
- Microsoft's ML lifecycle
- AWS well-architected ML

#### 2. Feature Stores:

- Feast (open source)
- Tecton
- AWS SageMaker Feature Store

#### 3. Model Monitoring:

- Evidently AI
- WhyLabs
- Arize AI

#### 4. Experiment Orchestration:

- Kubeflow Pipelines
- Apache Airflow
- Prefect

## 22.6 Final Thoughts

## The MLOps Journey

This guide represents a **production-ready foundation** for MLOps, but MLOps is a journey, not a destination.

### You now have:

- A reproducible workflow
- Team collaboration capability
- Experiment tracking infrastructure
- Clear path to production

### Continue improving:

- Add more automation
- Enhance monitoring
- Improve documentation
- Share knowledge with team

**Remember:** The best MLOps system is one that your team actually uses. Start simple, iterate, and scale as needed.

## 22.7 Project Checklist

Use this checklist to verify your MLOps implementation:

Component	Status
<b>Version Control</b>	
Git repository initialized	✓
Code properly committed	✓
.gitignore configured	✓
Remote repository connected	✓
<b>Data Versioning</b>	
DVC initialized	✓
Data tracked with DVC	✓
DVC remote configured	✓
Artifacts pushed to remote	✓
<b>Pipeline</b>	
dvc.yaml created	✓
All stages defined	✓
params.yaml configured	✓
Pipeline runs successfully	✓
dvc.lock committed	✓
<b>Experiment Tracking</b>	
MLflow integrated	✓
DagShub connected	✓
Metrics logged	✓
Parameters logged	✓
Multiple experiments run	✓
<b>Best Model</b>	
Best model identified	✓
Configuration frozen	✓
Git commit created	✓
DVC artifacts pushed	✓
MLflow run tagged	✓
<b>Documentation</b>	
README updated	✓
Commit messages clear	✓
Experiment notes recorded	✓
<b>Reproducibility</b>	
Tested on new machine	✓
Team can reproduce	✓
All dependencies listed	✓

## 22.8 Get In Touch

Questions or Feedback?

Email: sujil9480@gmail.com

*This guide is continuously updated. Check the repository for the latest version.*

Version 1.0 - December 2025